

Problem Set Lab 12, Dec 25, 2025 (Graph Embeddings)

Goals. The goals of this exercise are to:

- Understand the core components and principles of the Graph Encoder-Decoder Model (GraphEDM).
- Explain the mechanisms behind shallow unsupervised graph embedding methods like DeepWalk.
- Describe the functioning of spatial GNNs, particularly GraphSAGE, for inductive learning.
- Understand how deep unsupervised models like (V)GAE learn graph embeddings.
- Compare and contrast different graph embedding approaches for a specific downstream task.

Submission instructions:

- Please submit a PDF file to canvas.

Problem 1 (Label Propagation):

Consider the undirected graph $G = (V, E)$ with $V = \{1, 2, 3, 4\}$. The unweighted adjacency matrix A is given by:

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Nodes 1 and 4 are labeled: $y_1 = +1$ and $y_4 = -1$. Nodes 2 and 3 are unlabeled. We initialize the soft labels for unlabeled nodes as $f_2^{(0)} = 0$ and $f_3^{(0)} = 0$. The labels for labeled nodes are fixed: $f_1^{(t)} = y_1 = +1$ and $f_4^{(t)} = y_4 = -1$ for all iterations t .

1. Construct the row-normalized transition matrix P . Recall $P_{ij} = \frac{W_{ij}}{\sum_k W_{ik}}$, where W is the (weighted) adjacency matrix. Here, use A as W .
2. Perform one iteration of label propagation for the unlabeled nodes $i \in \{2, 3\}$:

$$f_i^{(t+1)} = \sum_j P_{ij} f_j^{(t)}$$

Compute $f_2^{(1)}$ and $f_3^{(1)}$.

Problem 2 (Shallow Embeddings: Matrix Factorization):

Consider a graph with $N = 3$ nodes. We have learned low-dimensional embeddings $Z \in \mathbb{R}^{N \times L}$ for these nodes using a shallow embedding method. Let $L = 2$ and the learned embedding matrix be:

$$Z = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 1 & 1 \end{pmatrix}$$

We are using an outer product-based decoder to reconstruct the similarity matrix $\hat{W} = ZZ^T$. The original (ground truth) similarity matrix (e.g., adjacency matrix) is:

$$W = \begin{pmatrix} 0 & 1 & 1 \\ 1 & 0 & 1 \\ 1 & 1 & 0 \end{pmatrix}$$

1. Compute the reconstructed similarity matrix $\hat{W} = ZZ^T$.
2. Compute the reconstruction loss $L_{G,RECON} = \|W - \hat{W}\|_F^2$ (squared Frobenius norm of the difference).

Problem 3 (Graph Convolutional Networks (GCN)):

A common form for a GCN layer is:

$$H^{(l+1)} = \sigma(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)})$$

where $\tilde{A} = A + I_N$ is the adjacency matrix with self-loops, \tilde{D} is the diagonal degree matrix of \tilde{A} (i.e., $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$), $H^{(l)}$ is the matrix of node features/embeddings at layer l , $W^{(l)}$ is the learnable weight matrix for layer l , and σ is an activation function.

Let A be the adjacency matrix from Problem 1:

$$A = \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Let $H^{(0)} = X$ be the initial node feature matrix.

1. Compute $\tilde{A} = A + I_N$.
2. Compute the degree matrix \tilde{D} for \tilde{A} .
3. Explain the role of the term $\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$ in the GCN layer. What does multiplying by $H^{(l)} W^{(l)}$ achieve?

Problem 4 (DeepWalk Intuition):

DeepWalk is a shallow graph embedding method inspired by Word2Vec from Natural Language Processing.

1. Briefly explain how DeepWalk generates "sentences" from a graph.
2. How does DeepWalk then use these "sentences" to learn node embeddings, drawing an analogy to Word2Vec? What is the objective function typically optimizing?

Problem 5 (Graph Autoencoder (GAE)):

A Graph Autoencoder (GAE) uses a GCN as an encoder and a simple inner product decoder for link prediction. The encoder is typically a two-layer GCN:

$$Z = \text{GCN}(X, A) = \tilde{A} \text{ReLU}(\tilde{A} X W^{(0)}) W^{(1)}$$

(Here, \tilde{A} is used as shorthand for the normalized adjacency matrix $\hat{A} = \tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$ for simplicity in the formula structure, though different normalization schemes can be used). The decoder reconstructs the adjacency matrix \hat{A} from the embeddings Z :

$$\hat{A}_{ij} = \sigma(Z_i Z_j^T)$$

where Z_i is the embedding of node i and σ is the sigmoid function. The reconstruction loss is often the binary cross-entropy:

$$L_{G,RECON} = - \sum_{i,j} (A_{ij} \log \hat{A}_{ij} + (1 - A_{ij}) \log (1 - \hat{A}_{ij}))$$

1. What is the primary purpose of the GCN encoder in the GAE framework? What information does it aim to capture in the embeddings Z ?
2. Explain why an inner product decoder ($Z_i Z_j^T$) followed by a sigmoid is a suitable choice for reconstructing the adjacency matrix.
3. What does the binary cross-entropy loss penalize in this context?