

Lecture 5: Linear Models for Classification

Tao LIN

SoE, Westlake University

March 24, 2025



1 Review of Last Week

- Data Model and Learning Algorithm
- Generalization Gap and Model Selection
- Bias-Variance Decomposition

2 Logistic Regression

- Classification
- Logistic Regression

Reading materials & Reference

Reading materials:

- Chapter 3, Stanford CS 229 Lecture Notes,
https://cs229.stanford.edu/notes2022fall/main_notes.pdf
- Lecture 4, Stanford CS 231n, <http://cs231n.stanford.edu/schedule.html>

Reference:

- EPFL, CS-433 Machine Learning, https://github.com/epfml/ML_course

Table of Contents

1 Review of Last Week

- Data Model and Learning Algorithm
- Generalization Gap and Model Selection
- Bias-Variance Decomposition

2 Logistic Regression

Table of Contents

1 Review of Last Week

- Data Model and Learning Algorithm
- Generalization Gap and Model Selection
- Bias-Variance Decomposition

2 Logistic Regression

- Classification
- Logistic Regression

How do we trade-off the under-fitting and over-fitting?

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$
- We see a dataset S of independent samples from \mathcal{D} :

$$S = \{(\mathbf{x}_n, y_n) \mid \text{i.i.d.} \sim \mathcal{D}\}_{n=1}^N. \quad (1)$$

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$
- We see a dataset S of independent samples from \mathcal{D} :

$$S = \{(\mathbf{x}_n, y_n) \text{ i.i.d. } \sim \mathcal{D}\}_{n=1}^N. \quad (1)$$

Learning algorithm:

$$f_S = \mathcal{A}(S) \quad (2)$$

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$
- We see a dataset S of independent samples from \mathcal{D} :

$$S = \{(\mathbf{x}_n, y_n) \text{ i.i.d. } \sim \mathcal{D}\}_{n=1}^N. \quad (1)$$

Learning algorithm:

$$f_S = \mathcal{A}(S) \quad (2)$$

- f_S denotes the output.

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$
- We see a dataset S of independent samples from \mathcal{D} :

$$S = \{(\mathbf{x}_n, y_n) \text{ i.i.d. } \sim \mathcal{D}\}_{n=1}^N. \quad (1)$$

Learning algorithm:

$$f_S = \mathcal{A}(S) \quad (2)$$

- f_S denotes the output.
- \mathcal{A} denotes the learning algorithm.

Probabilistic setup

Data model:

- We assume an (unknown) underlying distribution \mathcal{D} , with range $\mathcal{X} \times \mathcal{Y}$
- We see a dataset S of independent samples from \mathcal{D} :

$$S = \{(\mathbf{x}_n, y_n) \mid \text{i.i.d.} \sim \mathcal{D}\}_{n=1}^N. \quad (1)$$

Learning algorithm:

$$f_S = \mathcal{A}(S) \quad (2)$$

- f_S denotes the output.
- \mathcal{A} denotes the learning algorithm.
- S denotes the input (dataset).

True Error, Empirical Error, and Training Error

- **True Error** (the **expected error** over all samples chosen according to \mathcal{D}):

$$\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{\mathcal{D}} [\ell(y, f(\mathbf{x}))] . \quad (3)$$

We cannot estimate it due to the unknown \mathcal{D} .

True Error, Empirical Error, and Training Error

- **True Error** (the **expected error** over all samples chosen according to \mathcal{D}):

$$\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{\mathcal{D}} [\ell(y, f(\mathbf{x}))] . \quad (3)$$

We cannot estimate it due to the unknown \mathcal{D} .

- **Empirical Error** (what we can compute—the approximated true error via dataset S):

$$L_S(f) = \frac{1}{|S|} \sum_{(\mathbf{x}_n, y_n) \in S} \ell(y_n, f(\mathbf{x}_n)) . \quad (4)$$

True Error, Empirical Error, and Training Error

- **True Error** (the **expected error** over all samples chosen according to \mathcal{D}):

$$\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{\mathcal{D}} [\ell(y, f(\mathbf{x}))] . \quad (3)$$

We cannot estimate it due to the unknown \mathcal{D} .

- **Empirical Error** (what we can compute—the approximated true error via dataset S):

$$L_S(f) = \frac{1}{|S|} \sum_{(\mathbf{x}_n, y_n) \in S} \ell(y_n, f(\mathbf{x}_n)) . \quad (4)$$

- **Training Error** (what we are minimizing—when the model has been trained on the same data it is applied to):

$$L_S(f_S) = \frac{1}{|S|} \sum_{(\mathbf{x}_n, y_n) \in S} \ell(y_n, f_S(\mathbf{x}_n)) . \quad (5)$$

True Error, Empirical Error, and Training Error

- **True Error** (the **expected error** over all samples chosen according to \mathcal{D}):

$$\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{\mathcal{D}} [\ell(y, f(\mathbf{x}))] . \quad (3)$$

We cannot estimate it due to the unknown \mathcal{D} .

- **Empirical Error** (what we can compute—the approximated true error via dataset S):

$$L_S(f) = \frac{1}{|S|} \sum_{(\mathbf{x}_n, y_n) \in S} \ell(y_n, f(\mathbf{x}_n)) . \quad (4)$$

- **Training Error** (what we are minimizing—when the model has been trained on the same data it is applied to):

$$L_S(f_S) = \frac{1}{|S|} \sum_{(\mathbf{x}_n, y_n) \in S} \ell(y_n, f_S(\mathbf{x}_n)) . \quad (5)$$

The reason that $L_S(f_S)$ might not be close $L_{\mathcal{D}}(f_S)$ is of course over-fitting.

Problem: validating the model on the same data subset we trained it on!

Splitting the data and Test Error

- **Fix:** Split the data into a *training* set S_{train} and a *test* set S_{test} (a.k.a. *validation* set):

$$S = S_{\text{train}} \cup S_{\text{test}} \quad (6)$$

Splitting the data and Test Error

- **Fix:** Split the data into a *training* set S_{train} and a *test* set S_{test} (a.k.a. *validation* set):

$$S = S_{\text{train}} \cup S_{\text{test}} \quad (6)$$

- 1 We **learn** the function $f_{S_{\text{train}}}$ using the train set S_{train}

Splitting the data and Test Error

- **Fix:** Split the data into a *training* set S_{train} and a *test* set S_{test} (a.k.a. *validation* set):

$$S = S_{\text{train}} \cup S_{\text{test}} \quad (6)$$

- 1 We **learn** the function $f_{S_{\text{train}}}$ using the train set S_{train}
- 2 We **validate** it computing the error on the **test set** S_{test} :

$$L_{S_{\text{test}}}(f_{S_{\text{train}}}) = \frac{1}{|S_{\text{test}}|} \sum_{(y_n, \mathbf{x}_n) \in S_{\text{test}}} \ell(y_n, f_{S_{\text{train}}}(\mathbf{x}_n)). \quad (7)$$

Splitting the data and Test Error

- **Fix:** Split the data into a *training* set S_{train} and a *test* set S_{test} (a.k.a. *validation* set):

$$S = S_{\text{train}} \cup S_{\text{test}} \quad (6)$$

- 1 We **learn** the function $f_{S_{\text{train}}}$ using the train set S_{train}
- 2 We **validate** it computing the error on the **test set** S_{test} :

$$L_{S_{\text{test}}}(f_{S_{\text{train}}}) = \frac{1}{|S_{\text{test}}|} \sum_{(y_n, \mathbf{x}_n) \in S_{\text{test}}} \ell(y_n, f_{S_{\text{train}}}(\mathbf{x}_n)). \quad (7)$$

- Since S_{train} and S_{test} are independent, we hope that $L_{S_{\text{test}}}(f_{S_{\text{train}}}) \approx L_{\mathcal{D}}(f_{S_{\text{train}}})$

Splitting the data and Test Error

- **Fix:** Split the data into a *training* set S_{train} and a *test* set S_{test} (a.k.a. *validation* set):

$$S = S_{\text{train}} \cup S_{\text{test}} \quad (6)$$

- 1 We **learn** the function $f_{S_{\text{train}}}$ using the train set S_{train}
- 2 We **validate** it computing the error on the **test set** S_{test} :

$$L_{S_{\text{test}}}(f_{S_{\text{train}}}) = \frac{1}{|S_{\text{test}}|} \sum_{(y_n, \mathbf{x}_n) \in S_{\text{test}}} \ell(y_n, f_{S_{\text{train}}}(\mathbf{x}_n)). \quad (7)$$

- Since S_{train} and S_{test} are independent, we hope that $L_{S_{\text{test}}}(f_{S_{\text{train}}}) \approx L_{\mathcal{D}}(f_{S_{\text{train}}})$
- **Issues:** we have fewer data both for the learning and validation tasks (trade-off)

Table of Contents

1 Review of Last Week

- Data Model and Learning Algorithm
- Generalization Gap and Model Selection
- Bias-Variance Decomposition

2 Logistic Regression

- Classification
- Logistic Regression

Generalization gap: How far is the test from the true error?

- **True Error:**

$$L_{\mathcal{D}}(f) = \mathbb{E}_{(y, \mathbf{x}) \sim \mathcal{D}} [\ell(y, f(\mathbf{x}))]. \quad (8)$$

- **Test/Empirical Error:**

$$L_{S_{\text{test}}}(f) = \frac{1}{|S_{\text{test}}|} \sum_{(\mathbf{x}_n, y_n) \in S_{\text{test}}} \ell(y_n, f(\mathbf{x}_n)). \quad (9)$$

- **Generalization Error:**

$$|L_{\mathcal{D}}(f) - L_{S_{\text{test}}}(f)|. \quad (10)$$

Generalization gap: How far is the test from the true error?

- **True Error:**

$$L_{\mathcal{D}}(f) = \mathbb{E}_{(y, \mathbf{x}) \sim \mathcal{D}} [\ell(y, f(\mathbf{x}))] . \quad (8)$$

- **Test/Empirical Error:**

$$L_{S_{\text{test}}}(f) = \frac{1}{|S_{\text{test}}|} \sum_{(\mathbf{x}_n, y_n) \in S_{\text{test}}} \ell(y_n, f(\mathbf{x}_n)) . \quad (9)$$

- **Generalization Error:**

$$|L_{\mathcal{D}}(f) - L_{S_{\text{test}}}(f)| . \quad (10)$$

- **In Expectation:**

$$L_{\mathcal{D}}(f) = \mathbb{E}_{S_{\text{test}} \sim \mathcal{D}} [L_{S_{\text{test}}}(f)] , \quad (11)$$

Generalization gap: How far is the test from the true error?

- **True Error:**

$$L_{\mathcal{D}}(f) = \mathbb{E}_{(y, \mathbf{x}) \sim \mathcal{D}} [\ell(y, f(\mathbf{x}))]. \quad (8)$$

- **Test/Empirical Error:**

$$L_{S_{\text{test}}}(f) = \frac{1}{|S_{\text{test}}|} \sum_{(\mathbf{x}_n, y_n) \in S_{\text{test}}} \ell(y_n, f(\mathbf{x}_n)). \quad (9)$$

- **Generalization Error:**

$$|L_{\mathcal{D}}(f) - L_{S_{\text{test}}}(f)|. \quad (10)$$

- **In Expectation:**

$$L_{\mathcal{D}}(f) = \mathbb{E}_{S_{\text{test}} \sim \mathcal{D}} [L_{S_{\text{test}}}(f)], \quad (11)$$

The variation of the $|L_{\mathcal{D}}(f) - L_{S_{\text{test}}}(f)|$ matters.

Theorem 1

Given a model f and a test set $S_{test} \sim \mathcal{D}$ i.i.d. (not used to learn f) and a loss $\ell(\cdot, \cdot) \in [a, b]$:

$$\Pr \left[|L_{\mathcal{D}}(f) - L_{S_{test}}(f)| \geq \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2|S_{test}|}} \right] \leq \delta \quad (12)$$

Theorem 1

Given a model f and a test set $S_{test} \sim \mathcal{D}$ i.i.d. (not used to learn f) and a loss $\ell(\cdot, \cdot) \in [a, b]$:

$$\Pr \left[|L_{\mathcal{D}}(f) - L_{S_{test}}(f)| \geq \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2 |S_{test}|}} \right] \leq \delta \quad (12)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{test}|})$ with the number test points.

Theorem 1

Given a model f and a test set $S_{test} \sim \mathcal{D}$ i.i.d. (not used to learn f) and a loss $\ell(\cdot, \cdot) \in [a, b]$:

$$\Pr \left[|L_{\mathcal{D}}(f) - L_{S_{test}}(f)| \geq \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2|S_{test}|}} \right] \leq \delta \quad (12)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{test}|})$ with the number test points.

⇒ The more data points we have, the more confident we are that the empirical loss we measure is close to the true loss.

Theorem 1

Given a model f and a test set $S_{\text{test}} \sim \mathcal{D}$ i.i.d. (not used to learn f) and a loss $\ell(\cdot, \cdot) \in [a, b]$:

$$\Pr \left[|L_{\mathcal{D}}(f) - L_{S_{\text{test}}}(f)| \geq \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2 |S_{\text{test}}|}} \right] \leq \delta \quad (12)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{\text{test}}|})$ with the number test points.

⇒ The more data points we have, the more confident we are that the empirical loss we measure is close to the true loss.

Given a predictor f and a dataset S , we can control the expected risk:

$$\Pr \left[\underbrace{L_{\mathcal{D}}(f)}_{\text{not computable}} \geq \underbrace{L_{S_{\text{test}}}(f)}_{\text{computable}} + \underbrace{\sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2 |S_{\text{test}}|}}}_{\text{deviation}} \right] \leq \delta. \quad (13)$$

How far is each of the K test errors $L_{S_{\text{test}}}(f_k)$ from the true $L_{\mathcal{D}}(f_k)$?

Theorem 2

We can bound the maximum deviation for all K candidates, by

$$\Pr \left[\max_k |L_{\mathcal{D}}(f_k) - L_{S_{\text{test}}}(f_k)| \geq \sqrt{\frac{(b-a)^2 \ln(2K/\delta)}{2|S_{\text{test}}|}} \right] \leq \delta \quad (14)$$

How far is each of the K test errors $L_{S_{\text{test}}}(f_k)$ from the true $L_{\mathcal{D}}(f_k)$?

Theorem 2

We can bound the maximum deviation for all K candidates, by

$$\Pr \left[\max_k |L_{\mathcal{D}}(f_k) - L_{S_{\text{test}}}(f_k)| \geq \sqrt{\frac{(b-a)^2 \ln(2K/\delta)}{2|S_{\text{test}}|}} \right] \leq \delta \quad (14)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{\text{test}}|})$ with the number test points.

How far is each of the K test errors $L_{S_{\text{test}}}(f_k)$ from the true $L_{\mathcal{D}}(f_k)$?

Theorem 2

We can bound the maximum deviation for all K candidates, by

$$\Pr \left[\max_k |L_{\mathcal{D}}(f_k) - L_{S_{\text{test}}}(f_k)| \geq \sqrt{\frac{(b-a)^2 \ln(2K/\delta)}{2|S_{\text{test}}|}} \right] \leq \delta \quad (14)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{\text{test}}|})$ with the number test points.
- When testing K hyper-parameters, the error only goes up by $\sqrt{\ln(K)}$.

How far is each of the K test errors $L_{S_{\text{test}}}(f_k)$ from the true $L_{\mathcal{D}}(f_k)$?

Theorem 2

We can bound the maximum deviation for all K candidates, by


$$\Pr \left[\max_k |L_{\mathcal{D}}(f_k) - L_{S_{\text{test}}}(f_k)| \geq \sqrt{\frac{(b-a)^2 \ln(2K/\delta)}{2|S_{\text{test}}|}} \right] \leq \delta \quad (14)$$

- The error decreases as $\mathcal{O}(1/\sqrt{|S_{\text{test}}|})$ with the number test points.
 - When testing K hyper-parameters, the error only goes up by $\sqrt{\ln(K)}$.
- ⇒ So we can test many different models without incurring a large penalty.

Issues: Splitting the data once into two parts (one for training and one for testing)
is not the most efficient way to use the data!

K-fold cross-validation:

- 1 Randomly partition the data into K groups
- 2 Train K times. Each time leave out exactly one of the K groups for testing and use the remaining $K - 1$ groups for training.
- 3 Average the K results




Benefits:

Issues: Splitting the data once into two parts (one for training and one for testing)
is not the most efficient way to use the data!

K-fold cross-validation:

- 1 Randomly partition the data into K groups
- 2 Train K times. Each time leave out exactly one of the K groups for testing and use the remaining $K - 1$ groups for training.
- 3 Average the K results




Benefits:

- We have used all data for training, and all data for testing, and used each data point the same number of times.

Issues: Splitting the data once into two parts (one for training and one for testing)
is not the most efficient way to use the data!

K-fold cross-validation:

- 1 Randomly partition the data into K groups
- 2 Train K times. Each time leave out exactly one of the K groups for testing and use the remaining $K - 1$ groups for training.
- 3 Average the K results



Benefits:

- We have used all data for training, and all data for testing, and used each data point the same number of times.
- Cross-validation returns an unbiased estimate of the generalization error and its variance.

Table of Contents

1 Review of Last Week

- Data Model and Learning Algorithm
- Generalization Gap and Model Selection
- Bias-Variance Decomposition

2 Logistic Regression

- Classification
- Logistic Regression

Bias-Variance Decomposition

$$\begin{aligned}\mathbb{E}_{S_{\text{train}} \in \mathcal{D}} [L(f_{S_{\text{train}}})] &= \text{Var}_{\epsilon \sim \mathcal{D}_{\epsilon}} [\epsilon] && \text{(noise variance)} \\ &+ (f(\mathbf{x}_0) - \mathbb{E}_{S_{\text{train}}} [f_{S_{\text{train}}}(\mathbf{x}_0)])^2 && \text{(Bias)} \\ &+ \mathbb{E}_{S_{\text{train}} \sim \mathcal{D}} \left[(\mathbb{E}_{S_{\text{train}}} [f_{S_{\text{train}}}(\mathbf{x}_0)] - f_{S_{\text{train}}}(\mathbf{x}_0))^2 \right], && \text{(Variance)}\end{aligned}$$

which always lower-bounds the true error.


Bias-Variance Decomposition

$$\begin{aligned}\mathbb{E}_{S_{\text{train}} \in \mathcal{D}} [L(f_{S_{\text{train}}})] &= \text{Var}_{\epsilon \sim \mathcal{D}_{\epsilon}} [\epsilon] && \text{(noise variance)} \\ &+ (f(\mathbf{x}_0) - \mathbb{E}_{S_{\text{train}}} [f_{S_{\text{train}}}(\mathbf{x}_0)])^2 && \text{(Bias)} \\ &+ \mathbb{E}_{S_{\text{train}} \sim \mathcal{D}} \left[(\mathbb{E}_{S_{\text{train}}} [f_{S_{\text{train}}}(\mathbf{x}_0)] - f_{S_{\text{train}}}(\mathbf{x}_0))^2 \right], && \text{(Variance)}\end{aligned}$$

which always lower-bounds the true error.

\Rightarrow To minimize the true error, we need to select a method that **simultaneously achieves low bias and low variance**.

Double descent curve in Deep Learning



Last lecture:

- Generalization Gap and Model Selection
- Bias-Variance Decomposition

Last lecture:

- Generalization Gap and Model Selection
- Bias-Variance Decomposition

This lecture:

- Logistic Regression
- Exponential Families and Generalized Linear Models

Table of Contents

① Review of Last Week

② Logistic Regression

- Classification
- Logistic Regression

Table of Contents

- 1 Review of Last Week
 - Data Model and Learning Algorithm
 - Generalization Gap and Model Selection
 - Bias-Variance Decomposition
- 2 Logistic Regression
 - Classification
 - Logistic Regression

Definition of classification

We observe some data $\mathcal{S} = \{\mathbf{x}_n, y_n\}_{n=1}^N \in \mathcal{X} \times \underbrace{\mathcal{Y}}_{\text{Discrete set}}$

Definition of classification

We observe some data $\mathcal{S} = \{\mathbf{x}_n, y_n\}_{n=1}^N \in \mathcal{X} \times \underbrace{\mathcal{Y}}_{\text{Discrete set}}$

Goal: given a new observation \mathbf{x} , we want to predict its label y .

Definition of classification

We observe some data $\mathcal{S} = \{\mathbf{x}_n, y_n\}_{n=1}^N \in \mathcal{X} \times \underbrace{\mathcal{Y}}_{\text{Discrete set}}$

Goal: given a new observation \mathbf{x} , we want to predict its label y .

How: relates input to a categorical variable

Classifier


A classifier $f : \mathcal{X} \rightarrow \mathcal{Y}$

Classifier

A **classifier** $f : \mathcal{X} \rightarrow \mathcal{Y}$ divides the input space into a collection of regions belonging to each class.

Classifier


A **classifier** $f : \mathcal{X} \rightarrow \mathcal{Y}$ divides the input space into a collection of regions belonging to each class.




Linear Decision boundary

Classifier

A **classifier** $f: \mathcal{X} \rightarrow \mathcal{Y}$ divides the input space into a collection of regions belonging to each class.




Linear Decision boundary




Nonlinear Decision boundary

Classifier

A **classifier** $f: \mathcal{X} \rightarrow \mathcal{Y}$ divides the input space into a collection of regions belonging to each class.



Linear Decision boundary



Nonlinear Decision boundary

Classification: a special case of regression?

Classification is a **regression problem** with discrete labels:

$$(\mathbf{x}, y) \in \mathcal{X} \times \{0, 1\} \subset \mathcal{X} \times \mathbb{R} \quad (15)$$


Classification: a special case of regression?

Classification is a **regression problem** with discrete labels:

$$(\mathbf{x}, y) \in \mathcal{X} \times \{0, 1\} \subset \mathcal{X} \times \mathbb{R} \quad (15)$$


Could we use previously seen regression methods to solve it?

Is it a good idea to use some regression methods?




Classification is not just a special form of regression

- The predicted values are not probabilities (not in $[0, 1]$)




Classification is not just a special form of regression

- Sensitivity to unbalanced data



Classification is not just a special form of regression



- Sensitivity to unbalanced data



The position of the line depends crucially on how many points are in each class.



Classification is not just a special form of regression

- Sensitivity to extreme values:



Classification is not just a special form of regression

- Sensitivity to extreme values:



The position of the line depends crucially on where the points lie.

How to perform classification?

- Many approaches have been developed

How to perform classification?

- Many approaches have been developed
- We won't cover them in detail today

How to perform classification?

- Many approaches have been developed
- We won't cover them in detail today
- Instead, we will provide quick introductions


How to perform classification?

- Many approaches have been developed
- We won't cover them in detail today
- Instead, we will provide quick introductions
- Fundamental task of classification:

Divide the space into distinct decision regions


k-Nearest Neighbor

Assume that nearby points are likely to have similar labels



k-Nearest Neighbor

Assume that nearby points are likely to have similar labels



A new point x is classified based on the **majority vote of its k -nearest neighbors**.

k-Nearest Neighbor

Pros:

Cons:




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training

Cons:

1-Nearest Neighbor Classifier




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training
- **Easy** to implement

Cons:

1-Nearest Neighbor Classifier




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training
- **Easy** to implement
- Works well in **low dimensions**, allowing for very complex decision boundaries

Cons:

1-Nearest Neighbor Classifier




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training
- **Easy** to implement
- Works well in **low dimensions**, allowing for very complex decision boundaries

Cons:

- **Slow** at query time

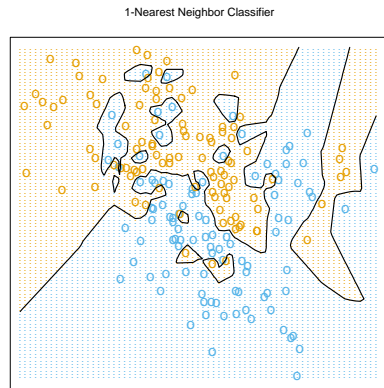


FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training
- **Easy** to implement
- Works well in **low dimensions**, allowing for very complex decision boundaries

Cons:

- **Slow** at query time
- Not suitable for high-dimensional data




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

k-Nearest Neighbor

Pros:

- No **optimization** or training
- **Easy** to implement
- Works well in **low dimensions**, allowing for very complex decision boundaries

Cons:

- **Slow** at query time
- Not suitable for high-dimensional data
- Choosing the right local distance is crucial

1-Nearest Neighbor Classifier




FIGURE 2.3. The same classification example in two dimensions as in Figure 2.1. The classes are coded as a binary variable (BLUE = 0, ORANGE = 1), and then predicted by 1-nearest-neighbor classification.

Linear Decision boundaries

Assume we restrict ourselves to linear decision boundaries (hyperplane):

Linear Decision boundaries


Assume we restrict ourselves to linear decision boundaries (hyperplane):

$$\Rightarrow \text{Prediction: } f(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$$

Linear Decision boundaries


Assume we restrict ourselves to linear decision boundaries (hyperplane):

$$\Rightarrow \text{Prediction: } f(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$$




Separating hyperplane

Assume we restrict ourselves to linear decision boundaries (hyperplane):



Separating hyperplane


Assume we restrict ourselves to linear decision boundaries (hyperplane):



Assume the data are linearly separable, i.e., a separating hyperplane exists

Separating hyperplane

Assume we restrict ourselves to linear decision boundaries (hyperplane):




Assume the data are linearly separable, i.e., a separating hyperplane exists

Which separating hyperplane would you pick?


Margin

Key concept: The margin is the distance from the hyperplane to the closest point.



Margin


Key concept: The margin is the distance from the hyperplane to the closest point.



Take the one with the largest margin!


Max-margin separating hyperplane

Choose the hyperplane which maximizes the margin



Max-margin separating hyperplane


Choose the hyperplane which maximizes the margin



WHY?

Max-margin separating hyperplane


Choose the hyperplane which maximizes the margin



WHY? If we slightly change the training set, the number of misclassifications will stay low

Max-margin separating hyperplane


Choose the hyperplane which maximizes the margin



WHY? If we slightly change the training set, the number of misclassifications will stay low


⇒ It will lead us to support vector machine (SVM) and logistic regression!

Nonlinear classifier




Nonlinear classifier

- Linear decision boundaries will not always work




Nonlinear classifier

- Linear decision boundaries will not always work
- Feature augmentation (x, x^2, x^3, x^4)



Nonlinear classifier

- Linear decision boundaries will not always work
- Feature augmentation (x, x^2, x^3, x^4)
- Kernel method



Formalizing binary classification

Formalizing binary classification

- Setting: $(X, Y) \sim \mathcal{D}$ with ranges $\mathcal{X}, \mathcal{Y} = \{-1, 1\}$

Formalizing binary classification

- Setting: $(X, Y) \sim \mathcal{D}$ with ranges $\mathcal{X}, \mathcal{Y} = \{-1, 1\}$
- Loss function: (0 – 1 loss) $\ell(y, y') = 1_{y \neq y'}$

Formalizing binary classification

- Setting: $(X, Y) \sim \mathcal{D}$ with ranges $\mathcal{X}, \mathcal{Y} = \{-1, 1\}$
- Loss function: (0 – 1 loss) $\ell(y, y') = 1_{y \neq y'}$
- True risk for the classification:

$$L_{\mathcal{D}}(f) = \underbrace{\mathbb{E}_{\mathcal{D}} [1_{Y \neq f(X)}]}_{\text{classification error}} = \underbrace{\Pr_{\mathcal{D}} [Y \neq f(X)]}_{\text{probability of making an error}}$$

Formalizing binary classification

- Setting: $(X, Y) \sim \mathcal{D}$ with ranges $\mathcal{X}, \mathcal{Y} = \{-1, 1\}$
- Loss function: (0 – 1 loss) $\ell(y, y') = 1_{y \neq y'}$
- True risk for the classification:

$$L_{\mathcal{D}}(f) = \underbrace{\mathbb{E}_{\mathcal{D}} [1_{Y \neq f(X)}]}_{\text{classification error}} = \underbrace{\Pr_{\mathcal{D}} [Y \neq f(X)]}_{\text{probability of making an error}}$$

- Goal: minimize $L_{\mathcal{D}}(f)$

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.
- **Maximum A-Posteriori** (MAP):

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.
- **Maximum A-Posteriori** (MAP):
If we want to maximize the probability of guessing the correct label,

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

$$\hat{y}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x}) \quad (16)$$

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , *the probability that the “correct” label is y* is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

$$\hat{y}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x}) \quad (16)$$

This classifier is also called the *Bayes classifier*: $f^* = \arg \min_f L_{\mathcal{D}}(f)$, where

$$f^*(\mathbf{x}) \in \arg \max_{\{-1,1\}} \Pr(Y = y|X = \mathbf{x}) \quad (17)$$

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , the probability that the “correct” label is y is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

$$\hat{y}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x}) \quad (16)$$

This classifier is also called the *Bayes classifier*: $f^* = \arg \min_f L_{\mathcal{D}}(f)$, where

$$f^*(\mathbf{x}) \in \arg \max_{\{-1,1\}} \Pr(Y = y|X = \mathbf{x}) \quad (17)$$

- In practice, we do not know the joint distribution $p(\mathbf{x}, y)$

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , the probability that the “correct” label is y is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

$$\hat{y}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x}) \quad (16)$$

This classifier is also called the *Bayes classifier*: $f^* = \arg \min_f L_{\mathcal{D}}(f)$, where

$$f^*(\mathbf{x}) \in \arg \max_{\{-1,1\}} \Pr(Y = y|X = \mathbf{x}) \quad (17)$$

- In practice, we do not know the joint distribution $p(\mathbf{x}, y)$
 \Rightarrow Bayes classifier is an unattainable gold standard.

Optimal classification for known generating model

What is the **optimal performance**, regardless of the finiteness of the training data?

Assume that we know the joint distribution $p(\mathbf{x}, y)$.

- For a given input \mathbf{x} , the probability that the “correct” label is y is $p(y|\mathbf{x})$.

- **Maximum A-Posteriori** (MAP):

If we want to maximize the probability of guessing the correct label, then we should choose the decision rule

$$\hat{y}(\mathbf{x}) = \arg \max_{y \in \mathcal{Y}} p(y|\mathbf{x}) \quad (16)$$

This classifier is also called the *Bayes classifier*: $f^* = \arg \min_f L_{\mathcal{D}}(f)$, where

$$f^*(\mathbf{x}) \in \arg \max_{\{-1,1\}} \Pr(Y = y|X = \mathbf{x}) \quad (17)$$

- In practice, we do not know the joint distribution $p(\mathbf{x}, y)$

⇒ Bayes classifier is an unattainable gold standard.

But we can use the data to learn the distribution (by assuming the data distribution).

Proof of the Bayes classifier

Claim 1: $\forall \mathbf{x} \in \mathcal{X}, f(\mathbf{x}^*) \in \arg \min_{y \in \mathcal{Y}} \Pr(Y \neq y | X = \mathbf{x}) \implies f^* \in \arg \min_{f: \mathcal{X} \rightarrow \mathcal{Y}} L_{\mathcal{D}}(f)$

$$L_{\mathcal{D}}(f) = \mathbb{E}_{X,Y} [1_{Y \neq f(X)}] = \mathbb{E}_X [\mathbb{E}_{Y|X} [1_{Y \neq f(X)} | X]] \quad (18)$$

$$= \mathbb{E}_X [\Pr(Y \neq f(X) | X)] \quad (19)$$

$$\geq \mathbb{E}_X \left[\min_{y \in \mathcal{Y}} \Pr(Y \neq y | X) \right] \quad (20)$$

$$= \mathbb{E}_X [\Pr(Y \neq f^*(X) | X)] = \mathbb{E}_{X,Y} [1_{Y \neq f^*(X)}] = L_{\mathcal{D}}(f^*) \quad (21)$$

Claim 2: $f^*(\mathbf{x}) \in \arg \min_{y \in \mathcal{Y}} \Pr(Y \neq y | X = \mathbf{x})$

$$f^*(\mathbf{x}) \in \arg \max_{y \in \mathcal{Y}} \Pr(Y = y | X = \mathbf{x}) = \arg \min_{y \in \mathcal{Y}} \Pr(Y \neq y | X = \mathbf{x}) \quad (22)$$

Two classes of classification algorithms

- **Non-parametric:**
- **Parametric:**

Two classes of classification algorithms

- **Non-parametric:** approximate the conditional distribution $\Pr(Y = y|X = \mathbf{x})$ via local averaging
- **Parametric:**

Two classes of classification algorithms

- **Non-parametric:** approximate the conditional distribution $\Pr(Y = y|X = \mathbf{x})$ via local averaging
⇒ Follow nearest neighbors' decisions (KNN).
- **Parametric:**

Two classes of classification algorithms

- **Non-parametric:** approximate the conditional distribution $\Pr(Y = y|X = \mathbf{x})$ via local averaging
⇒ Follow nearest neighbors' decisions (KNN).
- **Parametric:** approximate true distribution \mathcal{D} via training data

Two classes of classification algorithms

- **Non-parametric:** approximate the conditional distribution $\Pr(Y = y|X = \mathbf{x})$ via local averaging
⇒ Follow nearest neighbors' decisions (KNN).
- **Parametric:** approximate true distribution \mathcal{D} via training data
⇒ Minimize the empirical risk on training data (ERM)

Recap

- Classification:
 - Mapping inputs to discrete outputs (categorical)
 - Not a special form of regression!
- Ways to perform classification:
 - **Non-parametric**: K-Nearest-Neighbors
 - **Parametric**: learning X -to- Y mapping via ERM
 - More complex decision boundaries? Non-linear classifiers

Table of Contents

- ① Review of Last Week
 - Data Model and Learning Algorithm
 - Generalization Gap and Model Selection
 - Bias-Variance Decomposition
- ② Logistic Regression
 - Classification
 - Logistic Regression

Motivation for Logistic Regression

Rather than modeling the output Y directly,

Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Previously, we used a linear regression model $\Pr(Y = 1|X = x) = \mathbf{x}^\top \mathbf{w} + w_0$, but

Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Previously, we used a linear regression model $\Pr(Y = 1|X = x) = \mathbf{x}^\top \mathbf{w} + w_0$, but

- The predicted value is not in $[0, 1]$.

Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Previously, we used a linear regression model $\Pr(Y = 1|X = x) = \mathbf{x}^\top \mathbf{w} + w_0$, but



- The predicted value is not in $[0, 1]$.
- Very large ($y \gg 1$) or very small ($y \ll 0$) values of the prediction will contribute to the error if we use the squared loss.

Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Previously, we used a linear regression model $\Pr(Y = 1|X = x) = \mathbf{x}^\top \mathbf{w} + w_0$, but

- The predicted value is not in $[0, 1]$.
- Very large ($y \gg 1$) or very small ($y \ll 0$) values of the prediction will contribute to the error if we use the squared loss.





Motivation for Logistic Regression

Rather than modeling the output Y directly,
we can **model the probability** that Y belongs to a particular class.

Previously, we used a linear regression model $\Pr(Y = 1|X = x) = \mathbf{x}^\top \mathbf{w} + w_0$, but


- The predicted value is not in $[0, 1]$.
- Very large ($y \gg 1$) or very small ($y \ll 0$) values of the prediction will contribute to the error if we use the squared loss.



Solution: Transforming the predictions that take values in $(-\infty, \infty)$ into $[0, 1]$.

The logistic function

Consider first of all the case of two classes.




The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

(24)




The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} \quad (24)$$




The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$



The logistic function

Consider first of all the case of two classes.


The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (25)$$



The logistic function

Consider first of all the case of two classes.


The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$



The logistic function

Consider first of all the case of two classes.


The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$



Properties of the logistic function (Exercise):

The logistic function

Consider first of all the case of two classes.


The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$



Properties of the logistic function (Exercise):

- $1 - \sigma(\eta) = \sigma(-\eta)$

The logistic function

Consider first of all the case of two classes.


The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$



Properties of the logistic function (Exercise):

- $1 - \sigma(\eta) = \sigma(-\eta)$
- $\sigma'(\eta) = \sigma(\eta) (1 - \sigma(\eta))$

The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :


$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$

$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$

(Exercise) For the case of $K > 2$ classes:



Properties of the logistic function (Exercise):

- $1 - \sigma(\eta) = \sigma(-\eta)$
- $\sigma'(\eta) = \sigma(\eta) (1 - \sigma(\eta))$

The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$


$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$

(Exercise) For the case of $K > 2$ classes:

$$p(\mathcal{C}_k|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{\sum_j p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j)} \quad (26)$$



Properties of the logistic function (Exercise):

- $1 - \sigma(\eta) = \sigma(-\eta)$
- $\sigma'(\eta) = \sigma(\eta) (1 - \sigma(\eta))$

The logistic function

Consider first of all the case of two classes.

The posterior probability for class \mathcal{C}_1 :

$$p(\mathcal{C}_1|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1) + p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \quad (23)$$


$$= \frac{1}{1 + \exp(-\eta)} = \sigma(\eta) \quad (24)$$

where we have defined

$$\eta = \ln \frac{p(\mathbf{x}|\mathcal{C}_1)p(\mathcal{C}_1)}{p(\mathbf{x}|\mathcal{C}_2)p(\mathcal{C}_2)} \text{ and } \sigma(\eta) := \frac{e^\eta}{1 + e^\eta} \quad (25)$$

(Exercise) For the case of $K > 2$ classes:

$$p(\mathcal{C}_k|\mathbf{x}) = \frac{p(\mathbf{x}|\mathcal{C}_k)p(\mathcal{C}_k)}{\sum_j p(\mathbf{x}|\mathcal{C}_j)p(\mathcal{C}_j)} = \frac{\exp(\eta_k)}{\sum_j \exp(\eta_j)} \quad (26)$$



Properties of the logistic function (Exercise):

- $1 - \sigma(\eta) = \sigma(-\eta)$
- $\sigma'(\eta) = \sigma(\eta) (1 - \sigma(\eta))$

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] \quad , \quad (28)$$

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] \quad , \quad (28)$$

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

where we predict a real value (a probability) and not a label.

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

where we predict a real value (a probability) and not a label.

Label prediction: quantize the probability

$$\text{if } p(1|\mathbf{x}) \geq 1/2 \Rightarrow \text{predict the class 1} \quad (29)$$

$$\text{if } p(1|\mathbf{x}) < 1/2 \Rightarrow \text{predict the class 0} \quad (30)$$

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

where we predict a real value (a probability) and not a label.

Label prediction: quantize the probability

$$\text{if } p(1|\mathbf{x}) \geq 1/2 \Rightarrow \text{predict the class 1} \quad (29)$$

$$\text{if } p(1|\mathbf{x}) < 1/2 \Rightarrow \text{predict the class 0} \quad (30)$$

Interpretation:

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

where we predict a real value (a probability) and not a label.

Label prediction: quantize the probability

$$\text{if } p(1|\mathbf{x}) \geq 1/2 \Rightarrow \text{predict the class 1} \quad (29)$$

$$\text{if } p(1|\mathbf{x}) < 1/2 \Rightarrow \text{predict the class 0} \quad (30)$$

Interpretation:

- Very large $\mathbf{x}^\top \mathbf{w} + w_0$ corresponds to $p(1|\mathbf{x})$ very close to 0 or 1 (high confidence).

Logistic Regression

Given a “new” feature vector \mathbf{x} , we predict the (posterior) probability of the two class labels given \mathbf{x} by means of

$$p(1|\mathbf{x}) := \Pr[Y = 1|\mathbf{X} = \mathbf{x}] = \sigma(\mathbf{x}^\top \mathbf{w} + w_0) \quad (27)$$

$$p(0|\mathbf{x}) := \Pr[Y = 0|\mathbf{X} = \mathbf{x}] = 1 - \sigma(\mathbf{x}^\top \mathbf{w} + w_0) , \quad (28)$$

where we predict a real value (a probability) and not a label.

Label prediction: quantize the probability

$$\text{if } p(1|\mathbf{x}) \geq 1/2 \Rightarrow \text{predict the class 1} \quad (29)$$

$$\text{if } p(1|\mathbf{x}) < 1/2 \Rightarrow \text{predict the class 0} \quad (30)$$

Interpretation:

- Very large $\mathbf{x}^\top \mathbf{w} + w_0$ corresponds to $p(1|\mathbf{x})$ very close to 0 or 1 (high confidence).
- Small $|\mathbf{x}^\top \mathbf{w} + w_0|$ corresponds to $p(1|\mathbf{x})$ very close to 0.5 (low confidence).

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}).

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) \quad (31)$$

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) \quad (31)$$

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) = \prod_{n=1}^N p(\{\mathbf{x}_n, y_n\} | \mathbf{w}), \quad (31)$$

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) = \prod_{n=1}^N p(\{\mathbf{x}_n, y_n\} | \mathbf{w}) , \quad (31)$$

or equivalently,

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} [-\log \mathcal{L}(\mathbf{w})] = \arg \min_{\mathbf{w}} \sum_{n=1}^N -\log (p(\{\mathbf{x}_n, y_n\} | \mathbf{w})) . \quad (32)$$

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) = \prod_{n=1}^N p(\{\mathbf{x}_n, y_n\} | \mathbf{w}), \quad (31)$$

or equivalently,

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} [-\log \mathcal{L}(\mathbf{w})] = \arg \min_{\mathbf{w}} \sum_{n=1}^N -\log (p(\{\mathbf{x}_n, y_n\} | \mathbf{w})) . \quad (32)$$

This estimator is **consistent** (under mild condition):

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) = \prod_{n=1}^N p(\{\mathbf{x}_n, y_n\} | \mathbf{w}), \quad (31)$$

or equivalently,

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} [-\log \mathcal{L}(\mathbf{w})] = \arg \min_{\mathbf{w}} \sum_{n=1}^N -\log (p(\{\mathbf{x}_n, y_n\} | \mathbf{w})) . \quad (32)$$

This estimator is **consistent** (under mild condition):
 \implies if the data are generated according to the model,

MLE is a method of estimating the parameters of a statistical model

Assume a training set S_{train} , consisting of i.i.d. samples $\{\mathbf{x}_n, y_n\}_{n=1}^N$ (fixed but unknown \mathcal{D}). The MLE finds the parameters \mathbf{w}^* under which $\{\mathbf{x}_n, y_n\}$ are the most likely:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} \mathcal{L}(\mathbf{w}) := p(\{\mathbf{x}_n, y_n\}_{n=1}^N | \mathbf{w}) = \prod_{n=1}^N p(\{\mathbf{x}_n, y_n\} | \mathbf{w}), \quad (31)$$

or equivalently,

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} [-\log \mathcal{L}(\mathbf{w})] = \arg \min_{\mathbf{w}} \sum_{n=1}^N -\log (p(\{\mathbf{x}_n, y_n\} | \mathbf{w})) . \quad (32)$$

This estimator is **consistent** (under mild condition):

\implies if the data are generated according to the model,

the MLE converges to the true parameter when $N \rightarrow \infty$.

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

where \mathbf{X} does not depend on \mathbf{w} , i.e., $p(\mathbf{W})$ is a constant w.r.t. arbitrary \mathbf{w} .

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

where \mathbf{X} does not depend on \mathbf{w} , i.e., $p(\mathbf{W})$ is a constant w.r.t. arbitrary \mathbf{w} .

For Logistic Regression, we have:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n) \quad (35)$$

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

where \mathbf{X} does not depend on \mathbf{w} , i.e., $p(\mathbf{X})$ is a constant w.r.t. arbitrary \mathbf{w} .

For Logistic Regression, we have:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n) = \prod_{n:y_n=1} p(y_n = 1|\mathbf{x}_n) \prod_{n:y_n=0} p(y_n = 0|\mathbf{x}_n) \quad (34)$$

(35)

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

where \mathbf{X} does not depend on \mathbf{w} , i.e., $p(\mathbf{X})$ is a constant w.r.t. arbitrary \mathbf{w} .

For Logistic Regression, we have:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n) = \prod_{n:y_n=1} p(y_n = 1|\mathbf{x}_n) \prod_{n:y_n=0} p(y_n = 0|\mathbf{x}_n) \quad (34)$$

$$= \prod_{n=1}^N \sigma(\mathbf{x}_n^\top \mathbf{w})^{y_n} [1 - \sigma(\mathbf{x}_n^\top \mathbf{w})]^{1-y_n} \quad (35)$$

MLE for Logistic Regression

The likelihood of the data $\{\mathbf{y}, \mathbf{X}\}$ given the parameter \mathbf{w} , i.e., $p(\mathbf{y}, \mathbf{X}|\mathbf{w})$.

$$p(\mathbf{y}, \mathbf{X}|\mathbf{w}) = p(\mathbf{X}|\mathbf{w})p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = p(\mathbf{X})p(\mathbf{y}|\mathbf{X}, \mathbf{w}), \quad (33)$$

where \mathbf{X} does not depend on \mathbf{w} , i.e., $p(\mathbf{X})$ is a constant w.r.t. arbitrary \mathbf{w} .

For Logistic Regression, we have:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \prod_{n=1}^N p(y_n|\mathbf{x}_n) = \prod_{n:y_n=1} p(y_n=1|\mathbf{x}_n) \prod_{n:y_n=0} p(y_n=0|\mathbf{x}_n) \quad (34)$$

$$= \prod_{n=1}^N \sigma(\mathbf{x}_n^\top \mathbf{w})^{y_n} [1 - \sigma(\mathbf{x}_n^\top \mathbf{w})]^{1-y_n} \quad (35)$$

As a result,

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \arg \min_{\mathbf{w}} \left(-\log p(\mathbf{y}|\mathbf{X}, \mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (36)$$

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Let's minimize $\mathcal{L}(\mathbf{w})$ through the property of stationary points.

$$\nabla \mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \quad (38)$$

(39)

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Let's minimize $\mathcal{L}(\mathbf{w})$ through the property of stationary points.

$$\nabla \mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \left(\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n \right) \quad (38)$$

$$= \frac{1}{N} \mathbf{X}^\top [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}] , \quad (39)$$

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Let's minimize $\mathcal{L}(\mathbf{w})$ through the property of stationary points.

$$\nabla \mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \left(\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n \right) \quad (38)$$

$$= \frac{1}{N} \mathbf{X}^\top [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}] , \quad (39)$$

where $\mathbf{X} \in \mathbb{R}^{N \times d}$.

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Let's minimize $\mathcal{L}(\mathbf{w})$ through the property of stationary points.

$$\nabla \mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \quad (38)$$

$$= \frac{1}{N} \mathbf{X}^\top [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}] , \quad (39)$$

where $\mathbf{X} \in \mathbb{R}^{N \times d}$.

\Rightarrow It has no closed-form solution to $\nabla \mathcal{L}(\mathbf{w}) = 0$.

Gradient of the negative log likelihood

Recall that

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \right) \quad (37)$$

Let's minimize $\mathcal{L}(\mathbf{w})$ through the property of stationary points.

$$\nabla \mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \quad (38)$$

$$= \frac{1}{N} \mathbf{X}^\top [\sigma(\mathbf{X}\mathbf{w}) - \mathbf{y}] , \quad (39)$$

where $\mathbf{X} \in \mathbb{R}^{N \times d}$.

⇒ It has no closed-form solution to $\nabla \mathcal{L}(\mathbf{w}) = 0$.

⇒ Good news: the cost function L is convex.

Convexity of the loss function L

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

Convexity of the loss function L

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

- Proof: \mathcal{L} is obtained through simple convexity preserving operations:

Convexity of the loss function L

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

- Proof: \mathcal{L} is obtained through simple convexity preserving operations:
 - 1 Positive combinations of convex functions is convex

Convexity of the loss function L

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

- Proof: \mathcal{L} is obtained through simple convexity preserving operations:
 - 1 Positive combinations of convex functions is convex
 - 2 Composition of a convex and a linear functions is convex

Convexity of the loss function \mathcal{L}

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

- Proof: \mathcal{L} is obtained through simple convexity preserving operations:
 - 1 Positive combinations of convex functions is convex
 - 2 Composition of a convex and a linear functions is convex
 - 3 A linear function is both convex and concave

Convexity of the loss function \mathcal{L}

- Claim: The function

$$\mathcal{L}(\mathbf{w}) := \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \quad (40)$$

is convex in the weight vector \mathbf{w}

- Proof: \mathcal{L} is obtained through simple convexity preserving operations:
 - 1 Positive combinations of convex functions is convex
 - 2 Composition of a convex and a linear functions is convex
 - 3 A linear function is both convex and concave
 - 4 $\eta \rightarrow \log(1 + e^\eta)$ is convex (exercise)

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

$$\implies \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

$$\implies \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

- With the property of “A linear function is both convex and concave”

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

$$\implies \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

- With the property of “A linear function is both convex and concave”

$$\implies -y_n \mathbf{x}_n^\top \mathbf{w} \text{ is convex}$$

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

$$\implies \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

- With the property of “A linear function is both convex and concave”

$$\implies -y_n \mathbf{x}_n^\top \mathbf{w} \text{ is convex}$$

- Positive combinations of convex functions is convex

Proof of the convexity of \mathcal{L}

- Proof of “ $\eta \rightarrow \log(1 + e^\eta)$ is convex”

$$h(\eta) = \log(1 + e^\eta) \quad (41)$$

$$h'(\eta) = \frac{e^\eta}{1 + e^\eta} = \sigma(\eta) \quad (42)$$

$$h''(\eta) = \sigma'(\eta) = \frac{e^\eta}{(1 + e^\eta)^2} \geq 0 \quad (43)$$

- With the property of “composition of a convex and a linear functions is convex”:

$$\implies \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

- With the property of “A linear function is both convex and concave”

$$\implies -y_n \mathbf{x}_n^\top \mathbf{w} \text{ is convex}$$

- Positive combinations of convex functions is convex

$$\mathcal{L}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N -y_n \mathbf{x}_n^\top \mathbf{w} + \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) \text{ is convex}$$

Second proof: Hessian of \mathcal{L} is psd

The Hessian $\nabla^2 \mathcal{L}$ is the **matrix** whose entries are the **second derivatives** $\frac{\partial^2}{\partial w_i \partial w_j} \mathcal{L}(\mathbf{w})$.

Second proof: Hessian of \mathcal{L} is psd

The Hessian $\nabla^2 \mathcal{L}$ is the **matrix** whose entries are the **second derivatives** $\frac{\partial^2}{\partial w_i \partial w_j} \mathcal{L}(\mathbf{w})$.

$$\begin{aligned}\nabla^2 \mathcal{L}(\mathbf{w}) &= \nabla [\nabla \mathcal{L}(\mathbf{w})]^\top \\ &= \nabla \left[\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \right]^\top \\ &= \frac{1}{N} \sum_{n=1}^N \nabla \sigma(\mathbf{x}_n^\top \mathbf{w}) \mathbf{x}_n^\top = \frac{1}{N} \sum_{n=1}^N \sigma(\mathbf{x}_n^\top \mathbf{w}) (1 - \sigma(\mathbf{x}_n^\top \mathbf{w})) \mathbf{x}_n \mathbf{x}_n^\top\end{aligned}$$

Second proof: Hessian of \mathcal{L} is psd

The Hessian $\nabla^2 \mathcal{L}$ is the **matrix** whose entries are the **second derivatives** $\frac{\partial^2}{\partial w_i \partial w_j} \mathcal{L}(\mathbf{w})$.

$$\begin{aligned}\nabla^2 \mathcal{L}(\mathbf{w}) &= \nabla [\nabla \mathcal{L}(\mathbf{w})]^\top \\ &= \nabla \left[\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \right]^\top \\ &= \frac{1}{N} \sum_{n=1}^N \nabla \sigma(\mathbf{x}_n^\top \mathbf{w}) \mathbf{x}_n^\top = \frac{1}{N} \sum_{n=1}^N \sigma(\mathbf{x}_n^\top \mathbf{w}) (1 - \sigma(\mathbf{x}_n^\top \mathbf{w})) \mathbf{x}_n \mathbf{x}_n^\top\end{aligned}$$

It can be written under the matrix form:

$$\nabla^2 \mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{X}^\top \mathbf{S} \mathbf{X}, \quad \text{where } \mathbf{S} = \text{diag} [\sigma(\mathbf{x}_n^\top \mathbf{w}) (1 - \sigma(\mathbf{x}_n^\top \mathbf{w}))] \succeq 0 \quad (44)$$

Second proof: Hessian of \mathcal{L} is psd

The Hessian $\nabla^2 \mathcal{L}$ is the **matrix** whose entries are the **second derivatives** $\frac{\partial^2}{\partial w_i \partial w_j} \mathcal{L}(\mathbf{w})$.

$$\begin{aligned}\nabla^2 \mathcal{L}(\mathbf{w}) &= \nabla [\nabla \mathcal{L}(\mathbf{w})]^\top \\ &= \nabla \left[\frac{1}{N} \sum_{n=1}^N \mathbf{x}_n (\sigma(\mathbf{x}_n^\top \mathbf{w}) - y_n) \right]^\top \\ &= \frac{1}{N} \sum_{n=1}^N \nabla \sigma(\mathbf{x}_n^\top \mathbf{w}) \mathbf{x}_n^\top = \frac{1}{N} \sum_{n=1}^N \sigma(\mathbf{x}_n^\top \mathbf{w}) (1 - \sigma(\mathbf{x}_n^\top \mathbf{w})) \mathbf{x}_n \mathbf{x}_n^\top\end{aligned}$$

It can be written under the matrix form:

$$\nabla^2 \mathcal{L}(\mathbf{w}) = \frac{1}{N} \mathbf{X}^\top \mathbf{S} \mathbf{X}, \quad \text{where } \mathbf{S} = \text{diag} [\sigma(\mathbf{x}_n^\top \mathbf{w}) (1 - \sigma(\mathbf{x}_n^\top \mathbf{w}))] \succeq 0 \quad (44)$$

$\implies \mathcal{L}$ is convex since $\nabla^2 \mathcal{L}(\mathbf{w}) \succeq 0$.