

Lecture 8: Introduction to Unsupervised Learning, K-means, Gaussian Mixture Model (GMM), and EM algorithm

Tao LIN

SoE, Westlake University

November 4, 2025



1 EM Algorithm

- Expectation-Maximization Algorithm
- Supp.

Table of Contents

1 EM Algorithm

- Expectation-Maximization Algorithm
- Supp.

Table of Contents

1 EM Algorithm

- Expectation-Maximization Algorithm
- Supp.

Motivation of Expectation-Maximization (EM)

Recall the equation of maximum likelihood for GMM:

$$\max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}) := \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (1)$$

Motivation of Expectation-Maximization (EM)

Recall the equation of maximum likelihood for GMM:

$$\max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}) := \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (1)$$

Remarks:

Motivation of Expectation-Maximization (EM)

Recall the equation of maximum likelihood for GMM:

$$\max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}) := \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (1)$$

Remarks:

- Computing maximum likelihood for GMM is difficult, due to .

Motivation of Expectation-Maximization (EM)

Recall the equation of maximum likelihood for GMM:

$$\max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}) := \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (1)$$

Remarks:

- Computing maximum likelihood for GMM is difficult, due to the \log outside the sum.

EM provides an elegant solution by iteratively building & optimizing lower bounds.

Motivation of Expectation-Maximization (EM)

Recall the equation of maximum likelihood for GMM:

$$\max_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}) := \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (1)$$

Remarks:

- Computing maximum likelihood for GMM is difficult, due to the \log outside the sum.

EM provides an elegant solution by iteratively building & optimizing lower bounds.

- It uses an iterative two-step procedure, where individual steps usually involve problems that are easy to optimize.

An Overview of EM

- The EM algorithm provides an elegant approach to latent variable estimation problems.
- EM iteratively tries to overcome the challenge of not knowing the hidden variables z_n .

An Overview of EM

- The EM algorithm provides an elegant approach to latent variable estimation problems.
- EM iteratively tries to overcome the challenge of not knowing the hidden variables z_n .
- If we knew the values of the latent variables z_n , maximum likelihood estimation would be simple.

An Overview of EM

- The EM algorithm provides an elegant approach to latent variable estimation problems.
- EM iteratively tries to overcome the challenge of not knowing the hidden variables z_n .
- If we knew the values of the latent variables z_n , maximum likelihood estimation would be simple.
- Since we don't know them, we:

An Overview of EM

- The EM algorithm provides an elegant approach to latent variable estimation problems.
- EM iteratively tries to overcome the challenge of not knowing the hidden variables z_n .
- If we knew the values of the latent variables z_n , maximum likelihood estimation would be simple.
- Since we don't know them, we:
 - **E-step:** Compute the *distribution* of the latent variables given our current parameters.

An Overview of EM

- The EM algorithm provides an elegant approach to latent variable estimation problems.
- EM iteratively tries to overcome the challenge of not knowing the hidden variables z_n .
- If we knew the values of the latent variables z_n , maximum likelihood estimation would be simple.
- Since we don't know them, we:
 - **E-step:** Compute the *distribution* of the latent variables given our current parameters.
 - **M-step:** Update the parameters using these *expected* latent variable values.

A Note on Notation: z_n (scalar) vs. \mathbf{z}_n (vector)

The latent variable for x_n can be written in two ways:

- \mathbf{z}_n is the **one-hot vector** (formal notation).
 - A K -dimensional vector, e.g., $\mathbf{z}_n = [0, 1, 0, \dots, 0]^\top$.
 - Useful for writing the complete-data likelihood compactly:

$$p(\mathbf{x}_n, \mathbf{z}_n | \boldsymbol{\theta}) = \prod_{k=1}^K [\pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)]^{z_{nk}}$$

- z_n is the **scalar index** (convenient notation).
 - A single number, e.g., $z_n = 2$.
 - $z_n = k$ is equivalent to the vector \mathbf{z}_n having $z_{nk} = 1$.
- $q_n(z_n)$ is our **approximating posterior distribution**.
 - Since z_n takes one of K discrete values, $q_n(z_n)$ is a *categorical distribution*.
 - In the E-step, we set it to the true posterior: $q_n(z_n) = p(z_n | \mathbf{x}_n, \boldsymbol{\theta}^{(t)})$.
- $q_{kn}^{(t)}$ is the **probability** of component k for point n . This single value connects all notations:

$$q_{kn}^{(t)} := q_n(z_n = k) = p(z_n = k | \mathbf{x}_n, \boldsymbol{\theta}^{(t)}) = \mathbb{E}_{q_n}[\mathbf{z}_{nk}]$$

EM algorithm: Formal Summary

Start with $\theta^{(1)}$ and iterate:

EM algorithm: Formal Summary

Start with $\theta^{(1)}$ and iterate:

- ① **Expectation step:** Compute a lower bound $\underline{\mathcal{L}}$ to the true log-likelihood \mathcal{L} such that it is *tight* at the current parameters $\theta^{(t)}$:

$$\mathcal{L}(\theta) \geq \underline{\mathcal{L}}(\theta, \theta^{(t)}) \text{ and} \quad (2)$$

$$\mathcal{L}(\theta^{(t)}) = \underline{\mathcal{L}}(\theta^{(t)}, \theta^{(t)}). \quad (3)$$

EM algorithm: Formal Summary

Start with $\theta^{(1)}$ and iterate:

- ① **Expectation step:** Compute a lower bound $\underline{\mathcal{L}}$ to the true log-likelihood \mathcal{L} such that it is *tight* at the current parameters $\theta^{(t)}$:

$$\mathcal{L}(\theta) \geq \underline{\mathcal{L}}(\theta, \theta^{(t)}) \text{ and} \quad (2)$$

$$\mathcal{L}(\theta^{(t)}) = \underline{\mathcal{L}}(\theta^{(t)}, \theta^{(t)}). \quad (3)$$

- ② **Maximization step:** Find the new parameters $\theta^{(t+1)}$ by maximizing this lower bound:

$$\theta^{(t+1)} = \arg \max_{\theta} \underline{\mathcal{L}}(\theta, \theta^{(t)}). \quad (4)$$

EM algorithm: Formal Summary

Start with $\theta^{(1)}$ and iterate:

- ① **Expectation step:** Compute a lower bound $\underline{\mathcal{L}}$ to the true log-likelihood \mathcal{L} such that it is *tight* at the current parameters $\theta^{(t)}$:

$$\mathcal{L}(\theta) \geq \underline{\mathcal{L}}(\theta, \theta^{(t)}) \text{ and} \quad (2)$$

$$\mathcal{L}(\theta^{(t)}) = \underline{\mathcal{L}}(\theta^{(t)}, \theta^{(t)}). \quad (3)$$

- ② **Maximization step:** Find the new parameters $\theta^{(t+1)}$ by maximizing this lower bound:

$$\theta^{(t+1)} = \arg \max_{\theta} \underline{\mathcal{L}}(\theta, \theta^{(t)}). \quad (4)$$

Each iteration of EM is guaranteed to increase (or keep the same) the true log-likelihood $\mathcal{L}(\theta)$.

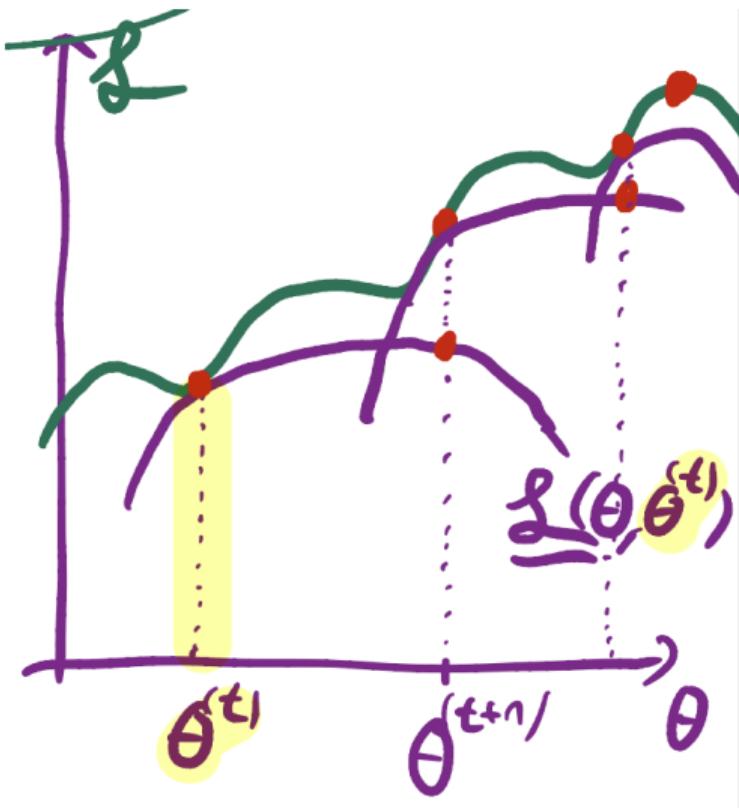


Figure: Visualization of EM iterations: Each M-step maximizes a new lower bound $\underline{\mathcal{L}}$ (purple), guaranteeing an increase in the true log-likelihood \mathcal{L} (green).

Background: What is KL Divergence?

The Kullback-Leibler (KL) Divergence is a measure of how one probability distribution $q(z)$ is different from a second, reference probability distribution $p(z)$.

For discrete distributions:

$$\text{KL}(q \parallel p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$$

Key Properties:

Background: What is KL Divergence?

The Kullback-Leibler (KL) Divergence is a measure of how one probability distribution $q(z)$ is different from a second, reference probability distribution $p(z)$.

For discrete distributions:

$$\text{KL}(q \parallel p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$$

Key Properties:

- It is not a true “distance”:

- It is not symmetric: $\text{KL}(q \parallel p) \neq \text{KL}(p \parallel q)$

Background: What is KL Divergence?

The Kullback-Leibler (KL) Divergence is a measure of how one probability distribution $q(z)$ is different from a second, reference probability distribution $p(z)$.

For discrete distributions:

$$\text{KL}(q \parallel p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$$

Key Properties:

- **It is not a true “distance”:**
 - It is not symmetric: $\text{KL}(q \parallel p) \neq \text{KL}(p \parallel q)$
- **It is always non-negative:**
 - $\text{KL}(q \parallel p) \geq 0$

Background: What is KL Divergence?

The Kullback-Leibler (KL) Divergence is a measure of how one probability distribution $q(z)$ is different from a second, reference probability distribution $p(z)$.

For discrete distributions:

$$\text{KL}(q \parallel p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$$

Key Properties:

- **It is not a true “distance”:**
 - It is not symmetric: $\text{KL}(q \parallel p) \neq \text{KL}(p \parallel q)$
- **It is always non-negative:**
 - $\text{KL}(q \parallel p) \geq 0$
- **It is zero only if the distributions are identical:**
 - $\text{KL}(q \parallel p) = 0 \iff q(z) = p(z) \text{ for all } z.$

Background: What is KL Divergence?

The Kullback-Leibler (KL) Divergence is a measure of how one probability distribution $q(z)$ is different from a second, reference probability distribution $p(z)$.

For discrete distributions:

$$\text{KL}(q \parallel p) = \sum_z q(z) \log \frac{q(z)}{p(z)}$$

Key Properties:

- **It is not a true “distance”:**
 - It is not symmetric: $\text{KL}(q \parallel p) \neq \text{KL}(p \parallel q)$
- **It is always non-negative:**
 - $\text{KL}(q \parallel p) \geq 0$
- **It is zero only if the distributions are identical:**
 - $\text{KL}(q \parallel p) = 0 \iff q(z) = p(z) \text{ for all } z.$

In EM, we use KL divergence to measure how “bad” our approximation $q(\mathbf{z})$ is compared to the true posterior $p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})$.

The Evidence Lower Bound (ELBO)

- This lower bound is known as the Evidence Lower Bound (ELBO).

The Evidence Lower Bound (ELBO)

- This lower bound is known as the Evidence Lower Bound (ELBO).
- The log-likelihood (evidence) can be decomposed for *any* distribution $q(\mathbf{z})$:

$$\log p(\mathbf{x}|\boldsymbol{\theta}) = \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})) \quad (5)$$

where $\mathcal{L}(q, \boldsymbol{\theta})$ is the ELBO:

$$\mathcal{L}(q, \boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\boldsymbol{\theta})] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (6)$$

The Evidence Lower Bound (ELBO)

- This lower bound is known as the Evidence Lower Bound (ELBO).
- The log-likelihood (evidence) can be decomposed for *any* distribution $q(\mathbf{z})$:

$$\log p(\mathbf{x}|\boldsymbol{\theta}) = \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})) \quad (5)$$

where $\mathcal{L}(q, \boldsymbol{\theta})$ is the ELBO:

$$\mathcal{L}(q, \boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\boldsymbol{\theta})] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (6)$$

- Since $\text{KL}(\cdot) \geq 0$, the ELBO $\mathcal{L}(q, \boldsymbol{\theta})$ is always a *lower bound* on the true log-likelihood.

The Evidence Lower Bound (ELBO)

- This lower bound is known as the Evidence Lower Bound (ELBO).
- The log-likelihood (evidence) can be decomposed for *any* distribution $q(\mathbf{z})$:

$$\log p(\mathbf{x}|\boldsymbol{\theta}) = \mathcal{L}(q, \boldsymbol{\theta}) + \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})) \quad (5)$$

where $\mathcal{L}(q, \boldsymbol{\theta})$ is the ELBO:

$$\mathcal{L}(q, \boldsymbol{\theta}) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\boldsymbol{\theta})] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (6)$$

- Since $\text{KL}(\cdot) \geq 0$, the ELBO $\mathcal{L}(q, \boldsymbol{\theta})$ is always a *lower bound* on the true log-likelihood.
- EM alternates between:
 - **E-step:** Maximize \mathcal{L} w.r.t. q . This happens when $\text{KL} = 0$, so we set $q(\mathbf{z}) = p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta}^{(t)})$.
 - **M-step:** Maximize \mathcal{L} w.r.t. $\boldsymbol{\theta}$, keeping q fixed.

ELBO: Intuition and Properties

- The ELBO can be rewritten in two equivalent forms:

$$\mathcal{L}(q, \theta) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\theta)] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (7)$$

$$= \log p(\mathbf{x}|\theta) - \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \theta)) \quad (8)$$

ELBO: Intuition and Properties

- The ELBO can be rewritten in two equivalent forms:

$$\mathcal{L}(q, \theta) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\theta)] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (7)$$

$$= \log p(\mathbf{x}|\theta) - \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \theta)) \quad (8)$$

- The first form (7) shows ELBO as:
 - Expected complete-data log-likelihood (tries to fit data)
 - Minus entropy of q (tries to keep q “spread out”)

ELBO: Intuition and Properties

- The ELBO can be rewritten in two equivalent forms:

$$\mathcal{L}(q, \theta) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\theta)] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})] \quad (7)$$

$$= \log p(\mathbf{x}|\theta) - \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z}|\mathbf{x}, \theta)) \quad (8)$$

- The first form (7) shows ELBO as:
 - Expected complete-data log-likelihood (tries to fit data)
 - Minus entropy of q (tries to keep q “spread out”)
- The second form (8) shows ELBO as:
 - The true log-likelihood (what we want)
 - Minus the KL divergence (how bad our q is)

ELBO: Intuition and Properties

- The ELBO can be rewritten in two equivalent forms:

$$\mathcal{L}(q, \theta) = \mathbb{E}_{q(\mathbf{z})} [\log p(\mathbf{x}, \mathbf{z} | \theta)] - \mathbb{E}_{q(\mathbf{z})} [\log q(\mathbf{z})] \quad (7)$$

$$= \log p(\mathbf{x} | \theta) - \text{KL}(q(\mathbf{z}) \parallel p(\mathbf{z} | \mathbf{x}, \theta)) \quad (8)$$

- The first form (7) shows ELBO as:
 - Expected complete-data log-likelihood (tries to fit data)
 - Minus entropy of q (tries to keep q “spread out”)
- The second form (8) shows ELBO as:
 - The true log-likelihood (what we want)
 - Minus the KL divergence (how bad our q is)
- EM optimization balances:
 - Finding a good fit to the data through θ .
 - Finding a good approximation to the posterior through q .

The Expectation step (E-step) in Practice

Goal: Set $q_n(z_n) = p(z_n | \mathbf{x}_n, \boldsymbol{\theta}^{(t)})$.

The Expectation step (E-step) in Practice

Goal: Set $q_n(z_n) = p(z_n | \mathbf{x}_n, \boldsymbol{\theta}^{(t)})$.

How? We just need to find the probability for each component k :

$$\begin{aligned} q_{kn}^{(t)} &:= q_n(z_n = k) = p(z_n = k | \mathbf{x}_n, \boldsymbol{\theta}^{(t)}) \\ &= \frac{p(\mathbf{x}_n | z_n = k, \boldsymbol{\theta}^{(t)}) p(z_n = k | \boldsymbol{\theta}^{(t)})}{p(\mathbf{x}_n | \boldsymbol{\theta}^{(t)})} \quad (\text{Bayes' rule}) \\ &= \frac{p(\mathbf{x}_n | z_n = k, \boldsymbol{\theta}^{(t)}) p(z_n = k | \boldsymbol{\theta}^{(t)})}{\sum_{j=1}^K p(\mathbf{x}_n | z_n = j, \boldsymbol{\theta}^{(t)}) p(z_n = j | \boldsymbol{\theta}^{(t)})} \end{aligned}$$

The Expectation step (E-step) in Practice

Goal: Set $q_n(z_n) = p(z_n | \mathbf{x}_n, \boldsymbol{\theta}^{(t)})$.

How? We just need to find the probability for each component k :

$$\begin{aligned} q_{kn}^{(t)} &:= q_n(z_n = k) = p(z_n = k | \mathbf{x}_n, \boldsymbol{\theta}^{(t)}) \\ &= \frac{p(\mathbf{x}_n | z_n = k, \boldsymbol{\theta}^{(t)}) p(z_n = k | \boldsymbol{\theta}^{(t)})}{p(\mathbf{x}_n | \boldsymbol{\theta}^{(t)})} \quad (\text{Bayes' rule}) \\ &= \frac{p(\mathbf{x}_n | z_n = k, \boldsymbol{\theta}^{(t)}) p(z_n = k | \boldsymbol{\theta}^{(t)})}{\sum_{j=1}^K p(\mathbf{x}_n | z_n = j, \boldsymbol{\theta}^{(t)}) p(z_n = j | \boldsymbol{\theta}^{(t)})} \end{aligned}$$

This gives the famous E-step update rule:

$$q_{kn}^{(t)} = \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})}{\sum_{j=1}^K \pi_j^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j^{(t)}, \boldsymbol{\Sigma}_j^{(t)})} \quad (9)$$

This is the “soft assignment” or “responsibility” of component k for point n .

The Maximization step (M-step)

Goal: Maximize the lower bound w.r.t. θ , using the $q_{kn}^{(t)}$ we just found.

$$\max_{\theta} \sum_{n=1}^N \sum_{k=1}^K q_{kn}^{(t)} [\log \pi_k + \log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)] \quad (10)$$

The Maximization step (M-step)

Goal: Maximize the lower bound w.r.t. θ , using the $q_{kn}^{(t)}$ we just found.

$$\max_{\theta} \sum_{n=1}^N \sum_{k=1}^K q_{kn}^{(t)} [\log \pi_k + \log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)] \quad (10)$$

This maximization can be solved independently for $\boldsymbol{\mu}_k$, $\boldsymbol{\Sigma}_k$, and π_k . This results in intuitive “weighted MLE” updates:

$$\boldsymbol{\mu}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}} \quad (11)$$

$$\boldsymbol{\Sigma}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{\sum_{n=1}^N q_{kn}^{(t)}} \quad (12)$$

The Maximization step (M-step)

Goal: Maximize the lower bound w.r.t. θ , using the $q_{kn}^{(t)}$ we just found.

$$\max_{\theta} \sum_{n=1}^N \sum_{k=1}^K q_{kn}^{(t)} [\log \pi_k + \log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)] \quad (10)$$

This maximization can be solved independently for $\boldsymbol{\mu}_k$, $\boldsymbol{\Sigma}_k$, and π_k . This results in intuitive “weighted MLE” updates:

$$\boldsymbol{\mu}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}} \quad (11)$$

$$\boldsymbol{\Sigma}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{\sum_{n=1}^N q_{kn}^{(t)}} \quad (12)$$

For π_k , we maximize (10) with a Lagrangian for the constraint $\sum_k \pi_k = 1$:

$$\pi_k^{(t+1)} := \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)} \quad (13)$$

(See derivation slides for full proof.)

M-step: The Objective Function (Why Eq. 10?)

The M-step maximizes the ELBO w.r.t. θ . From our ELBO definition:

$$\mathcal{L}(q, \theta) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{x}, \mathbf{z}|\theta)] - \mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})]$$

Since the M-step optimizes for θ , we can ignore the second term, $\mathbb{E}_{q(\mathbf{z})}[\log q(\mathbf{z})]$, as it does not depend on θ . We are left with maximizing the expected complete-data log-likelihood:

$$Q(\theta, \theta^{(t)}) = \mathbb{E}_{q(\mathbf{z})}[\log p(\mathbf{X}, \mathbf{z}|\theta)] = \sum_{n=1}^N \mathbb{E}_{q_n(\mathbf{z}_n)}[\log p(\mathbf{x}_n, \mathbf{z}_n|\theta)]$$

We expand $\log p(\mathbf{x}_n, \mathbf{z}_n|\theta) = \log p(\mathbf{x}_n|\mathbf{z}_n, \theta) + \log p(z_n|\theta)$.

- $p(z_n = k|\theta) = \pi_k$
- $p(\mathbf{x}_n|z_n = k, \theta) = \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$

Using $q_{kn}^{(t)} := q_n(z_n = k)$ as the expectation weight, we get:

$$Q(\theta, \theta^{(t)}) = \sum_{n=1}^N \sum_{k=1}^K q_{kn}^{(t)} [\log \mathcal{N}(\mathbf{x}_n|\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) + \log \pi_k]$$

This is exactly the objective from the previous slide.

M-step Derivation: Mean μ_k

Goal: Find μ_k that maximizes the objective. We only need terms that depend on μ_k :

$$\mathcal{L}(\boldsymbol{\mu}_k) = \sum_{n=1}^N q_{kn}^{(t)} \log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) + \text{const} = \sum_{n=1}^N q_{kn}^{(t)} \left[\text{const} - \frac{1}{2} (\mathbf{x}_n - \boldsymbol{\mu}_k)^\top \boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k) \right]$$

Take the gradient w.r.t. μ_k and set to 0 (using $\nabla_{\mathbf{z}}(\mathbf{a} - \mathbf{z})^\top \mathbf{C}(\mathbf{a} - \mathbf{z}) = 2\mathbf{C}(\mathbf{z} - \mathbf{a})$):

$$\nabla_{\boldsymbol{\mu}_k} \mathcal{L} = \sum_{n=1}^N q_{kn}^{(t)} \left[-\frac{1}{2} \cdot 2\boldsymbol{\Sigma}_k^{-1} (\boldsymbol{\mu}_k - \mathbf{x}_n) \right] = \sum_{n=1}^N q_{kn}^{(t)} \left[\boldsymbol{\Sigma}_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k) \right] \stackrel{\text{set}}{=} 0$$

Multiply by $\boldsymbol{\Sigma}_k$ (which is invertible) to simplify:

$$\begin{aligned} \sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k) &= 0, & \sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n &= \sum_{n=1}^N q_{kn}^{(t)} \boldsymbol{\mu}_k, & \sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n &= \boldsymbol{\mu}_k \left(\sum_{n=1}^N q_{kn}^{(t)} \right), \\ \implies \boldsymbol{\mu}_k^{(t+1)} &= \frac{\sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}} \end{aligned}$$

M-step Derivation: Covariance Σ_k

Goal: Find Σ_k that maximizes the objective. We only need terms that depend on Σ_k :

$$\begin{aligned}\mathcal{L}(\Sigma_k) &= \sum_{n=1}^N q_{kn}^{(t)} \log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \Sigma_k) + \text{const} = \sum_{n=1}^N q_{kn}^{(t)} \left[\log \left(\frac{1}{(2\pi)^{D/2} |\Sigma_k|^{1/2}} \right) - \frac{1}{2} (\dots)^\top \Sigma_k^{-1} (\dots) \right] \\ &= \sum_{n=1}^N q_{kn}^{(t)} \left[-\frac{1}{2} \log |\Sigma_k| - \frac{1}{2} (\mathbf{x}_n - \boldsymbol{\mu}_k)^\top \Sigma_k^{-1} (\mathbf{x}_n - \boldsymbol{\mu}_k) \right]\end{aligned}$$

This objective is a weighted log-likelihood for a Gaussian distribution.

- The Maximum Likelihood Estimate (MLE) for the covariance of a Gaussian is the sample covariance.
- For our *weighted* data (where \mathbf{x}_n has weight $q_{kn}^{(t)}$), the MLE is the *weighted sample covariance*.

By taking the derivative w.r.t. Σ_k (using matrix calculus), setting $\boldsymbol{\mu}_k = \boldsymbol{\mu}_k^{(t+1)}$ (the new mean we just found), and setting the result to 0, we get:

$$\Sigma_k^{(t+1)} = \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{\sum_{n=1}^N q_{kn}^{(t)}}$$

M-step Derivation: Mixture π_k

Goal: Maximize w.r.t. π_k , subject to the constraint $\sum_{k=1}^K \pi_k = 1$.

- Objective terms for π : $\mathcal{L}(\pi) = \sum_{n=1}^N \sum_{k=1}^K q_{kn}^{(t)} \log \pi_k$
- Constraint: $\sum_{k=1}^K \pi_k - 1 = 0$

We use a **Lagrange multiplier** λ , as suggested: $\mathcal{J}(\pi, \lambda) = \mathcal{L}(\pi) - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right)$.

Take the partial derivative w.r.t. a single π_k and set to 0:

$$\frac{\partial \mathcal{J}}{\partial \pi_k} = \sum_{n=1}^N \frac{q_{kn}^{(t)}}{\pi_k} - \lambda \stackrel{\text{set}}{=} 0 \implies \sum_{n=1}^N q_{kn}^{(t)} = \lambda \pi_k$$

Let $N_k = \sum_{n=1}^N q_{kn}^{(t)}$, then $N_k = \lambda \pi_k$. Sum over all k to find λ :

$$\sum_{k=1}^K N_k = \sum_{k=1}^K \lambda \pi_k \implies \sum_{k=1}^K \underbrace{\sum_{n=1}^N q_{kn}^{(t)}}_{=1} = \lambda \underbrace{\sum_{k=1}^K \pi_k}_{=1}, \quad \sum_{n=1}^N \underbrace{\sum_{k=1}^K q_{kn}^{(t)}}_{=1} = \lambda \implies N = \lambda.$$

Substitute $\lambda = N$ back into $N_k = \lambda \pi_k$: $N_k = N \pi_k \implies \pi_k^{(t+1)} = \frac{N_k}{N} = \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)}$.

Summary of EM for GMM

Initialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$ and iterate between the E and M step, until $\mathcal{L}(\theta)$ stabilizes.

Summary of EM for GMM

Initialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$ and iterate between the E and M step, until $\mathcal{L}(\theta)$ stabilizes.

- ① **E-step:** Compute assignments $q_{kn}^{(t)}$ (posterior probabilities):

$$q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})}{\sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})} \quad (14)$$

Summary of EM for GMM

Initialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$ and iterate between the E and M step, until $\mathcal{L}(\theta)$ stabilizes.

- ① **E-step:** Compute assignments $q_{kn}^{(t)}$ (posterior probabilities):

$$q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})}{\sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})} \quad (14)$$

- ② Compute the marginal likelihood (cost).

$$\mathcal{L}(\theta^{(t)}) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)}) \quad (15)$$

Summary of EM for GMM

Initialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$ and iterate between the E and M step, until $\mathcal{L}(\theta)$ stabilizes.

- ① **E-step:** Compute assignments $q_{kn}^{(t)}$ (posterior probabilities):

$$q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})}{\sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})} \quad (14)$$

- ② Compute the marginal likelihood (cost).

$$\mathcal{L}(\theta^{(t)}) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)}) \quad (15)$$

- ③ **M-step:** Update $\boldsymbol{\mu}_k^{(t+1)}, \boldsymbol{\Sigma}_k^{(t+1)}, \pi_k^{(t+1)}$.

$$\boldsymbol{\mu}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}}, \quad \boldsymbol{\Sigma}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{\sum_{n=1}^N q_{kn}^{(t)}}, \quad (16)$$

$$\pi_k^{(t+1)} := \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)} \quad (17)$$

Summary of EM for GMM

Initialize $\mu^{(1)}, \Sigma^{(1)}, \pi^{(1)}$ and iterate between the E and M step, until $\mathcal{L}(\theta)$ stabilizes.

- ① **E-step:** Compute assignments $q_{kn}^{(t)}$ (posterior probabilities):

$$q_{kn}^{(t)} := \frac{\pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})}{\sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)})} \quad (14)$$

- ② Compute the marginal likelihood (cost).

$$\mathcal{L}(\theta^{(t)}) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k^{(t)} \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k^{(t)}, \boldsymbol{\Sigma}_k^{(t)}) \quad (15)$$

- ③ **M-step:** Update $\boldsymbol{\mu}_k^{(t+1)}, \boldsymbol{\Sigma}_k^{(t+1)}, \pi_k^{(t+1)}$.

$$\boldsymbol{\mu}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} \mathbf{x}_n}{\sum_{n=1}^N q_{kn}^{(t)}}, \quad \boldsymbol{\Sigma}_k^{(t+1)} := \frac{\sum_{n=1}^N q_{kn}^{(t)} (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)}) (\mathbf{x}_n - \boldsymbol{\mu}_k^{(t+1)})^\top}{\sum_{n=1}^N q_{kn}^{(t)}}, \quad (16)$$

$$\pi_k^{(t+1)} := \frac{1}{N} \sum_{n=1}^N q_{kn}^{(t)} \quad (17)$$

If we let the covariance be diagonal i.e. $\boldsymbol{\Sigma}_k := \sigma^2 \mathbf{I}$, EM algorithm is same as K-means as $\sigma^2 \rightarrow 0$.

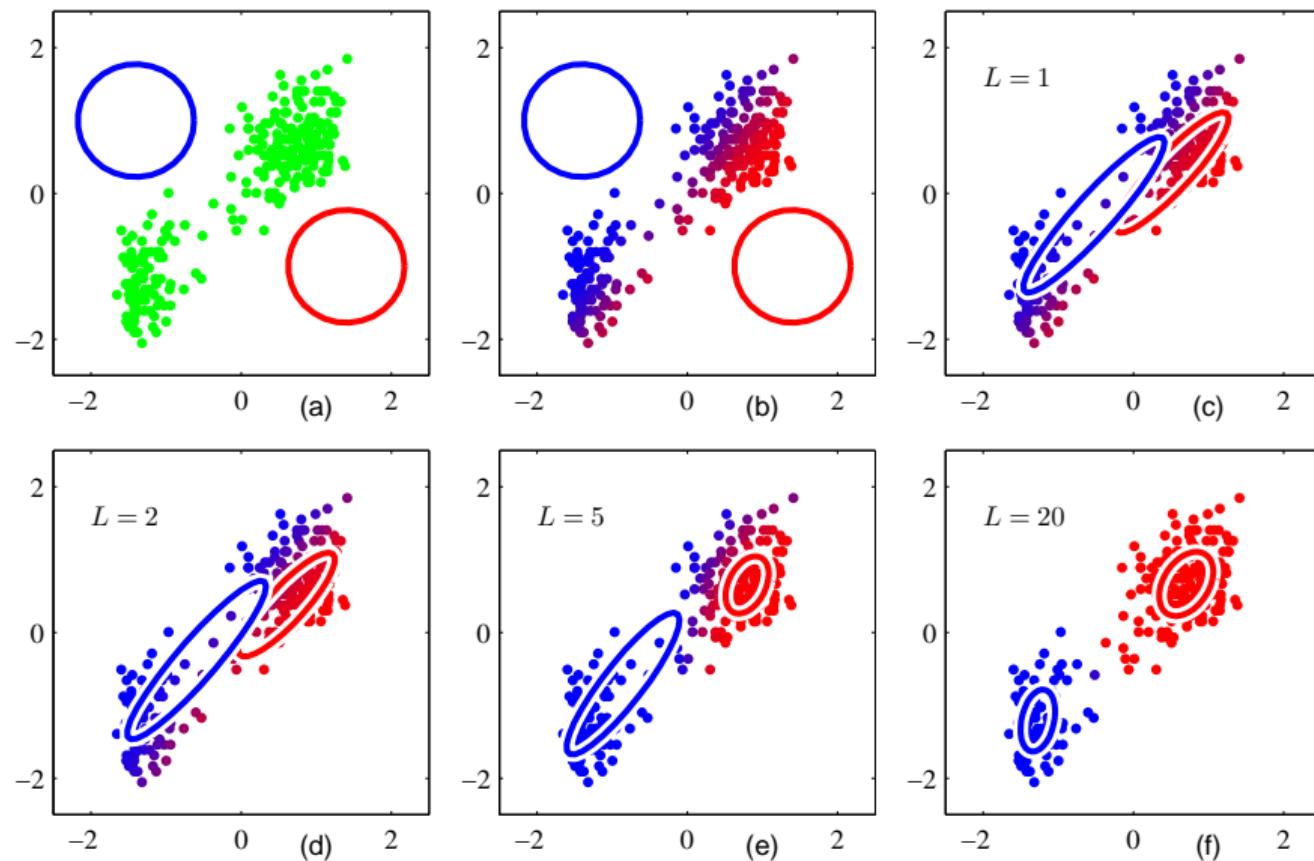
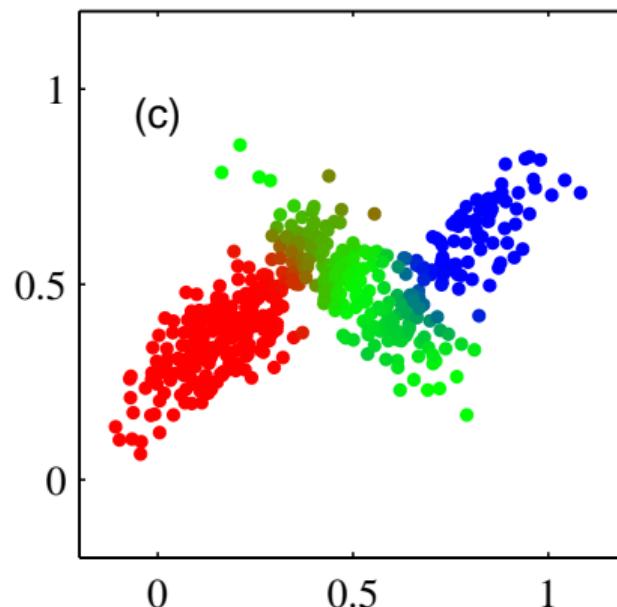


Figure: EM algorithm for GMM: Visualization of the iterative fitting process

Posterior distribution

We now show that $q_{kn}^{(t)}$ is the posterior distribution of the latent variable, i.e. $q_{kn}^{(t)} = p(z_n = k | \mathbf{x}_n, \boldsymbol{\theta}^{(t)})$

$$\underbrace{p(\mathbf{x}_n, z_n | \boldsymbol{\theta})}_{\text{joint}} = \underbrace{p(\mathbf{x}_n | z_n, \boldsymbol{\theta})}_{\text{likelihood}} \underbrace{p(z_n | \boldsymbol{\theta})}_{\text{prior}} = \underbrace{p(z_n | \mathbf{x}_n, \boldsymbol{\theta})}_{\text{posterior}} \underbrace{p(\mathbf{x}_n | \boldsymbol{\theta})}_{\text{marginal likelihood}} \quad (18)$$



The relationship to K-means

- The EM algorithm for GMMs is reminiscent of the K-means clustering algorithm.

The relationship to K-means

- The EM algorithm for GMMs is reminiscent of the K-means clustering algorithm.
- Instead of the “hard” cluster assignments in K-means, EM uses “soft” assignments q_{kn} .

$$q_{kn} = p(z_n = k \mid \mathbf{x}_n, \boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}) \quad (19)$$

The relationship to K-means

- The EM algorithm for GMMs is reminiscent of the K-means clustering algorithm.
- Instead of the “hard” cluster assignments in K-means, EM uses “soft” assignments q_{kn} .

$$q_{kn} = p(z_n = k \mid \mathbf{x}_n, \boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}) \quad (19)$$

- “Soft” means our assignments are probabilities taking values in $[0, 1]$.

The relationship to K-means

- The EM algorithm for GMMs is reminiscent of the K-means clustering algorithm.
- Instead of the “hard” cluster assignments in K-means, EM uses “soft” assignments q_{kn} .

$$q_{kn} = p(z_n = k \mid \mathbf{x}_n, \boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}) \quad (19)$$

- “Soft” means our assignments are probabilities taking values in $[0, 1]$.
- “Hard” assignments take values in $\{0, 1\}$ or $\{1, \dots, K\}$.

The relationship to K-means

- The EM algorithm for GMMs is reminiscent of the K-means clustering algorithm.
- Instead of the “hard” cluster assignments in K-means, EM uses “soft” assignments q_{kn} .

$$q_{kn} = p(z_n = k \mid \mathbf{x}_n, \boldsymbol{\mu}, \boldsymbol{\Sigma}, \boldsymbol{\pi}) \quad (19)$$

- “Soft” means our assignments are probabilities taking values in $[0, 1]$.
- “Hard” assignments take values in $\{0, 1\}$ or $\{1, \dots, K\}$.
- Like K-means, EM is also susceptible to local optima, so initializing at several different parameter values may be a good strategy.

EM Convergence Properties

- The EM algorithm is guaranteed to increase the likelihood at each iteration:

$$\mathcal{L}(\boldsymbol{\theta}^{(t+1)}) \geq \mathcal{L}(\boldsymbol{\theta}^{(t)}) \quad (20)$$

EM Convergence Properties

- The EM algorithm is guaranteed to increase the likelihood at each iteration:

$$\mathcal{L}(\boldsymbol{\theta}^{(t+1)}) \geq \mathcal{L}(\boldsymbol{\theta}^{(t)}) \quad (20)$$

- This guarantees convergence to a local maximum of the likelihood function.

EM Convergence Properties

- The EM algorithm is guaranteed to increase the likelihood at each iteration:

$$\mathcal{L}(\boldsymbol{\theta}^{(t+1)}) \geq \mathcal{L}(\boldsymbol{\theta}^{(t)}) \quad (20)$$

- This guarantees convergence to a local maximum of the likelihood function.
- However, like K-means, it may not find the global maximum.

EM Convergence Properties

- The EM algorithm is guaranteed to increase the likelihood at each iteration:

$$\mathcal{L}(\boldsymbol{\theta}^{(t+1)}) \geq \mathcal{L}(\boldsymbol{\theta}^{(t)}) \quad (20)$$

- This guarantees convergence to a local maximum of the likelihood function.
- However, like K-means, it may not find the global maximum.
- The rate of convergence depends on the fraction of information about the parameters contained in the latent variables.

EM Convergence Properties

- The EM algorithm is guaranteed to increase the likelihood at each iteration:

$$\mathcal{L}(\boldsymbol{\theta}^{(t+1)}) \geq \mathcal{L}(\boldsymbol{\theta}^{(t)}) \quad (20)$$

- This guarantees convergence to a local maximum of the likelihood function.
- However, like K-means, it may not find the global maximum.
- The rate of convergence depends on the fraction of information about the parameters contained in the latent variables.
- In practice, we monitor the change in log-likelihood and stop when it falls below a threshold.

EM and Variational Inference

- The ELBO perspective reveals that EM is a special case of variational inference.

EM and Variational Inference

- The ELBO perspective reveals that EM is a special case of [variational inference](#).
- Variational inference approximates intractable posterior distributions with a tractable distribution $q(\mathbf{z})$.

EM and Variational Inference

- The ELBO perspective reveals that EM is a special case of [variational inference](#).
- Variational inference approximates intractable posterior distributions with a tractable distribution $q(\mathbf{z})$.
- In standard EM:
 - We set $q(\mathbf{z}) = p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta}^{(t)})$ exactly in the E-step.
 - This is possible because the posterior is tractable for models like GMM.

EM and Variational Inference

- The ELBO perspective reveals that EM is a special case of **variational inference**.
- Variational inference approximates intractable posterior distributions with a tractable distribution $q(\mathbf{z})$.
- In standard EM:
 - We set $q(\mathbf{z}) = p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta}^{(t)})$ exactly in the E-step.
 - This is possible because the posterior is tractable for models like GMM.
- In **variational EM**:
 - The posterior $p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})$ is intractable.
 - We restrict $q(\mathbf{z})$ to a simpler family of distributions.
 - E-step: Find $q(\mathbf{z})$ that maximizes ELBO (minimizes KL-divergence).
 - M-step: Same as standard EM (maximize ELBO w.r.t. $\boldsymbol{\theta}$).

EM and Variational Inference

- The ELBO perspective reveals that EM is a special case of [variational inference](#).
- Variational inference approximates intractable posterior distributions with a tractable distribution $q(\mathbf{z})$.
- In standard EM:
 - We set $q(\mathbf{z}) = p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta}^{(t)})$ exactly in the E-step.
 - This is possible because the posterior is tractable for models like GMM.
- In [variational EM](#):
 - The posterior $p(\mathbf{z}|\mathbf{x}, \boldsymbol{\theta})$ is intractable.
 - We restrict $q(\mathbf{z})$ to a simpler family of distributions.
 - E-step: Find $q(\mathbf{z})$ that maximizes ELBO (minimizes KL-divergence).
 - M-step: Same as standard EM (maximize ELBO w.r.t. $\boldsymbol{\theta}$).
- This connection has led to powerful techniques like variational autoencoders (VAEs).

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:
 - Missing data problems

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:
 - Missing data problems
 - Hidden Markov Models

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:
 - Missing data problems
 - Hidden Markov Models
 - Mixture of experts models

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:
 - Missing data problems
 - Hidden Markov Models
 - Mixture of experts models
 - Bayesian networks with hidden variables

Applications of EM algorithm

- The EM algorithm can be applied to many problems beyond GMMs:
 - Missing data problems
 - Hidden Markov Models
 - Mixture of experts models
 - Bayesian networks with hidden variables
- The same principles apply: identify latent variables, compute their expected values, and maximize parameters.

Practical Considerations for GMM Implementation

- Initialization strategies:

Practical Considerations for GMM Implementation

- **Initialization strategies:**
 - Random initialization of means and covariances

Practical Considerations for GMM Implementation

- **Initialization strategies:**
 - Random initialization of means and covariances
 - K-means initialization (use K-means clusters to initialize GMM parameters)

Practical Considerations for GMM Implementation

- **Initialization strategies:**

- Random initialization of means and covariances
- K-means initialization (use K-means clusters to initialize GMM parameters)
- Multiple restarts to avoid local optima

Practical Considerations for GMM Implementation

- **Initialization strategies:**
 - Random initialization of means and covariances
 - K-means initialization (use K-means clusters to initialize GMM parameters)
 - Multiple restarts to avoid local optima
- **Numerical stability:**

Practical Considerations for GMM Implementation

- **Initialization strategies:**
 - Random initialization of means and covariances
 - K-means initialization (use K-means clusters to initialize GMM parameters)
 - Multiple restarts to avoid local optima
- **Numerical stability:**
 - Adding small values to diagonal of covariance matrices

Practical Considerations for GMM Implementation

- **Initialization strategies:**
 - Random initialization of means and covariances
 - K-means initialization (use K-means clusters to initialize GMM parameters)
 - Multiple restarts to avoid local optima
- **Numerical stability:**
 - Adding small values to diagonal of covariance matrices
 - Performing computations in log space to avoid underflow

Practical Considerations for GMM Implementation

- **Initialization strategies:**

- Random initialization of means and covariances
- K-means initialization (use K-means clusters to initialize GMM parameters)
- Multiple restarts to avoid local optima

- **Numerical stability:**

- Adding small values to diagonal of covariance matrices
- Performing computations in log space to avoid underflow

- **Model selection:**

Practical Considerations for GMM Implementation

- **Initialization strategies:**

- Random initialization of means and covariances
- K-means initialization (use K-means clusters to initialize GMM parameters)
- Multiple restarts to avoid local optima

- **Numerical stability:**

- Adding small values to diagonal of covariance matrices
- Performing computations in log space to avoid underflow

- **Model selection:**

- Determining the optimal number of components K

Practical Considerations for GMM Implementation

- **Initialization strategies:**

- Random initialization of means and covariances
- K-means initialization (use K-means clusters to initialize GMM parameters)
- Multiple restarts to avoid local optima

- **Numerical stability:**

- Adding small values to diagonal of covariance matrices
- Performing computations in log space to avoid underflow

- **Model selection:**

- Determining the optimal number of components K
- Using criteria like BIC (Bayesian Information Criterion) or AIC (Akaike Information Criterion)

Table of Contents

1 EM Algorithm

- Expectation-Maximization Algorithm
- Supp.

A Deeper Look: Unifying K-means and GMM

At first glance, K-means and GMM (via EM) seem very different:

K-means

- **Goal:** Minimize Sum of Squares
- **Loss:** $\mathcal{L} = \sum_{n,k} z_{nk} \|x_n - \mu_k\|^2$
- **Method:** Simple geometry, hard assignments.

GMM (EM)

- **Goal:** Maximize Log-Likelihood
- **Loss:** $\mathcal{L} = \sum_n \log \sum_k \pi_k \mathcal{N}(\dots)$
- **Method:** Complex, probabilistic, soft assignments.

A Deeper Look: Unifying K-means and GMM

At first glance, K-means and GMM (via EM) seem very different:

K-means

- **Goal:** Minimize Sum of Squares
- **Loss:** $\mathcal{L} = \sum_{n,k} z_{nk} \|x_n - \mu_k\|^2$
- **Method:** Simple geometry, hard assignments.

GMM (EM)

- **Goal:** Maximize Log-Likelihood
- **Loss:** $\mathcal{L} = \sum_n \log \sum_k \pi_k \mathcal{N}(\dots)$
- **Method:** Complex, probabilistic, soft assignments.

Key Idea: What if they are both instances of the *same* general algorithm, just using different “distance” functions?

The Unifying Tool: Bregman Divergence

A **Bregman divergence** is a “distance-like” function (not a true distance, as $D(p, q) \neq D(q, p)$) generated by any strictly convex function F .

Definition:

$$D_F(p, q) = \underbrace{F(p)}_{\text{Function at } p} - \left[\underbrace{F(q) + \nabla F(q)^\top (p - q)}_{\text{Taylor approx. at } q, \text{ eval. at } p} \right]$$

Geometric Intuition: It is the vertical gap between the function F at point p and the tangent line of F at point q .

Connection 1: K-means is Bregman Clustering

K-means minimizes the [Squared Euclidean Distance](#).

Claim: The Squared Euclidean Distance *is* a Bregman divergence.

Proof:

- Let the convex function be $F(\mathbf{x}) = \|\mathbf{x}\|^2 = \mathbf{x}^\top \mathbf{x}$.
- Its gradient is $\nabla F(\mathbf{x}) = 2\mathbf{x}$.

Now, plug this into the definition of $D_F(\mathbf{x}, \boldsymbol{\mu})$:

$$\begin{aligned} D_F(\mathbf{x}, \boldsymbol{\mu}) &= F(\mathbf{x}) - \left[F(\boldsymbol{\mu}) + \nabla F(\boldsymbol{\mu})^\top (\mathbf{x} - \boldsymbol{\mu}) \right] \\ &= \|\mathbf{x}\|^2 - \left[\|\boldsymbol{\mu}\|^2 + (2\boldsymbol{\mu})^\top (\mathbf{x} - \boldsymbol{\mu}) \right] = \|\mathbf{x}\|^2 - \|\boldsymbol{\mu}\|^2 - 2\boldsymbol{\mu}^\top \mathbf{x} + 2\|\boldsymbol{\mu}\|^2 \\ &= \|\mathbf{x}\|^2 - 2\boldsymbol{\mu}^\top \mathbf{x} + \|\boldsymbol{\mu}\|^2 = \|\mathbf{x} - \boldsymbol{\mu}\|_2^2 \end{aligned}$$

K-means is a clustering algorithm that
minimizes the Bregman divergence generated by $F(\mathbf{x}) = \|\mathbf{x}\|^2$.

Connection 2: GMM (EM) is Bregman Clustering

GMM-EM maximizes the log-likelihood, which is equivalent to minimizing the **negative log-likelihood**.

Key Insight:

- The **negative log-likelihood** of any distribution in the exponential family is a Bregman divergence.
- A Gaussian distribution $\mathcal{N}(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\Sigma})$ is in the exponential family.
- Therefore, the GMM “distance” (loss) for a point \mathbf{x}_n to a cluster k is also a Bregman divergence:

$$D_{\text{GMM}}(\mathbf{x}_n, \theta_k) = -\log \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

Connecting to K-means:

- If we set $\boldsymbol{\Sigma}_k = \sigma^2 \mathbf{I}$ (spherical clusters) and $\sigma^2 \rightarrow 0$, this divergence becomes:

$$D_{\text{GMM}} \propto \|\mathbf{x}_n - \boldsymbol{\mu}_k\|_2^2$$

- This exactly matches previous slide’s conclusion: GMM-EM becomes K-means when $\boldsymbol{\Sigma}_k$ is isotropic and $\sigma^2 \rightarrow 0$.

The Unification: Hard vs. Soft Clustering

Both algorithms are just “Bregman clustering”, differing only in the *divergence* used and the *assignment* type.

K-means (Hard Clustering)

- **Divergence:**

$$D_F(\mathbf{x}, \boldsymbol{\mu}) = \|\mathbf{x} - \boldsymbol{\mu}\|^2$$

- **E-step:** Hard Assignment. Find the *single k* that minimizes the divergence:

$$z_{nk} = 1 \iff k = \arg \min_j D_F(\mathbf{x}_n, \boldsymbol{\mu}_j)$$

- **M-step:** Update $\boldsymbol{\mu}_k$ to be the mean (centroid) of all points assigned to it.

GMM (EM) (Soft Clustering)

- **Divergence:**

$$D_{\text{GMM}}(\mathbf{x}, \theta_k) = -\log \mathcal{N}(\mathbf{x} | \theta_k)$$

- **E-step:** Soft Assignment. Compute $q_{kn}^{(t)}$ (responsibility) based on the divergence (and prior π_k).

$$q_{kn}^{(t)} \propto \pi_k \cdot \exp(-D_{\text{GMM}}(\mathbf{x}_n, \theta_k))$$

- **M-step:** Update θ_k to be the *weighted mean* (centroid) of all points.