

# Lecture 10: Kernel Methods and Support Vector Machines

**Tao LIN**

SoE, Westlake University

November 18, 2025



1 Support Vector Machines (SVM)

2 Kernel Methods and the Kernel Trick

# Table of Contents

- ① Support Vector Machines (SVM)
- ② Kernel Methods and the Kernel Trick

# A Brief History

- Support Vector Machines are a set of powerful supervised learning methods based on the concept of **maximizing the margin**.

# A Brief History

- Support Vector Machines are a set of powerful supervised learning methods based on the concept of **maximizing the margin**.
- The idea was developed by **Vladimir Vapnik** and colleagues.

# A Brief History

- Support Vector Machines are a set of powerful supervised learning methods based on the concept of **maximizing the margin**.
- The idea was developed by **Vladimir Vapnik** and colleagues.
- **Key Papers:**
  - **1992 (Linearly Separable):** Boser, Guyon, and Vapnik introduce the **optimal margin classifier**.
    - Solution is a linear combination of “**supporting patterns**” (Support Vectors).
  - **1995 (Non-Separable):** Cortes and Vapnik extend the idea to **non-separable data** (the “Soft Margin”).
    - Introduces non-linear mapping to a high-dimension feature space.

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .
- **Least Squares:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \mathbf{w})^2$

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .
- **Least Squares:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \mathbf{w})^2$
- **Logistic Regression:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w}$

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .
- **Least Squares:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \mathbf{w})^2$
- **Logistic Regression:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w}$
- It's often more convenient to use labels  $y_n \in \{-1, 1\}$ .

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .
- **Least Squares:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \mathbf{w})^2$
- **Logistic Regression:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w}$
- It's often more convenient to use labels  $y_n \in \{-1, 1\}$ .
- We use the mapping:  $y_n = 2\tilde{y}_n - 1 \leftrightarrow \tilde{y}_n = \frac{1}{2}(y_n + 1)$ .

# Motivation: Reframing Classification Losses

- Let's reconsider binary classification with data  $(\mathbf{x}_n, \tilde{y}_n)$ , where  $\tilde{y}_n \in \{0, 1\}$ .
- Least Squares:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \mathbf{w})^2$
- Logistic Regression:**  $\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w}$
- It's often more convenient to use labels  $y_n \in \{-1, 1\}$ .
- We use the mapping:  $y_n = 2\tilde{y}_n - 1 \leftrightarrow \tilde{y}_n = \frac{1}{2}(y_n + 1)$ .
- Let's see how our loss functions change with this new  $y_n$ .

## Motivation: Deriving the MSE Loss

- For Least Squares, let's assume bias  $e_0$  is component 0. Define  $\tilde{\mathbf{w}} = \frac{1}{2}(\mathbf{w} + e_0)$  (this is a one-to-one mapping).

# Motivation: Deriving the MSE Loss

- For Least Squares, let's assume bias  $e_0$  is component 0. Define  $\tilde{\mathbf{w}} = \frac{1}{2}(\mathbf{w} + e_0)$  (this is a one-to-one mapping).
- The derivation becomes:

$$\begin{aligned} 4 \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \tilde{\mathbf{w}})^2 &= 4 \sum_{n=1}^N \left( \frac{1}{2}(y_n + 1) - \frac{1}{2}\mathbf{x}_n^\top (\mathbf{w} + e_0) \right)^2 \\ &= \sum_{n=1}^N ((y_n + 1) - \mathbf{x}_n^\top (\mathbf{w} + e_0))^2 = \sum_{n=1}^N (y_n - \mathbf{x}_n^\top \mathbf{w})^2 \\ &\stackrel{(a)}{=} \sum_{n=1}^N (1 - y_n \mathbf{x}_n^\top \mathbf{w})^2 \end{aligned}$$

# Motivation: Deriving the MSE Loss

- For Least Squares, let's assume bias  $e_0$  is component 0. Define  $\tilde{\mathbf{w}} = \frac{1}{2}(\mathbf{w} + e_0)$  (this is a one-to-one mapping).
- The derivation becomes:

$$\begin{aligned}
 4 \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \tilde{\mathbf{w}})^2 &= 4 \sum_{n=1}^N \left( \frac{1}{2}(y_n + 1) - \frac{1}{2}\mathbf{x}_n^\top (\mathbf{w} + e_0) \right)^2 \\
 &= \sum_{n=1}^N ((y_n + 1) - \mathbf{x}_n^\top (\mathbf{w} + e_0))^2 = \sum_{n=1}^N (y_n - \mathbf{x}_n^\top \mathbf{w})^2 \\
 &\stackrel{(a)}{=} \sum_{n=1}^N (1 - y_n \mathbf{x}_n^\top \mathbf{w})^2
 \end{aligned}$$

- In step (a), we used  $y_n^2 = 1$  (since  $y_n \in \{\pm 1\}$ ).

# Motivation: Deriving the MSE Loss

- For Least Squares, let's assume bias  $e_0$  is component 0. Define  $\tilde{\mathbf{w}} = \frac{1}{2}(\mathbf{w} + e_0)$  (this is a one-to-one mapping).
- The derivation becomes:

$$\begin{aligned}
 4 \sum_{n=1}^N (\tilde{y}_n - \mathbf{x}_n^\top \tilde{\mathbf{w}})^2 &= 4 \sum_{n=1}^N \left( \frac{1}{2}(y_n + 1) - \frac{1}{2}\mathbf{x}_n^\top (\mathbf{w} + e_0) \right)^2 \\
 &= \sum_{n=1}^N ((y_n + 1) - \mathbf{x}_n^\top (\mathbf{w} + e_0))^2 = \sum_{n=1}^N (y_n - \mathbf{x}_n^\top \mathbf{w})^2 \\
 &\stackrel{(a)}{=} \sum_{n=1}^N (1 - y_n \mathbf{x}_n^\top \mathbf{w})^2
 \end{aligned}$$

- In step (a), we used  $y_n^2 = 1$  (since  $y_n \in \{\pm 1\}$ ).
- This gives the **MSE loss** in margin form:

$$\text{MSE}(z, y) = (1 - yz)^2 \quad \text{where } z = \mathbf{x}_n^\top \mathbf{w}$$

# Motivation: Deriving the Logistic Loss

- For Logistic Regression, we can also rewrite the loss:

$$\sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w} = \sum_{n=1}^N \log(1 + e^{-y_n \mathbf{x}_n^\top \mathbf{w}})$$

# Motivation: Deriving the Logistic Loss

- For Logistic Regression, we can also rewrite the loss:

$$\sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w} = \sum_{n=1}^N \log(1 + e^{-y_n \mathbf{x}_n^\top \mathbf{w}})$$

- This is easily seen by checking the two cases:

- $\tilde{y}_n = 1 \leftrightarrow y_n = 1$ :

$$\log(1 + e^{\mathbf{w}^\top \mathbf{x}}) - \mathbf{w}^\top \mathbf{x} = \log(e^{\mathbf{w}^\top \mathbf{x}}(e^{-\mathbf{w}^\top \mathbf{x}} + 1)) - \mathbf{w}^\top \mathbf{x} \quad (1)$$

$$= \log(e^{\mathbf{w}^\top \mathbf{x}}) + \log(1 + e^{-\mathbf{w}^\top \mathbf{x}}) - \mathbf{w}^\top \mathbf{x} \quad (2)$$

$$= \log(1 + e^{-\mathbf{w}^\top \mathbf{x}}) \quad (3)$$

- $\tilde{y}_n = 0 \leftrightarrow y_n = -1$ :

$$\log(1 + e^{\mathbf{w}^\top \mathbf{x}}) - 0 = \log(1 + e^{-(\mathbf{-1})^\top \mathbf{w}^\top \mathbf{x}}) \quad (4)$$

# Motivation: Deriving the Logistic Loss

- For Logistic Regression, we can also rewrite the loss:

$$\sum_{n=1}^N \log(1 + e^{\mathbf{x}_n^\top \mathbf{w}}) - \tilde{y}_n \mathbf{x}_n^\top \mathbf{w} = \sum_{n=1}^N \log(1 + e^{-y_n \mathbf{x}_n^\top \mathbf{w}})$$

- This is easily seen by checking the two cases:

- $\tilde{y}_n = 1 \leftrightarrow y_n = 1$ :

$$\log(1 + e^{\mathbf{w}^\top \mathbf{x}}) - \mathbf{w}^\top \mathbf{x} = \log(e^{\mathbf{w}^\top \mathbf{x}}(e^{-\mathbf{w}^\top \mathbf{x}} + 1)) - \mathbf{w}^\top \mathbf{x} \quad (1)$$

$$= \log(e^{\mathbf{w}^\top \mathbf{x}}) + \log(1 + e^{-\mathbf{w}^\top \mathbf{x}}) - \mathbf{w}^\top \mathbf{x} \quad (2)$$

$$= \log(1 + e^{-\mathbf{w}^\top \mathbf{x}}) \quad (3)$$

- $\tilde{y}_n = 0 \leftrightarrow y_n = -1$ :

$$\log(1 + e^{\mathbf{w}^\top \mathbf{x}}) - 0 = \log(1 + e^{-(\mathbf{-1})^\top \mathbf{w}^\top \mathbf{x}}) \quad (4)$$

- This gives the Logistic loss in margin form:

$$\text{LogisticLoss}(z, y) = \log(1 + \exp(-yz))$$

# Motivation: The Hinge Loss (SVM)

- Both MSE and Logistic Loss are functions of the functional margin  $yz$ .

# Motivation: The Hinge Loss (SVM)

- Both MSE and Logistic Loss are functions of the functional margin  $yz$ .
- Support Vector Machines (SVMs) are formed by using a third loss, the Hinge Loss:

$$\text{Hinge}(z, y) = [1 - yz]_+ = \max\{0, 1 - yz\}$$

# Motivation: The Hinge Loss (SVM)

- Both MSE and Logistic Loss are functions of the functional margin  $yz$ .
- Support Vector Machines (SVMs) are formed by using a third loss, the Hinge Loss:

$$\text{Hinge}(z, y) = [1 - yz]_+ = \max\{0, 1 - yz\}$$

- ...and adding a regularization term.

# Motivation: The Hinge Loss (SVM)

- Both MSE and Logistic Loss are functions of the functional margin  $yz$ .
- Support Vector Machines (SVMs) are formed by using a third loss, the Hinge Loss:


$$\text{Hinge}(z, y) = [1 - yz]_+ = \max\{0, 1 - yz\}$$

- ...and adding a regularization term.
- This leads directly to the SVM optimization problem we will study.

# The 0-1 Loss and Convex Relaxations

- The “ideal” loss for classification is the **Zero-one loss**:

$$L_{0-1}(\eta) = \mathbf{1}_{\eta < 0} \quad \text{where } \eta = y\mathbf{x}^\top \mathbf{w}$$




# The 0-1 Loss and Convex Relaxations

- The “ideal” loss for classification is the [Zero-one loss](#):

$$L_{0-1}(\eta) = \mathbf{1}_{\eta < 0} \quad \text{where } \eta = y\mathbf{x}^\top \mathbf{w}$$

- This corresponds to [Empirical Risk Minimization \(ERM\)](#).




# The 0-1 Loss and Convex Relaxations

- The “ideal” loss for classification is the **Zero-one loss**:

$$L_{0-1}(\eta) = \mathbf{1}_{\eta < 0} \quad \text{where } \eta = y\mathbf{x}^\top \mathbf{w}$$

- This corresponds to **Empirical Risk Minimization (ERM)**.
- Problem:** The 0-1 loss is **not convex** and **not continuous**. It is computationally hard to optimize.




# The 0-1 Loss and Convex Relaxations


- The “ideal” loss for classification is the **Zero-one loss**:

$$L_{0-1}(\eta) = \mathbf{1}_{\eta < 0} \quad \text{where } \eta = y\mathbf{x}^\top \mathbf{w}$$

- This corresponds to **Empirical Risk Minimization (ERM)**.
- Problem:** The 0-1 loss is **not convex** and **not continuous**. It is computationally hard to optimize.
- Solution:** We replace the 0-1 loss with a **convex surrogate** (an upper bound) that is easier to minimize.




# Comparison of Loss Functions




- **MSE:** Penalizes all points, even correctly classified ones. Symmetric.

# Comparison of Loss Functions



- **MSE:** Penalizes all points, even correctly classified ones. Symmetric.
- **Logistic:** Always incurs a cost, but it's asymmetric.

# Comparison of Loss Functions



- **MSE:** Penalizes all points, even correctly classified ones. Symmetric.
- **Logistic:** Always incurs a cost, but it's asymmetric.
- **Hinge:** No cost if a point is “sufficiently far” on the correct side (i.e.,  $y_n \mathbf{x}_n^\top \mathbf{w} \geq 1$ ). Creates a margin.

# Support Vector Machine Optimization Problem

SVMs solve the following optimization problem:

$$\min_{\mathbf{w}} \underbrace{\frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+}_{\text{Hinge Loss (Empirical Risk)}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Regularization}}$$

# Support Vector Machine Optimization Problem

SVMs solve the following optimization problem:

$$\min_{\mathbf{w}} \underbrace{\frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+}_{\text{Hinge Loss (Empirical Risk)}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Regularization}}$$

- The **Hinge Loss** pushes for correct classification, penalizing points in the margin ( $yz < 1$ ).
- The **Regularization** ( $\|\mathbf{w}\|^2$ ) pushes for a simpler model.

# Support Vector Machine Optimization Problem

SVMs solve the following optimization problem:

$$\min_{\mathbf{w}} \underbrace{\frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+}_{\text{Hinge Loss (Empirical Risk)}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Regularization}}$$

- The **Hinge Loss** pushes for correct classification, penalizing points in the margin ( $yz < 1$ ).
- The **Regularization** ( $\|\mathbf{w}\|^2$ ) pushes for a simpler model.
- **Geometric Meaning:** Minimizing  $\|\mathbf{w}\|^2$  is equivalent to **maximizing the margin**.

# Support Vector Machine Optimization Problem

SVMs solve the following optimization problem:

$$\min_{\mathbf{w}} \underbrace{\frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+}_{\text{Hinge Loss (Empirical Risk)}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Regularization}}$$

- The **Hinge Loss** pushes for correct classification, penalizing points in the margin ( $yz < 1$ ).
- The **Regularization** ( $\|\mathbf{w}\|^2$ ) pushes for a simpler model.
- **Geometric Meaning:** Minimizing  $\|\mathbf{w}\|^2$  is equivalent to **maximizing the margin**.
- $\lambda$  is a hyperparameter that controls the trade-off:
  - **Small**  $\lambda$ : “Harder” margin. Focus on maximizing margin. Low bias, high variance.
  - **Large**  $\lambda$ : “Softer” margin. Focus on minimizing classification errors. High bias, low variance.


# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:

$$H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}.$$


# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:  
 $H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}$ .
- The prediction is  $g(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$ .




# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:  
 $H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}$ .
- The prediction is  $g(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$ .
- The **geometric margin** for a point  $\mathbf{x}_0$  is its Euclidean distance to the hyperplane  $H$ .




# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:  
 $H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}$ .
- The prediction is  $g(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$ .
- The **geometric margin** for a point  $\mathbf{x}_0$  is its Euclidean distance to the hyperplane  $H$ .
- **Claim:** The distance from  $\mathbf{x}_0$  to  $H$  is  $\frac{|\mathbf{x}_0^\top \mathbf{w}|}{\|\mathbf{w}\|}$ .
  - (This is found by solving  $\min_{\mathbf{x}} \|\mathbf{x}_0 - \mathbf{x}\|$  s.t.  $\mathbf{x}^\top \mathbf{w} = 0$ ).




# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:  
 $H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}$ .
- The prediction is  $g(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$ .
- The **geometric margin** for a point  $\mathbf{x}_0$  is its Euclidean distance to the hyperplane  $H$ .
- **Claim:** The distance from  $\mathbf{x}_0$  to  $H$  is  $\frac{|\mathbf{x}_0^\top \mathbf{w}|}{\|\mathbf{w}\|}$ .
  - (This is found by solving  $\min_{\mathbf{x}} \|\mathbf{x}_0 - \mathbf{x}\|$  s.t.  $\mathbf{x}^\top \mathbf{w} = 0$ ).
- The **functional margin** for a point  $(\mathbf{x}_n, y_n)$  is  $y_n(\mathbf{x}_n^\top \mathbf{w})$ .
  - $g(\mathbf{x}_n) = y_n \iff$  functional margin is positive.




# Geometric Preliminaries: Margins

- A linear classifier is defined by a hyperplane:  
 $H = \{\mathbf{x} : \mathbf{x}^\top \mathbf{w} = 0\}$ .
- The prediction is  $g(\mathbf{x}) = \text{sign}(\mathbf{x}^\top \mathbf{w})$ .
- The **geometric margin** for a point  $\mathbf{x}_0$  is its Euclidean distance to the hyperplane  $H$ .
- **Claim:** The distance from  $\mathbf{x}_0$  to  $H$  is  $\frac{|\mathbf{x}_0^\top \mathbf{w}|}{\|\mathbf{w}\|}$ .
  - (This is found by solving  $\min_{\mathbf{x}} \|\mathbf{x}_0 - \mathbf{x}\|$  s.t.  $\mathbf{x}^\top \mathbf{w} = 0$ ).
- The **functional margin** for a point  $(\mathbf{x}_n, y_n)$  is  $y_n(\mathbf{x}_n^\top \mathbf{w})$ .
  - $g(\mathbf{x}_n) = y_n \iff$  functional margin is positive.
- We can rescale  $\mathbf{w}$  without changing the hyperplane, but it changes the functional margin. The geometric margin is invariant.




# Geometric Interpretation: The Margin




- The separating hyperplane is  $x^T w = 0$ .

# Geometric Interpretation: The Margin




- The separating hyperplane is  $x^T w = 0$ .
- The “margin” is the “street” defined by the two hyperplanes  $x^T w = 1$  and  $x^T w = -1$ .

# Geometric Interpretation: The Margin




- The separating hyperplane is  $x^\top w = 0$ .
- The “margin” is the “street” defined by the two hyperplanes  $x^\top w = 1$  and  $x^\top w = -1$ .
- The Hinge Loss penalizes any point  $n$  that is on the wrong side of its margin boundary (i.e.,  $y_n x_n^\top w < 1$ ).

# Geometric Interpretation: The Margin



- The separating hyperplane is  $x^\top \mathbf{w} = 0$ .
- The “margin” is the “street” defined by the two hyperplanes  $x^\top \mathbf{w} = 1$  and  $x^\top \mathbf{w} = -1$ .
- The Hinge Loss penalizes any point  $n$  that is on the wrong side of its margin boundary (i.e.,  $y_n x_n^\top \mathbf{w} < 1$ ).
- The width of this margin is  $\frac{2}{\|\mathbf{w}\|}$ . The margin does not only depend on the direction of the vector  $\mathbf{w}$  but also its norm.

# Geometric Interpretation: The Margin



- The separating hyperplane is  $x^\top w = 0$ .
- The “margin” is the “street” defined by the two hyperplanes  $x^\top w = 1$  and  $x^\top w = -1$ .
- The Hinge Loss penalizes any point  $n$  that is on the wrong side of its margin boundary (i.e.,  $y_n x_n^\top w < 1$ ).
- The width of this margin is  $\frac{2}{\|w\|}$ . The margin does not only depend on the direction of the vector  $w$  but also its norm.
- **Maximizing the margin**  $\left(\frac{2}{\|w\|}\right)$  is the same as **minimizing**  $\|w\|^2$ .

# The Hard-Margin Formulation

- Let's first assume the data is **linearly separable**.

# The Hard-Margin Formulation

- Let's first assume the data is **linearly separable**.
- The goal is to find the **max-margin separating hyperplane**.

**(I) Maximize Geometric Margin** Maximize the geometric margin  $\frac{M}{\|\mathbf{w}\|}$ , while ensuring all points are correctly classified:  $\max_{\mathbf{w}, \|\mathbf{w}\|=1} \min_n |\mathbf{x}_n^\top \mathbf{w}| \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq 0, \forall n.$

# The Hard-Margin Formulation

- Let's first assume the data is **linearly separable**.
- The goal is to find the **max-margin separating hyperplane**.
- This problem can be formulated in (at least) three equivalent ways:

**(I) Maximize Geometric Margin** Maximize the geometric margin  $\frac{M}{\|\mathbf{w}\|}$ , while ensuring all points are correctly classified:  $\max_{\mathbf{w}, \|\mathbf{w}\|=1} \min_n |\mathbf{x}_n^\top \mathbf{w}| \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq 0, \forall n.$

**(II) Maximize Functional Margin** Fix  $\|\mathbf{w}\| = 1$  and maximize the functional margin  $M$ :

$$\max_{M \in \mathbb{R}, \mathbf{w}, \|\mathbf{w}\|=1} M \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq M, \forall n$$

# The Hard-Margin Formulation

- Let's first assume the data is **linearly separable**.
- The goal is to find the **max-margin separating hyperplane**.
- This problem can be formulated in (at least) three equivalent ways:

**(I) Maximize Geometric Margin** Maximize the geometric margin  $\frac{M}{\|\mathbf{w}\|}$ , while ensuring all points are correctly classified:  $\max_{\mathbf{w}, \|\mathbf{w}\|=1} \min_n |\mathbf{x}_n^\top \mathbf{w}| \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq 0, \forall n.$


**(II) Maximize Functional Margin** Fix  $\|\mathbf{w}\| = 1$  and maximize the functional margin  $M$ :

$$\max_{M \in \mathbb{R}, \mathbf{w}, \|\mathbf{w}\|=1} M \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq M, \forall n$$

**(III) Minimize  $\|\mathbf{w}\|^2$  (Standard Form)** Rescale  $\mathbf{w}$  so the functional margin for the closest points is 1. Then minimize  $\|\mathbf{w}\|^2$  (which maximizes geometric margin  $1/\|\mathbf{w}\|$ ).

$$\min_{\mathbf{w}} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{s.t. } y_n(\mathbf{x}_n^\top \mathbf{w}) \geq 1, \forall n$$

# From Hard-Margin to Soft-Margin SVM




## Hard-Margin SVM

- Assumes data is **linearly separable**.

$$\min_w \frac{1}{2} \|w\|^2 \text{ s.t. } y_n(x_n^\top w) \geq 1, \forall n$$

- Fails if data is not separable.




## Soft-Margin SVM

- Allows for violations of the margin.
- Introduce positive “slack” variables  $\xi_n \geq 0$ .

$$\min_{w,\xi} \frac{\lambda}{2} \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \xi_n \text{ s.t. } y_n(x_n^\top w) \geq 1 - \xi_n, \forall n$$

$$\xi_n \geq 0, \forall n$$

# From Hard-Margin to Soft-Margin SVM




## Hard-Margin SVM

- Assumes data is **linearly separable**.

$$\min_w \frac{1}{2} \|w\|^2 \text{ s.t. } y_n(x_n^\top w) \geq 1, \forall n$$

- Fails if data is not separable.



## Soft-Margin SVM

- Allows for violations of the margin.
- Introduce positive “slack” variables  $\xi_n \geq 0$ .

$$\min_{w,\xi} \frac{\lambda}{2} \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \xi_n \text{ s.t. } y_n(x_n^\top w) \geq 1 - \xi_n, \forall n$$

$$\xi_n \geq 0, \forall n$$

The two formulations are equivalent. The optimal  $\xi_n$  is exactly the Hinge Loss:

$$\xi_n^* = \max(0, 1 - y_n x_n^\top w) = [1 - y_n x_n^\top w]_+$$

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? **Yes.** (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? **Yes.** (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).
- Is it differentiable? (The Hinge loss has a “kink” at  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? **Yes.** (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).
- Is it differentiable? **No.** (The Hinge loss has a “kink” at  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? **Yes.** (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).
- Is it differentiable? **No.** (The Hinge loss has a “kink” at  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).
- Because it is convex, we can optimize it using **subgradient** methods, like Stochastic Subgradient Descent (SGD).

# Optimization (Primal)

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N [1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- Is this function convex? **Yes.** (Sum of convex functions: Hinge loss is convex,  $\|\mathbf{w}\|^2$  is convex).
- Is it differentiable? **No.** (The Hinge loss has a “kink” at  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).
- Because it is convex, we can optimize it using **subgradient** methods, like Stochastic Subgradient Descent (SGD).
- But this can be slow. An alternative, and very important, approach is to use **duality**.

# Duality: The Big Picture

- We rewrite our minimization problem (**Primal**)  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$  into equivalent “saddle-point” problem.

# Duality: The Big Picture

- We rewrite our minimization problem (**Primal**)  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$  into equivalent “saddle-point” problem.
- We introduce an auxiliary function  $G(\mathbf{w}, \boldsymbol{\alpha})$  and dual variables  $\boldsymbol{\alpha}$ :

$$\mathcal{L}(\mathbf{w}) = \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

# Duality: The Big Picture

- We rewrite our minimization problem (**Primal**)  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$  into equivalent “saddle-point” problem.
- We introduce an auxiliary function  $G(\mathbf{w}, \boldsymbol{\alpha})$  and dual variables  $\boldsymbol{\alpha}$ :

$$\mathcal{L}(\mathbf{w}) = \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

- **Primal Problem:**

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

# Duality: The Big Picture

- We rewrite our minimization problem (**Primal**)  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$  into equivalent “saddle-point” problem.
- We introduce an auxiliary function  $G(\mathbf{w}, \boldsymbol{\alpha})$  and dual variables  $\boldsymbol{\alpha}$ :

$$\mathcal{L}(\mathbf{w}) = \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

- **Primal Problem:**

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha}} G(\mathbf{w}, \boldsymbol{\alpha})$$

- **Dual Problem:** (We swap the min and max)

$$\max_{\boldsymbol{\alpha}} \min_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha})$$

# Duality: The Big Picture

- We rewrite our minimization problem (**Primal**)  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w})$  into equivalent “saddle-point” problem.
- We introduce an auxiliary function  $G(\mathbf{w}, \alpha)$  and dual variables  $\alpha$ :

$$\mathcal{L}(\mathbf{w}) = \max_{\alpha} G(\mathbf{w}, \alpha)$$


- **Primal Problem:**

$$\min_{\mathbf{w}} \max_{\alpha} G(\mathbf{w}, \alpha)$$

- **Dual Problem:** (We swap the min and max)

$$\max_{\alpha} \min_{\mathbf{w}} G(\mathbf{w}, \alpha)$$

- For SVM, the domains are convex and  $G$  is convex in  $\mathbf{w}$  and concave in  $\alpha$ , so **strong duality** holds:



# Duality: Finding $G(\mathbf{w}, \alpha)$ for SVM

- We use a key trick for the Hinge loss:

$$[z]_+ = \max_{\alpha \in [0,1]} \alpha z$$

# Duality: Finding $G(\mathbf{w}, \alpha)$ for SVM

- We use a key trick for the Hinge loss:

$$[z]_+ = \max_{\alpha \in [0,1]} \alpha z$$

- Applying this to our Hinge loss for each data point:

$$[1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ = \max_{\alpha_n \in [0,1]} \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w})$$

# Duality: Finding $G(\mathbf{w}, \boldsymbol{\alpha})$ for SVM

- We use a key trick for the Hinge loss:

$$[z]_+ = \max_{\alpha \in [0,1]} \alpha z$$

- Applying this to our Hinge loss for each data point:

$$[1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ = \max_{\alpha_n \in [0,1]} \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w})$$

- We can now rewrite the *entire* SVM problem:

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha} \in [0,1]^N} \underbrace{\frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2}_{G(\mathbf{w}, \boldsymbol{\alpha})}$$

# Duality: Finding $G(\mathbf{w}, \alpha)$ for SVM

- We use a key trick for the Hinge loss:

$$[z]_+ = \max_{\alpha \in [0,1]} \alpha z$$

- Applying this to our Hinge loss for each data point:

$$[1 - y_n \mathbf{x}_n^\top \mathbf{w}]_+ = \max_{\alpha_n \in [0,1]} \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w})$$

- We can now rewrite the *entire* SVM problem:

$$\min_{\mathbf{w}} \max_{\boldsymbol{\alpha} \in [0,1]^N} \underbrace{\frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}) + \frac{\lambda}{2} \|\mathbf{w}\|^2}_{G(\mathbf{w}, \boldsymbol{\alpha})}$$

- This  $G(\mathbf{w}, \boldsymbol{\alpha})$  is convex in  $\mathbf{w}$  and concave (linear) in  $\boldsymbol{\alpha}$ .

# Duality: Solving the Dual (Step 1)

We want to solve:  $\max_{\alpha \in [0,1]^N} \min_w G(\mathbf{w}, \alpha)$

- **Step 1: Minimize w.r.t.  $\mathbf{w}$  (for a fixed  $\alpha$ )**

# Duality: Solving the Dual (Step 1)

We want to solve:  $\max_{\alpha \in [0,1]^N} \min_w G(\mathbf{w}, \boldsymbol{\alpha})$

- **Step 1: Minimize w.r.t.  $\mathbf{w}$  (for a fixed  $\boldsymbol{\alpha}$ )**
- We find the  $\mathbf{w}$  that minimizes  $G$  by taking the gradient and setting it to zero:

$$\begin{aligned}\nabla_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha}) &= \nabla_{\mathbf{w}} \left( \frac{1}{N} \sum_{n=1}^N \alpha_n - \frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n^\top \mathbf{w} + \frac{\lambda}{2} \|\mathbf{w}\|^2 \right) \\ &= -\frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w} \stackrel{\text{set}}{=} 0\end{aligned}$$

# Duality: Solving the Dual (Step 1)

We want to solve:  $\max_{\alpha \in [0,1]^N} \min_w G(\mathbf{w}, \boldsymbol{\alpha})$

- **Step 1: Minimize w.r.t.  $\mathbf{w}$  (for a fixed  $\boldsymbol{\alpha}$ )**
- We find the  $\mathbf{w}$  that minimizes  $G$  by taking the gradient and setting it to zero:

$$\begin{aligned}\nabla_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha}) &= \nabla_{\mathbf{w}} \left( \frac{1}{N} \sum_{n=1}^N \alpha_n - \frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n^\top \mathbf{w} + \frac{\lambda}{2} \|\mathbf{w}\|^2 \right) \\ &= -\frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w} \stackrel{\text{set}}{=} 0\end{aligned}$$

- This gives us the optimal  $\mathbf{w}$  as a function of  $\boldsymbol{\alpha}$ :

$$\mathbf{w}(\boldsymbol{\alpha}) = \frac{1}{\lambda N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = \frac{1}{\lambda N} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}$$

# Duality: Solving the Dual (Step 1)

We want to solve:  $\max_{\alpha \in [0,1]^N} \min_w G(\mathbf{w}, \boldsymbol{\alpha})$

- **Step 1: Minimize w.r.t.  $\mathbf{w}$  (for a fixed  $\boldsymbol{\alpha}$ )**
- We find the  $\mathbf{w}$  that minimizes  $G$  by taking the gradient and setting it to zero:

$$\begin{aligned}\nabla_{\mathbf{w}} G(\mathbf{w}, \boldsymbol{\alpha}) &= \nabla_{\mathbf{w}} \left( \frac{1}{N} \sum_{n=1}^N \alpha_n - \frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n^\top \mathbf{w} + \frac{\lambda}{2} \|\mathbf{w}\|^2 \right) \\ &= -\frac{1}{N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n + \lambda \mathbf{w} \stackrel{\text{set}}{=} 0\end{aligned}$$

- This gives us the optimal  $\mathbf{w}$  as a function of  $\boldsymbol{\alpha}$ :

$$\mathbf{w}(\boldsymbol{\alpha}) = \frac{1}{\lambda N} \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n = \frac{1}{\lambda N} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}$$

- **Crucial Insight:** The optimal weight vector  $\mathbf{w}$  is a linear combination of the data points  $\mathbf{x}_n$ .

## Duality: Solving the Dual (Step 2)

- **Step 2: Plug  $w(\alpha)$  back into  $G(w, \alpha)$ .**

## Duality: Solving the Dual (Step 2)

- **Step 2: Plug  $\mathbf{w}(\boldsymbol{\alpha})$  back into  $G(\mathbf{w}, \boldsymbol{\alpha})$ .**
- We substitute  $\mathbf{w}(\boldsymbol{\alpha}) = \frac{1}{\lambda N} \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j$  into  $G$ :

$$\begin{aligned}
 G(\mathbf{w}(\boldsymbol{\alpha}), \boldsymbol{\alpha}) &= \frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}(\boldsymbol{\alpha})) + \frac{\lambda}{2} \|\mathbf{w}(\boldsymbol{\alpha})\|^2 \\
 &= \frac{1}{N} \sum_n \alpha_n - \frac{1}{N} \sum_n \alpha_n y_n \mathbf{x}_n^\top \left( \frac{1}{\lambda N} \sum_j \alpha_j y_j \mathbf{x}_j \right) + \frac{\lambda}{2} \left\| \frac{1}{\lambda N} \sum_n \alpha_n y_n \mathbf{x}_n \right\|^2
 \end{aligned}$$

## Duality: Solving the Dual (Step 2)

- **Step 2: Plug  $\mathbf{w}(\boldsymbol{\alpha})$  back into  $G(\mathbf{w}, \boldsymbol{\alpha})$ .**
- We substitute  $\mathbf{w}(\boldsymbol{\alpha}) = \frac{1}{\lambda N} \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j$  into  $G$ :

$$\begin{aligned} G(\mathbf{w}(\boldsymbol{\alpha}), \boldsymbol{\alpha}) &= \frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}(\boldsymbol{\alpha})) + \frac{\lambda}{2} \|\mathbf{w}(\boldsymbol{\alpha})\|^2 \\ &= \frac{1}{N} \sum_n \alpha_n - \frac{1}{N} \sum_n \alpha_n y_n \mathbf{x}_n^\top \left( \frac{1}{\lambda N} \sum_j \alpha_j y_j \mathbf{x}_j \right) + \frac{\lambda}{2} \left\| \frac{1}{\lambda N} \sum_n \alpha_n y_n \mathbf{x}_n \right\|^2 \end{aligned}$$

- Let's simplify the terms. Note  $\sum_n \sum_j \alpha_n \alpha_j y_n y_j (\mathbf{x}_n^\top \mathbf{x}_j) = \boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{K} \mathbf{Y} \boldsymbol{\alpha}$  where  $\mathbf{K} = \mathbf{X} \mathbf{X}^\top$ .

$$\begin{aligned} &= \frac{1}{N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{\lambda N^2} (\boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{X} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}) + \frac{\lambda}{2\lambda^2 N^2} (\boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{X} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}) \\ &= \frac{1}{N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{2\lambda N^2} \boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{K} \mathbf{Y} \boldsymbol{\alpha} \end{aligned}$$

## Duality: Solving the Dual (Step 2)

- **Step 2: Plug  $\mathbf{w}(\boldsymbol{\alpha})$  back into  $G(\mathbf{w}, \boldsymbol{\alpha})$ .**
- We substitute  $\mathbf{w}(\boldsymbol{\alpha}) = \frac{1}{\lambda N} \sum_{j=1}^N \alpha_j y_j \mathbf{x}_j$  into  $G$ :

$$\begin{aligned} G(\mathbf{w}(\boldsymbol{\alpha}), \boldsymbol{\alpha}) &= \frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n \mathbf{x}_n^\top \mathbf{w}(\boldsymbol{\alpha})) + \frac{\lambda}{2} \|\mathbf{w}(\boldsymbol{\alpha})\|^2 \\ &= \frac{1}{N} \sum_n \alpha_n - \frac{1}{N} \sum_n \alpha_n y_n \mathbf{x}_n^\top \left( \frac{1}{\lambda N} \sum_j \alpha_j y_j \mathbf{x}_j \right) + \frac{\lambda}{2} \left\| \frac{1}{\lambda N} \sum_n \alpha_n y_n \mathbf{x}_n \right\|^2 \end{aligned}$$

- Let's simplify the terms. Note  $\sum_n \sum_j \alpha_n \alpha_j y_n y_j (\mathbf{x}_n^\top \mathbf{x}_j) = \boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{K} \mathbf{Y} \boldsymbol{\alpha}$  where  $\mathbf{K} = \mathbf{X} \mathbf{X}^\top$ .

$$\begin{aligned} &= \frac{1}{N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{\lambda N^2} (\boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{X} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}) + \frac{\lambda}{2\lambda^2 N^2} (\boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{X} \mathbf{X}^\top \mathbf{Y} \boldsymbol{\alpha}) \\ &= \frac{1}{N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{2\lambda N^2} \boldsymbol{\alpha}^\top \mathbf{Y} \mathbf{K} \mathbf{Y} \boldsymbol{\alpha} \end{aligned}$$

- This gives the final **Dual Problem**:

$$\max_{\boldsymbol{\alpha} \in [0,1]^N} \frac{1}{N} \boldsymbol{\alpha}^\top \mathbf{1} - \frac{1}{2\lambda N^2} \boldsymbol{\alpha}^\top \mathbf{Q} \boldsymbol{\alpha}$$

where  $\mathbf{Q} = \mathbf{Y} \mathbf{K} \mathbf{Y} = \text{diag}(\mathbf{y}) \mathbf{X} \mathbf{X}^\top \text{diag}(\mathbf{y})$

# Duality: Why Bother?

The dual problem is often better:

$$\max_{\alpha \in [0,1]^N} \alpha^\top \mathbf{1} - \frac{1}{2\lambda N^2} \alpha^\top \mathbf{Q} \alpha \quad \text{where } \mathbf{Q} = \mathbf{Y}(\mathbf{X}\mathbf{X}^\top) \mathbf{Y}$$

- 1. It's a **constrained quadratic problem**. We can solve it efficiently with [Coordinate Ascent](#).

# Duality: Why Bother?

The dual problem is often better:

$$\max_{\alpha \in [0,1]^N} \alpha^\top \mathbf{1} - \frac{1}{2\lambda N^2} \alpha^\top \mathbf{Q} \alpha \quad \text{where } \mathbf{Q} = \mathbf{Y}(\mathbf{X}\mathbf{X}^\top)\mathbf{Y}$$

- 1. It's a constrained quadratic problem. We can solve it efficiently with [Coordinate Ascent](#).
- 2. The data only appears as  $\mathbf{K} = \mathbf{X}\mathbf{X}^\top$ .
  - This  $N \times N$  matrix is called the [Kernel](#).
  - The problem complexity now depends on  $N$  (samples), not  $D$  (features).
  - **This is a huge win when  $D \gg N$  (e.g., text, genomics).**
  - This property is what enables the [Kernel Trick](#).

# Duality: Why Bother?

The dual problem is often better:

$$\max_{\alpha \in [0,1]^N} \alpha^\top \mathbf{1} - \frac{1}{2\lambda N^2} \alpha^\top \mathbf{Q} \alpha \quad \text{where } \mathbf{Q} = \mathbf{Y}(\mathbf{X}\mathbf{X}^\top) \mathbf{Y}$$


- 1. **It's a constrained quadratic problem.** We can solve it efficiently with [Coordinate Ascent](#).
- 2. **The data only appears as  $\mathbf{K} = \mathbf{X}\mathbf{X}^\top$ .**
  - This  $N \times N$  matrix is called the [Kernel](#).
  - The problem complexity now depends on  $N$  (samples), not  $D$  (features).
  - **This is a huge win when  $D \gg N$  (e.g., text, genomics).**
  - This property is what enables the [Kernel Trick](#).
- 3. **The solution  $\alpha$  is sparse.**
  - Most  $\alpha_n$  will be 0.
  - The  $\mathbf{w}$  vector ( $\mathbf{w} = \frac{1}{\lambda N} \sum \alpha_n y_n \mathbf{x}_n$ ) only depends on the points  $\mathbf{x}_n$  where  $\alpha_n > 0$ .
  - These points are the [Support Vectors](#).

# Interpretation: Support Vectors

From the KKT conditions (the optimality conditions for the dual):

- **Case 1: Point is outside margin.** ( $\alpha_n = 0$ )

- $y_n \mathbf{x}_n^\top \mathbf{w} > 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} < 0$ ).
- The point is correctly classified with room to spare.
- The Hinge Loss is 0.
- This point has **no influence** on the decision boundary.



# Interpretation: Support Vectors


From the KKT conditions (the optimality conditions for the dual):

- **Case 1: Point is outside margin.** ( $\alpha_n = 0$ )

- $y_n \mathbf{x}_n^\top \mathbf{w} > 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} < 0$ ).
- The point is correctly classified with room to spare.
- The Hinge Loss is 0.
- This point has **no influence** on the decision boundary.

- **Case 2: Point is on the margin.** ( $\alpha_n \in (0, 1]$ )

- $y_n \mathbf{x}_n^\top \mathbf{w} = 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).
- This is an **Essential Support Vector**. It lies exactly on the margin and holds the “street” in place.



# Interpretation: Support Vectors

From the KKT conditions (the optimality conditions for the dual):

- **Case 1: Point is outside margin.** ( $\alpha_n = 0$ )


- $y_n \mathbf{x}_n^\top \mathbf{w} > 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} < 0$ ).
- The point is correctly classified with room to spare.
- The Hinge Loss is 0.
- This point has **no influence** on the decision boundary.

- **Case 2: Point is on the margin.** ( $\alpha_n \in (0, 1]$ )

- $y_n \mathbf{x}_n^\top \mathbf{w} = 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} = 0$ ).
- This is an **Essential Support Vector**. It lies exactly on the margin and holds the “street” in place.

- **Case 3: Point is inside margin or misclassified.** ( $\alpha_n = 1$ )

- $y_n \mathbf{x}_n^\top \mathbf{w} < 1$  (or  $1 - y_n \mathbf{x}_n^\top \mathbf{w} > 0$ ).
- This is a **Bound Support Vector**. It is either inside the margin or on the wrong side of the hyperplane.



## A Deeper Look: KKT Conditions

- KKT (Karush-Kuhn-Tucker) conditions are the formal rules that an optimal solution to a constrained problem (like SVM) must satisfy.

## A Deeper Look: KKT Conditions

- KKT (Karush-Kuhn-Tucker) conditions are the formal rules that an optimal solution to a constrained problem (like SVM) must satisfy.
- They provide the direct link between the primal problem (finding  $w$ ) and the dual problem (finding  $\alpha$ ).

## A Deeper Look: KKT Conditions

- KKT (Karush-Kuhn-Tucker) conditions are the formal rules that an optimal solution to a constrained problem (like SVM) must satisfy.
- They provide the direct link between the primal problem (finding  $w$ ) and the dual problem (finding  $\alpha$ ).
- The most important rule for interpretation is **Complementary Slackness**.

## A Deeper Look: KKT Conditions

- KKT (Karush-Kuhn-Tucker) conditions are the formal rules that an optimal solution to a constrained problem (like SVM) must satisfy.
- They provide the direct link between the primal problem (finding  $w$ ) and the dual problem (finding  $\alpha$ ).
- The most important rule for interpretation is **Complementary Slackness**.
- This rule connects the dual variable  $\alpha_n$  to the primal constraint for point  $n$ . In simple terms, it means:

*Either the constraint is active (the point is on/inside the margin),  
or the dual variable  $\alpha_n$  for that constraint must be zero.*

# A Deeper Look: KKT Conditions

- KKT (Karush-Kuhn-Tucker) conditions are the formal rules that an optimal solution to a constrained problem (like SVM) must satisfy.
- They provide the direct link between the primal problem (finding  $\mathbf{w}$ ) and the dual problem (finding  $\alpha$ ).
- The most important rule for interpretation is **Complementary Slackness**.
- This rule connects the dual variable  $\alpha_n$  to the primal constraint for point  $n$ . In simple terms, it means:

*Either the constraint is active (the point is on/inside the margin),  
or the dual variable  $\alpha_n$  for that constraint must be zero.*

- This is the formal reason for the three cases:

- **Case 1: Point is outside margin.**
  - $y_n \mathbf{x}_n^\top \mathbf{w} > 1$ . The constraint is *inactive* (slack).
  - KKT conditions force  $\alpha_n = 0$ .
- **Case 2 & 3: Point is on or inside margin.**
  - $y_n \mathbf{x}_n^\top \mathbf{w} \leq 1$ . The constraint is *active* (tight).
  - KKT conditions allow  $\alpha_n > 0$ .
  - These are the **Support Vectors**.

# Optimization: Coordinate Ascent

**Goal:** Maximize the dual  $g(\alpha) = \alpha^\top \mathbf{1} - \frac{1}{2\lambda N^2} \alpha^\top Q \alpha$  subject to  $\alpha_n \in [0, 1]$ .

**Idea:** Update one coordinate at a time, keeping others fixed.

- 1: initialize  $\alpha^{(0)} \in \mathbb{R}^N$  (e.g., all zeros)
- 2: **for**  $t = 0$  to maxIter **do**
- 3:     sample a coordinate  $n$  randomly from  $1 \dots N$ .
- 4:     optimize  $g$  w.r.t. that single coordinate  $\alpha_n$  (which is just a 1D quadratic problem).
- 5:      $u^* \leftarrow \arg \max_{u \in \mathbb{R}} g(\alpha_1^{(t)}, \dots, \alpha_{n-1}^{(t)}, u, \alpha_{n+1}^{(t)}, \dots, \alpha_N^{(t)})$
- 6:     clip the result to respect the “box” constraint:
- 7:          $u^* \leftarrow \min(1, \max(0, u^*))$
- 8:     update  $\alpha_n^{(t+1)} \leftarrow u^*$

# Issues with SVM

- There is no obvious probabilistic interpretation (unlike Logistic Regression).
  - The output is a “score”  $\mathbf{x}^\top \mathbf{w}$ , not a probability  $p(y = 1|\mathbf{x})$ .
  - (Though this can be mitigated with “Platt Scaling”).

# Issues with SVM

- There is no obvious probabilistic interpretation (unlike Logistic Regression).
  - The output is a “score”  $\mathbf{x}^\top \mathbf{w}$ , not a probability  $p(y = 1|\mathbf{x})$ .
  - (Though this can be mitigated with “Platt Scaling”).
- Extension to multi-class classification is non-trivial. Common approaches:
  - **One-vs-Rest (OvR):** Train  $K$  binary classifiers (class  $k$  vs. all others).
  - **One-vs-One (OvO):** Train  $\frac{K(K-1)}{2}$  binary classifiers (class  $i$  vs. class  $j$ ).

# Table of Contents

- ① Support Vector Machines (SVM)
- ② Kernel Methods and the Kernel Trick

# Motivation: Why Kernelize?

- We just saw that the SVM dual depends only on inner products  $\mathbf{K} = \mathbf{XX}^\top$ .

# Motivation: Why Kernelize?

- We just saw that the SVM dual depends only on inner products  $\mathbf{K} = \mathbf{XX}^\top$ .
- This  $N \times N$  matrix is called a **kernel matrix**.

# Motivation: Why Kernelize?

- We just saw that the SVM dual depends only on inner products  $\mathbf{K} = \mathbf{XX}^\top$ .
- This  $N \times N$  matrix is called a [kernel matrix](#).
- Let's see if this property holds for other algorithms, like Ridge Regression.

# Motivation: Why Kernelize?

- We just saw that the SVM dual depends only on inner products  $\mathbf{K} = \mathbf{XX}^\top$ .
- This  $N \times N$  matrix is called a **kernel matrix**.
- Let's see if this property holds for other algorithms, like Ridge Regression.
- This will lead us to a powerful general idea: the **Kernel Trick**.

# Motivation: Why Kernelize?

- We just saw that the SVM dual depends only on inner products  $\mathbf{K} = \mathbf{XX}^\top$ .
- This  $N \times N$  matrix is called a **kernel matrix**.
- Let's see if this property holds for other algorithms, like Ridge Regression.
- This will lead us to a powerful general idea: the **Kernel Trick**.
- This trick allows us to use highly complex, non-linear, and even infinite-dimensional feature maps *without* paying the computational cost.

# Alternative Formulation of Ridge Regression

Recall the **primal** problem for Ridge Regression:

$$\min_{\mathbf{w}} \quad \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

# Alternative Formulation of Ridge Regression

Recall the **primal** problem for Ridge Regression:

$$\min_{\mathbf{w}} \quad \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

The solution  $\mathbf{w}^* \in \mathbb{R}^D$  is:

$$\mathbf{w}^* = \frac{1}{N} \left( \frac{1}{N} \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D \right)^{-1} \mathbf{X}^\top \mathbf{y}$$

# Alternative Formulation of Ridge Regression

Recall the **primal** problem for Ridge Regression:

$$\min_{\mathbf{w}} \quad \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

The solution  $\mathbf{w}^* \in \mathbb{R}^D$  is:

$$\mathbf{w}^* = \frac{1}{N} \left( \frac{1}{N} \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D \right)^{-1} \mathbf{X}^\top \mathbf{y}$$

- The main computation is inverting a  $D \times D$  matrix.

# Alternative Formulation of Ridge Regression

Recall the **primal** problem for Ridge Regression:

$$\min_{\mathbf{w}} \quad \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

The solution  $\mathbf{w}^* \in \mathbb{R}^D$  is:

$$\mathbf{w}^* = \frac{1}{N} \left( \frac{1}{N} \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D \right)^{-1} \mathbf{X}^\top \mathbf{y}$$

- The main computation is inverting a  $D \times D$  matrix.
- This costs  $O(D^3 + ND^2)$ . This is great if  $D \ll N$ .

# Alternative Formulation of Ridge Regression

Recall the **primal** problem for Ridge Regression:

$$\min_{\mathbf{w}} \quad \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

The solution  $\mathbf{w}^* \in \mathbb{R}^D$  is:

$$\mathbf{w}^* = \frac{1}{N} \left( \frac{1}{N} \mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I}_D \right)^{-1} \mathbf{X}^\top \mathbf{y}$$

- The main computation is inverting a  $D \times D$  matrix.
- This costs  $O(D^3 + ND^2)$ . This is great if  $D \ll N$ .
- But what if  $D \gg N$ ? (e.g., 100 samples, 1,000,000 features)

## Alternative Formulation (Dual)

We claim the solution can be written alternatively as:

$$\mathbf{w}^* = \frac{1}{N} \mathbf{X}^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \quad (5)$$

## Alternative Formulation (Dual)

We claim the solution can be written alternatively as:

$$\mathbf{w}^* = \frac{1}{N} \mathbf{X}^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \quad (5)$$

- The main computation is now inverting an  $N \times N$  matrix.

## Alternative Formulation (Dual)

We claim the solution can be written alternatively as:

$$\mathbf{w}^* = \frac{1}{N} \mathbf{X}^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \quad (5)$$

- The main computation is now inverting an  $N \times N$  matrix.
- This costs  $O(N^3 + DN^2)$ .

## Alternative Formulation (Dual)

We claim the solution can be written alternatively as:

$$\mathbf{w}^* = \frac{1}{N} \mathbf{X}^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \quad (5)$$

- The main computation is now inverting an  $N \times N$  matrix.
- This costs  $O(N^3 + DN^2)$ .
- This is much better if  $N \ll D$ !

# Alternative Formulation (Dual)

We claim the solution can be written alternatively as:

$$\mathbf{w}^* = \frac{1}{N} \mathbf{X}^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N)^{-1} \mathbf{y} \quad (5)$$

- The main computation is now inverting an  $N \times N$  matrix.
- This costs  $O(N^3 + DN^2)$ .
- This is much better if  $N \ll D$ !

This relies on the “push-through” matrix identity:

$$(\mathbf{P}\mathbf{Q} + \mathbf{I}_N)^{-1}\mathbf{P} = \mathbf{P}(\mathbf{Q}\mathbf{P} + \mathbf{I}_M)^{-1}$$

(Here,  $\mathbf{P} = \mathbf{X}^\top$  and  $\mathbf{Q} = \frac{1}{\lambda N} \mathbf{X}$ )

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).
- Let  $\mathbf{P} \in \mathbb{R}^{m \times n}$  and  $\mathbf{Q} \in \mathbb{R}^{n \times m}$ .

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).
- Let  $\mathbf{P} \in \mathbb{R}^{m \times n}$  and  $\mathbf{Q} \in \mathbb{R}^{n \times m}$ .
- We have:  $\mathbf{P}(\mathbf{QP} + \mathbf{I}_n) = \mathbf{PQP} + \mathbf{P} = (\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}$ .

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).
- Let  $\mathbf{P} \in \mathbb{R}^{m \times n}$  and  $\mathbf{Q} \in \mathbb{R}^{n \times m}$ .
- We have:  $\mathbf{P}(\mathbf{QP} + \mathbf{I}_n) = \mathbf{PQP} + \mathbf{P} = (\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}$ .
- Assuming the inverses exist, we can multiply by  $(\mathbf{PQ} + \mathbf{I}_m)^{-1}$  on the left and  $(\mathbf{QP} + \mathbf{I}_n)^{-1}$  on the right:

$$\begin{aligned} (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)(\mathbf{QP} + \mathbf{I}_n)^{-1} &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}(\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \\ &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P} = \mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \end{aligned}$$

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).
- Let  $\mathbf{P} \in \mathbb{R}^{m \times n}$  and  $\mathbf{Q} \in \mathbb{R}^{n \times m}$ .
- We have:  $\mathbf{P}(\mathbf{QP} + \mathbf{I}_n) = \mathbf{PQP} + \mathbf{P} = (\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}$ .
- Assuming the inverses exist, we can multiply by  $(\mathbf{PQ} + \mathbf{I}_m)^{-1}$  on the left and  $(\mathbf{QP} + \mathbf{I}_n)^{-1}$  on the right:

$$\begin{aligned} (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)(\mathbf{QP} + \mathbf{I}_n)^{-1} &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}(\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \\ &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P} = \mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \end{aligned}$$

- We get our result by setting  $\mathbf{P} = \mathbf{X}^\top$  and  $\mathbf{Q} = \frac{1}{\lambda N} \mathbf{X}$ .

# Proof of Alternative Formulation

- This relies on a common matrix identity (sometimes called the “push-through” identity).
- Let  $\mathbf{P} \in \mathbb{R}^{m \times n}$  and  $\mathbf{Q} \in \mathbb{R}^{n \times m}$ .
- We have:  $\mathbf{P}(\mathbf{QP} + \mathbf{I}_n) = \mathbf{PQP} + \mathbf{P} = (\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}$ .
- Assuming the inverses exist, we can multiply by  $(\mathbf{PQ} + \mathbf{I}_m)^{-1}$  on the left and  $(\mathbf{QP} + \mathbf{I}_n)^{-1}$  on the right:

$$\begin{aligned} (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)(\mathbf{QP} + \mathbf{I}_n)^{-1} &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}(\mathbf{PQ} + \mathbf{I}_m)\mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \\ &= (\mathbf{PQ} + \mathbf{I}_m)^{-1}\mathbf{P} = \mathbf{P}(\mathbf{QP} + \mathbf{I}_n)^{-1} \end{aligned}$$

- We get our result by setting  $\mathbf{P} = \mathbf{X}^\top$  and  $\mathbf{Q} = \frac{1}{\lambda N} \mathbf{X}$ .
- This switches a  $D \times D$  inverse for an  $N \times N$  inverse.

# The Representer Theorem

The alternative form  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$  is not a coincidence.

Define  $\boldsymbol{\alpha}^* = \frac{1}{N}(\frac{1}{N}\mathbf{X}\mathbf{X}^\top + \lambda\mathbf{I}_N)^{-1}\mathbf{y}$ . Then  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$ .

# The Representer Theorem

The alternative form  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$  is not a coincidence.

Define  $\boldsymbol{\alpha}^* = \frac{1}{N}(\frac{1}{N}\mathbf{X}\mathbf{X}^\top + \lambda\mathbf{I}_N)^{-1}\mathbf{y}$ . Then  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$ .

This shows  $\mathbf{w}^*$  lies in the space spanned by the feature vectors (the rows of  $\mathbf{X}$ ).

# The Representer Theorem

The alternative form  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$  is not a coincidence.

Define  $\boldsymbol{\alpha}^* = \frac{1}{N}(\frac{1}{N}\mathbf{X}\mathbf{X}^\top + \lambda\mathbf{I}_N)^{-1}\mathbf{y}$ . Then  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$ .

This shows  $\mathbf{w}^*$  lies in the space spanned by the feature vectors (the rows of  $\mathbf{X}$ ).

**The Representer Theorem (General)** For any loss function  $\mathcal{L}_n$ , the solution  $\mathbf{w}^*$  to:

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{x}_n^\top \mathbf{w}, y_n) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

can be written as a linear combination of the data points:

$$\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^* = \sum_{n=1}^N \alpha_n^* \mathbf{x}_n$$

# The Representer Theorem

The alternative form  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$  is not a coincidence.

Define  $\boldsymbol{\alpha}^* = \frac{1}{N}(\frac{1}{N}\mathbf{X}\mathbf{X}^\top + \lambda\mathbf{I}_N)^{-1}\mathbf{y}$ . Then  $\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^*$ .

This shows  $\mathbf{w}^*$  lies in the space spanned by the feature vectors (the rows of  $\mathbf{X}$ ).

**The Representer Theorem (General)** For any loss function  $\mathcal{L}_n$ , the solution  $\mathbf{w}^*$  to:

$$\min_{\mathbf{w}} \frac{1}{N} \sum_{n=1}^N \mathcal{L}_n(\mathbf{x}_n^\top \mathbf{w}, y_n) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

can be written as a linear combination of the data points:

$$\mathbf{w}^* = \mathbf{X}^\top \boldsymbol{\alpha}^* = \sum_{n=1}^N \alpha_n^* \mathbf{x}_n$$

**Proof (Intuition):** Decompose  $\mathbf{w} = \mathbf{w}_{\parallel} + \mathbf{w}_{\perp}$ , where  $\mathbf{w}_{\parallel}$  is in the span of  $\mathbf{X}$  and  $\mathbf{w}_{\perp}$  is orthogonal to it.

- The loss term  $\mathcal{L}_n$  only depends on  $\mathbf{x}_n^\top \mathbf{w} = \mathbf{x}_n^\top \mathbf{w}_{\parallel}$ .
- The regularization term  $\|\mathbf{w}\|^2 = \|\mathbf{w}_{\parallel}\|^2 + \|\mathbf{w}_{\perp}\|^2$ .
- To minimize the cost, we must set  $\mathbf{w}_{\perp} = 0$ . Thus,  $\mathbf{w}^* = \mathbf{w}_{\parallel}^* \in \text{span}(\mathbf{X})$ .

# Representer Theorem (Formal Proof)

- Let  $w^*$  be any optimal solution.

# Representer Theorem (Formal Proof)

- Let  $\mathbf{w}^*$  be any optimal solution.
- We can always decompose  $\mathbf{w}^*$  into two orthogonal components:

$$\mathbf{w}^* = \underbrace{\sum_{n=1}^N \alpha_n \mathbf{x}_n}_{\mathbf{w}_{\parallel} \in \text{span}(\mathbf{X})} + \underbrace{\mathbf{u}}_{\mathbf{w}_{\perp} \perp \text{span}(\mathbf{X})}$$

where  $\mathbf{u}^\top \mathbf{x}_n = 0$  for all  $n$ .

# Representer Theorem (Formal Proof)

- Let  $\mathbf{w}^*$  be any optimal solution.
- We can always decompose  $\mathbf{w}^*$  into two orthogonal components:

$$\mathbf{w}^* = \underbrace{\sum_{n=1}^N \alpha_n \mathbf{x}_n}_{\mathbf{w}_{\parallel} \in \text{span}(\mathbf{X})} + \underbrace{\mathbf{u}}_{\mathbf{w}_{\perp} \perp \text{span}(\mathbf{X})}$$

where  $\mathbf{u}^\top \mathbf{x}_n = 0$  for all  $n$ .

- Part 1: The Loss Term.** The loss only depends on  $\mathbf{w}_{\parallel}$ :

$$\mathbf{x}_n^\top \mathbf{w}^* = \mathbf{x}_n^\top (\mathbf{w}_{\parallel} + \mathbf{u}) = \mathbf{x}_n^\top \mathbf{w}_{\parallel} + \mathbf{x}_n^\top \mathbf{u} = \mathbf{x}_n^\top \mathbf{w}_{\parallel}$$

So,  $l(\mathbf{x}_n^\top \mathbf{w}^*, y_n) = l(\mathbf{x}_n^\top \mathbf{w}_{\parallel}, y_n)$ .

# Representer Theorem (Formal Proof)

- Let  $\mathbf{w}^*$  be any optimal solution.
- We can always decompose  $\mathbf{w}^*$  into two orthogonal components:

$$\mathbf{w}^* = \underbrace{\sum_{n=1}^N \alpha_n \mathbf{x}_n}_{\mathbf{w}_{\parallel} \in \text{span}(\mathbf{X})} + \underbrace{\mathbf{u}}_{\mathbf{w}_{\perp} \perp \text{span}(\mathbf{X})}$$

where  $\mathbf{u}^\top \mathbf{x}_n = 0$  for all  $n$ .

- Part 1: The Loss Term.** The loss only depends on  $\mathbf{w}_{\parallel}$ :

$$\mathbf{x}_n^\top \mathbf{w}^* = \mathbf{x}_n^\top (\mathbf{w}_{\parallel} + \mathbf{u}) = \mathbf{x}_n^\top \mathbf{w}_{\parallel} + \mathbf{x}_n^\top \mathbf{u} = \mathbf{x}_n^\top \mathbf{w}_{\parallel}$$

So,  $l(\mathbf{x}_n^\top \mathbf{w}^*, y_n) = l(\mathbf{x}_n^\top \mathbf{w}_{\parallel}, y_n)$ .

- Part 2: The Regularization Term.** By Pythagorean theorem (since  $\mathbf{w}_{\parallel} \perp \mathbf{u}$ ):

$$\|\mathbf{w}^*\|^2 = \|\mathbf{w}_{\parallel} + \mathbf{u}\|^2 = \|\mathbf{w}_{\parallel}\|^2 + \|\mathbf{u}\|^2$$

This means  $\|\mathbf{w}_{\parallel}\|^2 \leq \|\mathbf{w}^*\|^2$ .

# Representer Theorem (Formal Proof)

- Let  $\mathbf{w}^*$  be any optimal solution.
- We can always decompose  $\mathbf{w}^*$  into two orthogonal components:

$$\mathbf{w}^* = \underbrace{\sum_{n=1}^N \alpha_n \mathbf{x}_n}_{\mathbf{w}_{\parallel} \in \text{span}(\mathbf{X})} + \underbrace{\mathbf{u}}_{\mathbf{w}_{\perp} \perp \text{span}(\mathbf{X})}$$

where  $\mathbf{u}^\top \mathbf{x}_n = 0$  for all  $n$ .

- Part 1: The Loss Term.** The loss only depends on  $\mathbf{w}_{\parallel}$ :

$$\mathbf{x}_n^\top \mathbf{w}^* = \mathbf{x}_n^\top (\mathbf{w}_{\parallel} + \mathbf{u}) = \mathbf{x}_n^\top \mathbf{w}_{\parallel} + \mathbf{x}_n^\top \mathbf{u} = \mathbf{x}_n^\top \mathbf{w}_{\parallel}$$

So,  $l(\mathbf{x}_n^\top \mathbf{w}^*, y_n) = l(\mathbf{x}_n^\top \mathbf{w}_{\parallel}, y_n)$ .

- Part 2: The Regularization Term.** By Pythagorean theorem (since  $\mathbf{w}_{\parallel} \perp \mathbf{u}$ ):

$$\|\mathbf{w}^*\|^2 = \|\mathbf{w}_{\parallel} + \mathbf{u}\|^2 = \|\mathbf{w}_{\parallel}\|^2 + \|\mathbf{u}\|^2$$

This means  $\|\mathbf{w}_{\parallel}\|^2 \leq \|\mathbf{w}^*\|^2$ .

- Conclusion:** The total cost for  $\mathbf{w}_{\parallel}$  is always less than or equal to the cost for  $\mathbf{w}^*$ .

# Kernelized Ridge Regression

- The Representer Theorem lets us write an equivalent optimization problem in terms of  $\alpha \in \mathbb{R}^N$  instead of  $\mathbf{w} \in \mathbb{R}^D$ .

# Kernelized Ridge Regression

- The Representer Theorem lets us write an equivalent optimization problem in terms of  $\alpha \in \mathbb{R}^N$  instead of  $\mathbf{w} \in \mathbb{R}^D$ .
- **Primal Problem (in  $\mathbf{w}$ ):**

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

# Kernelized Ridge Regression

- The Representer Theorem lets us write an equivalent optimization problem in terms of  $\alpha \in \mathbb{R}^N$  instead of  $\mathbf{w} \in \mathbb{R}^D$ .
- **Primal Problem (in  $\mathbf{w}$ ):**

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- **Dual Problem (in  $\alpha$ ):**

$$\alpha^* = \arg \min_{\alpha} \frac{1}{2} \alpha^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N) \alpha - \frac{1}{N} \alpha^\top \mathbf{y}$$

# Kernelized Ridge Regression

- The Representer Theorem lets us write an equivalent optimization problem in terms of  $\alpha \in \mathbb{R}^N$  instead of  $\mathbf{w} \in \mathbb{R}^D$ .
- **Primal Problem (in  $\mathbf{w}$ ):**

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \frac{1}{2N} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

- **Dual Problem (in  $\alpha$ ):**

$$\alpha^* = \arg \min_{\alpha} \frac{1}{2} \alpha^\top (\frac{1}{N} \mathbf{X} \mathbf{X}^\top + \lambda \mathbf{I}_N) \alpha - \frac{1}{N} \alpha^\top \mathbf{y}$$

- The dual problem only depends on the data through the **kernel matrix  $\mathbf{K} = \mathbf{X}\mathbf{X}^\top$** .

# The Kernel Matrix

- The dual formulations for SVM and Ridge Regression depend only on the  $N \times N$  matrix  $\mathbf{K} = \mathbf{XX}^\top$ .

# The Kernel Matrix

- The dual formulations for SVM and Ridge Regression depend only on the  $N \times N$  matrix  $\mathbf{K} = \mathbf{XX}^\top$ .
- This is called the **Kernel Matrix** or Gram Matrix.

# The Kernel Matrix

- The dual formulations for SVM and Ridge Regression depend only on the  $N \times N$  matrix  $\mathbf{K} = \mathbf{XX}^\top$ .
- This is called the **Kernel Matrix** or **Gram Matrix**.
- Each element  $\mathbf{K}_{ij}$  is the inner product between two data points:

$$\mathbf{K}_{ij} = \mathbf{x}_i^\top \mathbf{x}_j$$

# The Kernel Matrix

- The dual formulations for SVM and Ridge Regression depend only on the  $N \times N$  matrix  $\mathbf{K} = \mathbf{XX}^\top$ .
- This is called the **Kernel Matrix** or **Gram Matrix**.
- Each element  $\mathbf{K}_{ij}$  is the inner product between two data points:

$$\mathbf{K}_{ij} = \mathbf{x}_i^\top \mathbf{x}_j$$

- $\mathbf{K} = \mathbf{XX}^\top = \begin{pmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \mathbf{x}_1^\top \mathbf{x}_2 & \cdots & \mathbf{x}_1^\top \mathbf{x}_N \\ \mathbf{x}_2^\top \mathbf{x}_1 & \mathbf{x}_2^\top \mathbf{x}_2 & \cdots & \mathbf{x}_2^\top \mathbf{x}_N \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{x}_N^\top \mathbf{x}_1 & \mathbf{x}_N^\top \mathbf{x}_2 & \cdots & \mathbf{x}_N^\top \mathbf{x}_N \end{pmatrix}$

# The Kernel Matrix

- The dual formulations for SVM and Ridge Regression depend only on the  $N \times N$  matrix  $\mathbf{K} = \mathbf{XX}^\top$ .
- This is called the **Kernel Matrix** or **Gram Matrix**.
- Each element  $\mathbf{K}_{ij}$  is the inner product between two data points:

$$\mathbf{K}_{ij} = \mathbf{x}_i^\top \mathbf{x}_j$$

- $\mathbf{K} = \mathbf{XX}^\top = \begin{pmatrix} \mathbf{x}_1^\top \mathbf{x}_1 & \mathbf{x}_1^\top \mathbf{x}_2 & \cdots & \mathbf{x}_1^\top \mathbf{x}_N \\ \mathbf{x}_2^\top \mathbf{x}_1 & \mathbf{x}_2^\top \mathbf{x}_2 & \cdots & \mathbf{x}_2^\top \mathbf{x}_N \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{x}_N^\top \mathbf{x}_1 & \mathbf{x}_N^\top \mathbf{x}_2 & \cdots & \mathbf{x}_N^\top \mathbf{x}_N \end{pmatrix}$
- This matrix is symmetric ( $\mathbf{K} = \mathbf{K}^\top$ ) and positive semi-definite ( $\mathbf{K} \geq 0$ ).

# The Kernel Trick: Feature Maps

- What if our features  $x$  aren't good enough? We can augment them with a **feature map**  $\phi(x)$ .

# The Kernel Trick: Feature Maps

- What if our features  $x$  aren't good enough? We can augment them with a **feature map**  $\phi(x)$ .
- Example:  $x = [x_1, x_2] \rightarrow \phi(x) = [x_1^2, x_2^2, \sqrt{2}x_1x_2]$

# The Kernel Trick: Feature Maps


- What if our features  $x$  aren't good enough? We can augment them with a **feature map**  $\phi(x)$ .
- Example:  $x = [x_1, x_2] \rightarrow \phi(x) = [x_1^2, x_2^2, \sqrt{2}x_1x_2]$
- Our algorithms now run in this new, higher-dimensional feature space.
  - $w$  is now in this new space.
  - The kernel matrix becomes  $K_{ij} = \phi(x_i)^\top \phi(x_j)$ .

# The Kernel Trick: Feature Maps

- What if our features  $x$  aren't good enough? We can augment them with a **feature map**  $\phi(x)$ .
- Example:  $x = [x_1, x_2] \rightarrow \phi(x) = [x_1^2, x_2^2, \sqrt{2}x_1x_2]$
- Our algorithms now run in this new, higher-dimensional feature space.
  - $w$  is now in this new space.
  - The kernel matrix becomes  $K_{ij} = \phi(x_i)^\top \phi(x_j)$ .
- **Problem:** If  $\phi(x)$  is very high-dimensional (or infinite!), computing  $\phi(x)$  for all data and then  $K$  is extremely expensive or impossible.

# The Kernel Trick: Feature Maps

- What if our features  $x$  aren't good enough? We can augment them with a **feature map**  $\phi(x)$ .
- Example:  $x = [x_1, x_2] \rightarrow \phi(x) = [x_1^2, x_2^2, \sqrt{2}x_1x_2]$
- Our algorithms now run in this new, higher-dimensional feature space.
  - $w$  is now in this new space.
  - The kernel matrix becomes  $K_{ij} = \phi(x_i)^\top \phi(x_j)$ .
- **Problem:** If  $\phi(x)$  is very high-dimensional (or infinite!), computing  $\phi(x)$  for all data and then  $K$  is extremely expensive or impossible.
- **Example:** Data not linearly separable in 1D.



# Kernel Matrix with Feature Maps

- When we use a feature map  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{\tilde{D}}$ , our data matrix  $\mathbf{X}$  becomes  $\Phi$ .

# Kernel Matrix with Feature Maps

- When we use a feature map  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{\tilde{D}}$ , our data matrix  $\mathbf{X}$  becomes  $\Phi$ .
- The kernel matrix is now computed in the new feature space:

$$\mathbf{K}_{ij} = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

# Kernel Matrix with Feature Maps

- When we use a feature map  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{\tilde{D}}$ , our data matrix  $\mathbf{X}$  becomes  $\Phi$ .
- The kernel matrix is now computed in the new feature space:

$$\mathbf{K}_{ij} = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

- This is  $\mathbf{K} = \Phi \Phi^\top$ .

# Kernel Matrix with Feature Maps

- When we use a feature map  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{\tilde{D}}$ , our data matrix  $\mathbf{X}$  becomes  $\Phi$ .
- The kernel matrix is now computed in the new feature space:

$$\mathbf{K}_{ij} = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

- This is  $\mathbf{K} = \Phi \Phi^\top$ .
- **The Problem:** If the new dimension  $\tilde{D}$  is very large (e.g.,  $D \tilde{D}$ ), computing \*one\* inner product  $\phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$  costs  $O(\tilde{D})$ .

# Kernel Matrix with Feature Maps


- When we use a feature map  $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^{\tilde{D}}$ , our data matrix  $\mathbf{X}$  becomes  $\Phi$ .
- The kernel matrix is now computed in the new feature space:

$$\mathbf{K}_{ij} = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

- This is  $\mathbf{K} = \Phi \Phi^\top$ .
- **The Problem:** If the new dimension  $\tilde{D}$  is very large (e.g.,  $D \tilde{D}$ ), computing \*one\* inner product  $\phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$  costs  $O(\tilde{D})$ .
- This is too expensive, or impossible if  $\tilde{D} = \infty$ .


# Usefulness of Feature Spaces

Data that is not linearly separable in 2D...



# Usefulness of Feature Spaces

Data that is not linearly separable in 2D...



...can become linearly separable when mapped to a higher dimension (3D).

# The Kernel Trick: The “Shortcut”

**The Kernel Trick:** If we can find a **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$  that computes the inner product in the feature space *directly*, we never need to build  $\phi(\mathbf{x})$ .

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

# The Kernel Trick: The “Shortcut”

**The Kernel Trick:** If we can find a **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$  that computes the inner product in the feature space *directly*, we never need to build  $\phi(\mathbf{x})$ .

$$\kappa(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^\top \phi(\mathbf{x}_j)$$

- We only need to compute the  $N \times N$  kernel matrix  $\mathbf{K}_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$ .
- All “kernelized” algorithms (like SVM dual, Ridge dual) only need  $\mathbf{K}$ .
- We can work in an incredibly high-dimensional space, as long as we have an efficient  $\kappa(\cdot, \cdot)$ .

# Example 1: Polynomial Kernel

- **Kernel Function:**  $\kappa(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}')^2$

## Example 1: Polynomial Kernel

- **Kernel Function:**  $\kappa(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}')^2$
- Let  $\mathbf{x} \in \mathbb{R}^3$ .  $\kappa(\mathbf{x}, \mathbf{x}') = (x_1 x'_1 + x_2 x'_2 + x_3 x'_3)^2$

# Example 1: Polynomial Kernel

- **Kernel Function:**  $\kappa(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}')^2$
- Let  $\mathbf{x} \in \mathbb{R}^3$ .  $\kappa(\mathbf{x}, \mathbf{x}') = (x_1x'_1 + x_2x'_2 + x_3x'_3)^2$
- **Implicit Feature Map**  $\phi(\mathbf{x}) \in \mathbb{R}^6$ :

$$\phi(\mathbf{x})^\top = [x_1^2, x_2^2, x_3^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \sqrt{2}x_2x_3]$$

# Example 1: Polynomial Kernel

- **Kernel Function:**  $\kappa(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}')^2$
- Let  $\mathbf{x} \in \mathbb{R}^3$ .  $\kappa(\mathbf{x}, \mathbf{x}') = (x_1x'_1 + x_2x'_2 + x_3x'_3)^2$
- **Implicit Feature Map**  $\phi(\mathbf{x}) \in \mathbb{R}^6$ :

$$\phi(\mathbf{x})^\top = [x_1^2, x_2^2, x_3^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \sqrt{2}x_2x_3]$$

- **Cost:**
  - Computing  $\kappa(\mathbf{x}, \mathbf{x}')$ : 1 inner product ( $O(D)$ ), 1 multiplication.
  - Computing  $\phi(\mathbf{x})$  and then the dot product:  $O(D^2)$  operations.

# Example 1: Polynomial Kernel

- **Kernel Function:**  $\kappa(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^\top \mathbf{x}')^2$
- Let  $\mathbf{x} \in \mathbb{R}^3$ .  $\kappa(\mathbf{x}, \mathbf{x}') = (x_1 x'_1 + x_2 x'_2 + x_3 x'_3)^2$
- **Implicit Feature Map**  $\phi(\mathbf{x}) \in \mathbb{R}^6$ :

$$\phi(\mathbf{x})^\top = [x_1^2, x_2^2, x_3^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \sqrt{2}x_2x_3]$$

- **Cost:**
  - Computing  $\kappa(\mathbf{x}, \mathbf{x}')$ : 1 inner product ( $O(D)$ ), 1 multiplication.
  - Computing  $\phi(\mathbf{x})$  and then the dot product:  $O(D^2)$  operations.
- The kernel function is a huge computational shortcut!

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- This is a very popular and powerful kernel. (Note:  $\gamma = 1$  in your notes, e.g.,  $\exp(-(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}'))$ )

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- This is a very popular and powerful kernel. (Note:  $\gamma = 1$  in your notes, e.g.,  $\exp(-(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}'))$ )
- **Implicit Feature Map  $\phi(\mathbf{x})$ :**
  - The corresponding feature space is **infinite-dimensional!**

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- This is a very popular and powerful kernel. (Note:  $\gamma = 1$  in your notes, e.g.,  $\exp(-(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}'))$ )
- **Implicit Feature Map  $\phi(\mathbf{x})$ :**
  - The corresponding feature space is **infinite-dimensional!**
- We can see this by expanding the simple 1D case ( $x \in \mathbb{R}$ ,  $\gamma = 1$ ):

$$\begin{aligned}\kappa(x, x') &= e^{-(x-x')^2} = e^{-x^2} e^{-(x')^2} e^{2xx'} \\ &= e^{-x^2} e^{-(x')^2} \sum_{k=0}^{\infty} \frac{2^k (x)^k (x')^k}{k!}\end{aligned}$$

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- This is a very popular and powerful kernel. (Note:  $\gamma = 1$  in your notes, e.g.,  $\exp(-(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}'))$ )
- **Implicit Feature Map  $\phi(\mathbf{x})$ :**
  - The corresponding feature space is **infinite-dimensional!**
- We can see this by expanding the simple 1D case ( $x \in \mathbb{R}$ ,  $\gamma = 1$ ):

$$\begin{aligned}\kappa(x, x') &= e^{-(x-x')^2} = e^{-x^2} e^{-(x')^2} e^{2xx'} \\ &= e^{-x^2} e^{-(x')^2} \sum_{k=0}^{\infty} \frac{2^k (x)^k (x')^k}{k!}\end{aligned}$$

- This is an inner product between two infinite vectors:

$$\phi(x)^\top = e^{-x^2} \left[ \dots, \sqrt{\frac{2^k}{k!}} x^k, \dots \right]$$

## Example 2: RBF (Gaussian) Kernel

- **Kernel Function:**

$$\kappa(\mathbf{x}, \mathbf{x}') = \exp\left(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2\right)$$

- This is a very popular and powerful kernel. (Note:  $\gamma = 1$  in your notes, e.g.,  $\exp(-(\mathbf{x} - \mathbf{x}')^\top (\mathbf{x} - \mathbf{x}'))$ )
- **Implicit Feature Map  $\phi(\mathbf{x})$ :**
  - The corresponding feature space is **infinite-dimensional!**
- We can see this by expanding the simple 1D case ( $x \in \mathbb{R}$ ,  $\gamma = 1$ ):

$$\begin{aligned}\kappa(x, x') &= e^{-(x-x')^2} = e^{-x^2} e^{-(x')^2} e^{2xx'} \\ &= e^{-x^2} e^{-(x')^2} \sum_{k=0}^{\infty} \frac{2^k (x)^k (x')^k}{k!}\end{aligned}$$

- This is an inner product between two infinite vectors:

$$\phi(x)^\top = e^{-x^2} \left[ \dots, \sqrt{\frac{2^k}{k!}} x^k, \dots \right]$$

- The Kernel Trick lets us compute this impossible-to-build inner product.

# Building New Kernels from Old

If  $\kappa_1$  and  $\kappa_2$  are valid kernels, so are:

- **Positive Scaling:**  $\kappa(\mathbf{x}, \mathbf{x}') = c \kappa_1(\mathbf{x}, \mathbf{x}')$  for  $c \geq 0$ .

# Building New Kernels from Old

If  $\kappa_1$  and  $\kappa_2$  are valid kernels, so are:

- **Positive Scaling:**  $\kappa(\mathbf{x}, \mathbf{x}') = c \kappa_1(\mathbf{x}, \mathbf{x}')$  for  $c \geq 0$ .
- **Sum:**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}') + \kappa_2(\mathbf{x}, \mathbf{x}')$ .
  - Proof:  $\phi(\cdot) = [\phi_1(\cdot), \phi_2(\cdot)]$

# Building New Kernels from Old

If  $\kappa_1$  and  $\kappa_2$  are valid kernels, so are:

- **Positive Scaling:**  $\kappa(\mathbf{x}, \mathbf{x}') = c \kappa_1(\mathbf{x}, \mathbf{x}')$  for  $c \geq 0$ .
- **Sum:**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}') + \kappa_2(\mathbf{x}, \mathbf{x}')$ .
  - Proof:  $\phi(\cdot) = [\phi_1(\cdot), \phi_2(\cdot)]$
- **Product:**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}')\kappa_2(\mathbf{x}, \mathbf{x}')$ .
  - Proof:  $\phi(\cdot)$  has entries  $(\phi_1(\cdot))_i(\phi_2(\cdot))_j$

# Building New Kernels from Old

If  $\kappa_1$  and  $\kappa_2$  are valid kernels, so are:

- **Positive Scaling:**  $\kappa(\mathbf{x}, \mathbf{x}') = c \kappa_1(\mathbf{x}, \mathbf{x}')$  for  $c \geq 0$ .
- **Sum:**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}') + \kappa_2(\mathbf{x}, \mathbf{x}')$ .
  - Proof:  $\phi(\cdot) = [\phi_1(\cdot), \phi_2(\cdot)]$
- **Product:**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}')\kappa_2(\mathbf{x}, \mathbf{x}')$ .
  - Proof:  $\phi(\cdot)$  has entries  $(\phi_1(\cdot))_i(\phi_2(\cdot))_j$
- **Function mapping:**  $\kappa(\mathbf{x}, \mathbf{x}') = f(\mathbf{x})f(\mathbf{x}')$  for any real-valued  $f$ .

# Proofs for Building Kernels

- **Proof (Sum):**  $\kappa(\mathbf{x}, \mathbf{x}') = \alpha\kappa_1(\mathbf{x}, \mathbf{x}') + \beta\kappa_2(\mathbf{x}, \mathbf{x}')$ 
  - We can define a new feature map  $\phi(\mathbf{x})$  by concatenating the scaled original feature maps:

$$\phi(\mathbf{x}) = \begin{bmatrix} \sqrt{\alpha}\phi_1(\mathbf{x}) \\ \sqrt{\beta}\phi_2(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^{D_1+D_2}$$

- The inner product is then:

$$\phi(\mathbf{x})^\top \phi(\mathbf{x}') = \alpha\phi_1(\mathbf{x})^\top \phi_1(\mathbf{x}') + \beta\phi_2(\mathbf{x})^\top \phi_2(\mathbf{x}')$$

# Proofs for Building Kernels

- **Proof (Sum):**  $\kappa(\mathbf{x}, \mathbf{x}') = \alpha\kappa_1(\mathbf{x}, \mathbf{x}') + \beta\kappa_2(\mathbf{x}, \mathbf{x}')$

- We can define a new feature map  $\phi(\mathbf{x})$  by concatenating the scaled original feature maps:

$$\phi(\mathbf{x}) = \begin{bmatrix} \sqrt{\alpha}\phi_1(\mathbf{x}) \\ \sqrt{\beta}\phi_2(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^{D_1+D_2}$$

- The inner product is then:

$$\phi(\mathbf{x})^\top \phi(\mathbf{x}') = \alpha\phi_1(\mathbf{x})^\top \phi_1(\mathbf{x}') + \beta\phi_2(\mathbf{x})^\top \phi_2(\mathbf{x}')$$

- **Proof (Product):**  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa_1(\mathbf{x}, \mathbf{x}')\kappa_2(\mathbf{x}, \mathbf{x}')$

- We define a new feature map  $\phi(\mathbf{x})$  as the tensor (outer) product of the original features.
- The new feature vector contains all possible products  $(\phi_1(\mathbf{x}))_i(\phi_2(\mathbf{x}))_j$ .
- The inner product becomes:

$$\begin{aligned} \phi(\mathbf{x})^\top \phi(\mathbf{x}') &= \sum_{i,j} (\phi_1(\mathbf{x}))_i (\phi_2(\mathbf{x}))_j (\phi_1(\mathbf{x}'))_i (\phi_2(\mathbf{x}'))_j \\ &= \left( \sum_i (\phi_1(\mathbf{x}))_i (\phi_1(\mathbf{x}'))_i \right) \left( \sum_j (\phi_2(\mathbf{x}))_j (\phi_2(\mathbf{x}'))_j \right) \\ &= \kappa_1(\mathbf{x}, \mathbf{x}')\kappa_2(\mathbf{x}, \mathbf{x}') \end{aligned}$$

# Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.

# Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.
- A new data point  $x$  is classified by  $y = \mathbf{w}^* \cdot \phi(x)$ .

## Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.
- A new data point  $x$  is classified by  $y = \mathbf{w}^* \cdot \phi(x)$ .
- We can't compute  $\mathbf{w}^*$  or  $\phi(x)$  if they are infinite!

# Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.
- A new data point  $x$  is classified by  $y = \mathbf{w}^* \cdot \phi(x)$ .
- We can't compute  $\mathbf{w}^*$  or  $\phi(x)$  if they are infinite!
- But we can use the Representer Theorem:

$$\mathbf{w}^* = \sum_{n=1}^N \alpha_n^* \phi(\mathbf{x}_n) \quad (\text{using } \phi \text{ in } \mathbf{w} = \mathbf{X}^\top \boldsymbol{\alpha})$$

# Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.
- A new data point  $\mathbf{x}$  is classified by  $y = \mathbf{w}^* \cdot \phi(\mathbf{x})$ .
- We can't compute  $\mathbf{w}^*$  or  $\phi(\mathbf{x})$  if they are infinite!
- But we can use the Representer Theorem:

$$\mathbf{w}^* = \sum_{n=1}^N \alpha_n^* \phi(\mathbf{x}_n) \quad (\text{using } \phi \text{ in } \mathbf{w} = \mathbf{X}^\top \boldsymbol{\alpha})$$

- Now, the prediction becomes:

$$\begin{aligned} y &= \left( \sum_{n=1}^N \alpha_n^* \phi(\mathbf{x}_n) \right)^\top \phi(\mathbf{x}) \\ &= \sum_{n=1}^N \alpha_n^* (\phi(\mathbf{x}_n)^\top \phi(\mathbf{x})) \\ &= \sum_{n=1}^N \alpha_n^* \kappa(\mathbf{x}_n, \mathbf{x}) \end{aligned}$$

# Prediction with Kernels

- This whole enterprise is useless if we can't make predictions on new data.
- A new data point  $\mathbf{x}$  is classified by  $y = \mathbf{w}^* \cdot \phi(\mathbf{x})$ .
- We can't compute  $\mathbf{w}^*$  or  $\phi(\mathbf{x})$  if they are infinite!
- But we can use the Representer Theorem:

$$\mathbf{w}^* = \sum_{n=1}^N \alpha_n^* \phi(\mathbf{x}_n) \quad (\text{using } \phi \text{ in } \mathbf{w} = \mathbf{X}^\top \boldsymbol{\alpha})$$

- Now, the prediction becomes:

$$\begin{aligned} y &= \left( \sum_{n=1}^N \alpha_n^* \phi(\mathbf{x}_n) \right)^\top \phi(\mathbf{x}) \\ &= \sum_{n=1}^N \alpha_n^* (\phi(\mathbf{x}_n)^\top \phi(\mathbf{x})) \\ &= \sum_{n=1}^N \alpha_n^* \kappa(\mathbf{x}_n, \mathbf{x}) \end{aligned}$$

# Properties of Kernels: Mercer's Condition

How do we know if a function  $\kappa(x, x')$  is a valid kernel?  
(i.e., corresponds to an inner product in \*some\* feature space?)

# Properties of Kernels: Mercer's Condition

How do we know if a function  $\kappa(x, x')$  is a valid kernel?  
(i.e., corresponds to an inner product in \*some\* feature space?)

**Mercer's Condition:** A function  $\kappa$  is a valid kernel if and only if:

# Properties of Kernels: Mercer's Condition

How do we know if a function  $\kappa(\mathbf{x}, \mathbf{x}')$  is a valid kernel?  
(i.e., corresponds to an inner product in \*some\* feature space?)

**Mercer's Condition:** A function  $\kappa$  is a valid kernel if and only if:

- 1 It is **symmetric**:

$$\kappa(\mathbf{x}, \mathbf{x}') = \kappa(\mathbf{x}', \mathbf{x})$$

# Properties of Kernels: Mercer's Condition

How do we know if a function  $\kappa(\mathbf{x}, \mathbf{x}')$  is a valid kernel?  
(i.e., corresponds to an inner product in \*some\* feature space?)

**Mercer's Condition:** A function  $\kappa$  is a valid kernel if and only if:

- 1 It is **symmetric**:

$$\kappa(\mathbf{x}, \mathbf{x}') = \kappa(\mathbf{x}', \mathbf{x})$$

- 2 The kernel matrix  $\mathbf{K}$  is **positive semi-definite** for *any* set of points  $\{\mathbf{x}_n\}_{n=1}^N$ .

$$\mathbf{K}_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j) \implies \mathbf{c}^\top \mathbf{K} \mathbf{c} \geq 0 \text{ for all } \mathbf{c} \in \mathbb{R}^N$$

## Recap: Key Ideas

- Many algorithms (SVM, Ridge Regression, PCA) can be “kernelized” because their solutions can be written to depend only on inner products,  $\mathbf{X}\mathbf{X}^\top$ .

## Recap: Key Ideas

- Many algorithms (SVM, Ridge Regression, PCA) can be “kernelized” because their solutions can be written to depend only on inner products,  $\mathbf{X}\mathbf{X}^\top$ .
- This allows us to generalize a linear model by replacing the simple inner product  $\mathbf{x}^\top \mathbf{x}'$  with a non-linear **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$ .

## Recap: Key Ideas

- Many algorithms (SVM, Ridge Regression, PCA) can be “kernelized” because their solutions can be written to depend only on inner products,  $\mathbf{X}\mathbf{X}^\top$ .
- This allows us to generalize a linear model by replacing the simple inner product  $\mathbf{x}^\top \mathbf{x}'$  with a non-linear **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$ .
- This creates a non-linear model in the original input space, but one that is still linear (and convex) in a new, high-dimensional feature space.

## Recap: Key Ideas

- Many algorithms (SVM, Ridge Regression, PCA) can be “kernelized” because their solutions can be written to depend only on inner products,  $\mathbf{X}\mathbf{X}^\top$ .
- This allows us to generalize a linear model by replacing the simple inner product  $\mathbf{x}^\top \mathbf{x}'$  with a non-linear **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$ .
- This creates a non-linear model in the original input space, but one that is still linear (and convex) in a new, high-dimensional feature space.
- The **Kernel Trick** is that we can do this *without ever computing* the feature maps  $\phi(\mathbf{x})$ , which can be very high-dim or even infinite-dim.

## Recap: Key Ideas

- Many algorithms (SVM, Ridge Regression, PCA) can be “kernelized” because their solutions can be written to depend only on inner products,  $\mathbf{X}\mathbf{X}^\top$ .
- This allows us to generalize a linear model by replacing the simple inner product  $\mathbf{x}^\top \mathbf{x}'$  with a non-linear **kernel function**  $\kappa(\mathbf{x}, \mathbf{x}')$ .
- This creates a non-linear model in the original input space, but one that is still linear (and convex) in a new, high-dimensional feature space.
- The **Kernel Trick** is that we can do this *without ever computing* the feature maps  $\phi(\mathbf{x})$ , which can be very high-dim or even infinite-dim.
- To make a new prediction, we just need to compute the kernel (similarity) between the new point  $\mathbf{x}$  and all the training points  $\mathbf{x}_n$ .