

# Lecture 11: Neural Network (Basic Structure, Representation Power)

**Tao LIN**

SoE, Westlake University

December 2, 2025



- 1 Introduction and Motivation
- 2 Neural Networks: Basic Structure
- 3 Universal Approximation Theorem
- 4 Training Neural Networks

# Table of Contents

- ① Introduction and Motivation
- ② Neural Networks: Basic Structure
- ③ Universal Approximation Theorem
- ④ Training Neural Networks

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.
- **1958:** Frank Rosenblatt invents the [Perceptron](#) (a linear classifier).

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.
- **1958:** Frank Rosenblatt invents the [Perceptron](#) (a linear classifier).
- **1969:** Minsky & Papert publish "Perceptrons", showing that linear models cannot solve XOR. [AI Winter I](#).

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.
- **1958:** Frank Rosenblatt invents the [Perceptron](#) (a linear classifier).
- **1969:** Minsky & Papert publish "Perceptrons", showing that linear models cannot solve XOR. [AI Winter I](#).
- **1986:** Rumelhart, Hinton, & Williams popularize [Backpropagation](#), enabling training of multi-layer nets.

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.
- **1958:** Frank Rosenblatt invents the [Perceptron](#) (a linear classifier).
- **1969:** Minsky & Papert publish "Perceptrons", showing that linear models cannot solve XOR. [AI Winter I](#).
- **1986:** Rumelhart, Hinton, & Williams popularize [Backpropagation](#), enabling training of multi-layer nets.
- **1990s:** SVMs and Random Forests dominate. NNs are seen as hard to train and prone to overfitting.

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# A Brief History of Neural Networks

- **1943:** McCulloch & Pitts propose the first mathematical model of a neuron.
- **1958:** Frank Rosenblatt invents the [Perceptron](#) (a linear classifier).
- **1969:** Minsky & Papert publish "Perceptrons", showing that linear models cannot solve XOR. [AI Winter I](#).
- **1986:** Rumelhart, Hinton, & Williams popularize [Backpropagation](#), enabling training of multi-layer nets.
- **1990s:** SVMs and Random Forests dominate. NNs are seen as hard to train and prone to overfitting.
- **2012:** AlexNet wins ImageNet. The [Deep Learning](#) revolution begins (GPUs + Big Data + Algorithms).

**Takeaway** The **Deep Learning Revolution** (2012+) was enabled by three key factors: powerful GPUs, massive datasets, and improved algorithms (ReLU, dropout, batch normalization).

# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y = 1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y = 1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

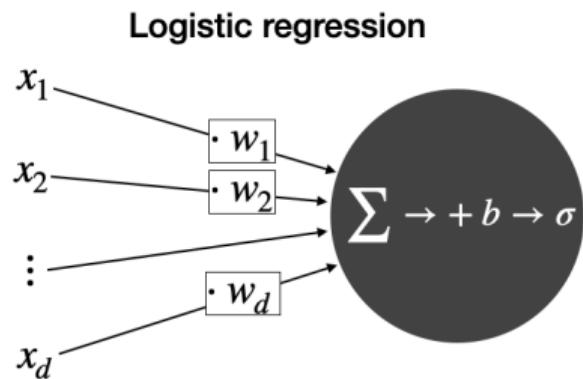
- **Problem:** They can only learn linear decision boundaries.

# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y = 1 | \mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

- Problem:** They can only learn **linear decision boundaries**.
- Example:** Consider data where  $y = 0$  is inside a unit circle and  $y = 1$  is outside.



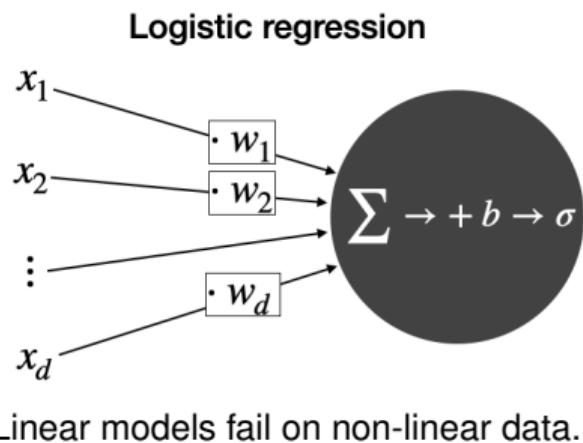
Linear models fail on non-linear data.

# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y=1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

- Problem:** They can only learn **linear decision boundaries**.
- Example:** Consider data where  $y = 0$  is inside a unit circle and  $y = 1$  is outside.
- A linear classifier cannot separate this!

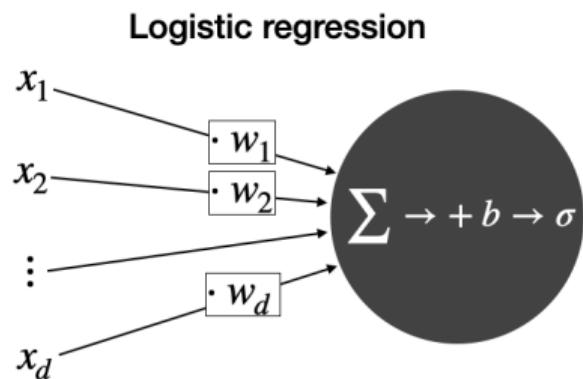


# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y=1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

- Problem:** They can only learn **linear decision boundaries**.
- Example:** Consider data where  $y = 0$  is inside a unit circle and  $y = 1$  is outside.
- A linear classifier cannot separate this!
- Old Solution:** Hand-craft features (e.g.,  $\phi(\mathbf{x}) = [x_1, x_2, x_1^2 + x_2^2]^\top$ ).



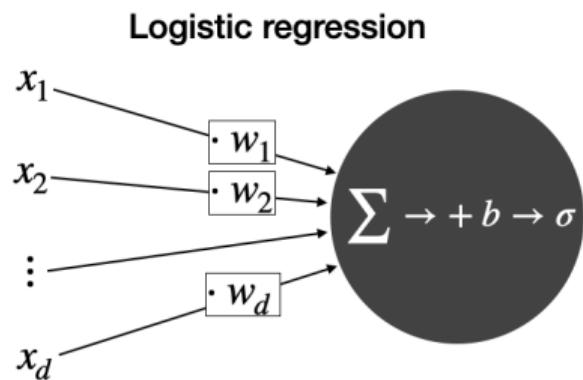
Linear models fail on non-linear data.

# Motivation: Limitations of Linear Models

- We have seen linear models like Logistic Regression:

$$p(y=1|\mathbf{x}) = \sigma(\mathbf{x}^\top \mathbf{w} + b)$$

- Problem:** They can only learn **linear decision boundaries**.
- Example:** Consider data where  $y = 0$  is inside a unit circle and  $y = 1$  is outside.
- A linear classifier cannot separate this!
- Old Solution:** Hand-craft features (e.g.,  $\phi(\mathbf{x}) = [x_1, x_2, x_1^2 + x_2^2]^\top$ ).
- New Solution:** Learn the features automatically using a **Neural Network**.



Linear models fail on non-linear data.

# Table of Contents

- ① Introduction and Motivation
- ② Neural Networks: Basic Structure
- ③ Universal Approximation Theorem
- ④ Training Neural Networks

# The Basic Building Block: The Neuron

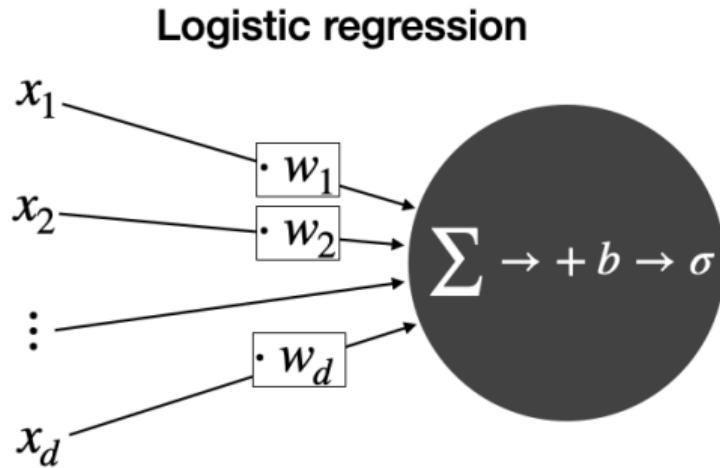
A single neuron  $j$  in layer  $l$  computes:

$$z_j^{(l)} = \sum_i w_{i,j}^{(l)} x_i^{(l-1)} + b_j^{(l)}$$

$$x_j^{(l)} = \phi(z_j^{(l)})$$

Vectorized form:  $\mathbf{x}^{(l)} = \phi(\mathbf{W}^{(l)\top} \mathbf{x}^{(l-1)} + \mathbf{b}^{(l)})$ .

- $w_{i,j}^{(l)}$ : Weights from layer  $l - 1$  to  $l$ .
- $b_j^{(l)}$ : Bias term.
- $\phi(\cdot)$ : Activation Function.

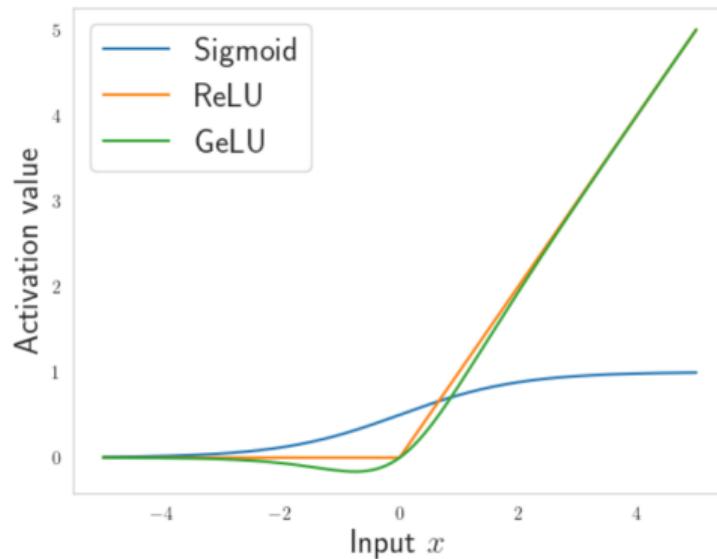


**Crucial Insight**  $\phi$  must be **non-linear**. If  $\phi$  were linear, the deep network would collapse into a single linear transformation:

$$\mathbf{W}_2(\mathbf{W}_1 \mathbf{x}) = (\mathbf{W}_2 \mathbf{W}_1) \mathbf{x}$$

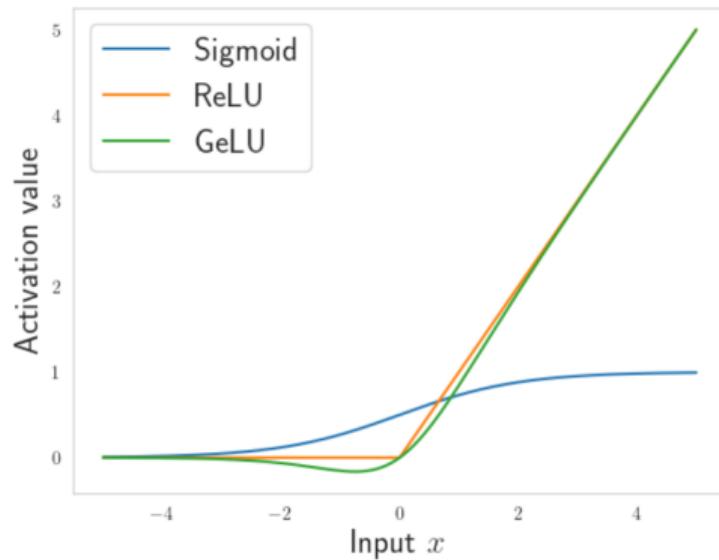
# Common Activation Functions

- **Sigmoid:**  $\sigma(z) = \frac{1}{1+e^{-z}}$ 
  - Squashes output to  $(0, 1)$ .
  - Interpretable as probability.
  - **Issue:** Vanishing gradients for large  $|z|$ .



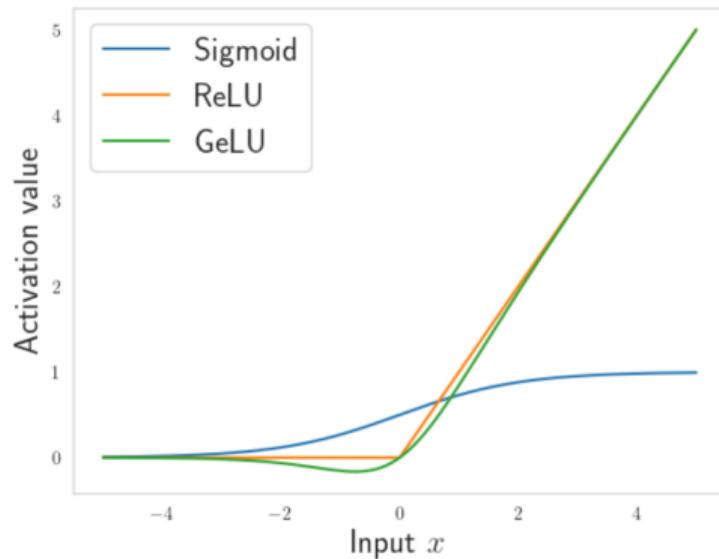
# Common Activation Functions

- **Sigmoid:**  $\sigma(z) = \frac{1}{1+e^{-z}}$ 
  - Squashes output to  $(0, 1)$ .
  - Interpretable as probability.
  - **Issue:** Vanishing gradients for large  $|z|$ .
- **ReLU:**  $\text{ReLU}(z) = \max(0, z)$ 
  - Default choice for modern deep learning.
  - Efficient, avoids vanishing gradients for  $z > 0$ .

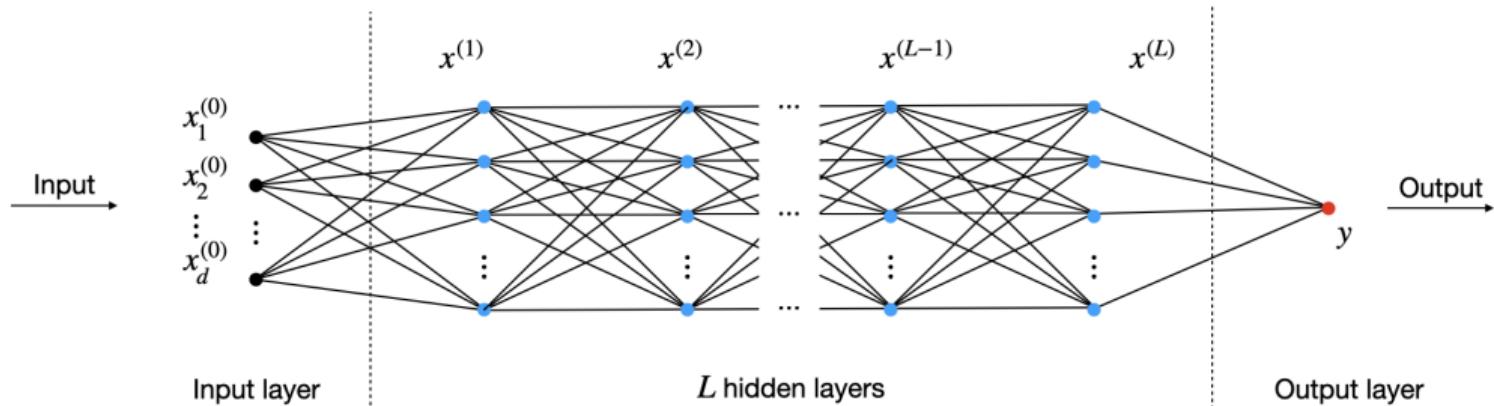


# Common Activation Functions

- **Sigmoid:**  $\sigma(z) = \frac{1}{1+e^{-z}}$ 
  - Squashes output to  $(0, 1)$ .
  - Interpretable as probability.
  - **Issue:** Vanishing gradients for large  $|z|$ .
- **ReLU:**  $\text{ReLU}(z) = \max(0, z)$ 
  - Default choice for modern deep learning.
  - Efficient, avoids vanishing gradients for  $z > 0$ .
- **Others:** Tanh, Leaky ReLU, GeLU, Swish...



# Multi-Layer Perceptron (MLP)



- **Input Layer:** Raw features  $\mathbf{x} \in \mathbb{R}^D$ .
- **Hidden Layers:** Learn abstract representations.
  - $f : \mathbb{R}^D \rightarrow \mathbb{R}^K$ .
  - Replaces the "domain expert" feature engineering.
- **Output Layer:** Final prediction (e.g., linear classifier on top of learned features).

# Table of Contents

- ① Introduction and Motivation
- ② Neural Networks: Basic Structure
- ③ Universal Approximation Theorem
- ④ Training Neural Networks

# Barron's Universal Approximation Theorem

Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  and define  $\tilde{f}(\omega) = \int_{\mathbb{R}^d} f(\mathbf{x}) e^{-i\omega^\top \mathbf{x}} d\mathbf{x}$ , its Fourier transform.

**Assumption:**  $\int_{\mathbb{R}^d} |\omega| |\tilde{f}(\omega)| d\omega \leq C$  (smoothness assumption)

**Claim:** For all  $n \geq 1$  and  $r > 0$ , there exists a function  $f_n$  of the form

$$f_n(\mathbf{x}) = \sum_{j=1}^n c_j \phi(\mathbf{x}^\top \mathbf{w}_j + b_j) + c_0$$

such that

$$\int_{|\mathbf{x}| \leq r} (f(\mathbf{x}) - f_n(\mathbf{x}))^2 d\mathbf{x} \leq \frac{(2Cr)^2}{n}$$

A. R. Barron, *Universal approximation bounds for superpositions of a sigmoidal function*, IEEE Trans. Inf. Theory, 1993

# All Sufficiently Smooth Functions Can Be Approximated

$$\int_{|\mathbf{x}| \leq r} (f(\mathbf{x}) - f_n(\mathbf{x}))^2 d\mathbf{x} \leq \frac{(2Cr)^2}{n} \quad (1)$$

- **The more neurons allowed, the smaller the error.**
  - Error decreases as  $\mathcal{O}(1/n)$  — very fast convergence!

**Takeaway** One-hidden-layer neural networks can approximate *any* smooth function arbitrarily well, with error decreasing as  $\mathcal{O}(1/n)$ . This is the theoretical foundation for deep learning!

# All Sufficiently Smooth Functions Can Be Approximated

$$\int_{|x| \leq r} (f(x) - f_n(x))^2 dx \leq \frac{(2Cr)^2}{n} \quad (1)$$

- **The more neurons allowed, the smaller the error.**
  - Error decreases as  $\mathcal{O}(1/n)$  — **very fast convergence!**
- **The smoother the function (the smaller  $C$ ), the smaller the error.**
  - The constant  $C$  measures the "complexity" of  $f$  in the frequency domain.

**Takeaway** One-hidden-layer neural networks can approximate *any* smooth function arbitrarily well, with error decreasing as  $\mathcal{O}(1/n)$ . This is the theoretical foundation for deep learning!

# All Sufficiently Smooth Functions Can Be Approximated

$$\int_{|x| \leq r} (f(x) - f_n(x))^2 dx \leq \frac{(2Cr)^2}{n} \quad (1)$$

- **The more neurons allowed, the smaller the error.**
  - Error decreases as  $\mathcal{O}(1/n)$  — very fast convergence!
- **The smoother the function (the smaller  $C$ ), the smaller the error.**
  - The constant  $C$  measures the "complexity" of  $f$  in the frequency domain.
- **The larger the domain (the larger  $r$ ), the greater the error.**
  - Need more neurons to approximate over larger regions.

**Takeaway** One-hidden-layer neural networks can approximate *any* smooth function arbitrarily well, with error decreasing as  $\mathcal{O}(1/n)$ . This is the theoretical foundation for deep learning!

# All Sufficiently Smooth Functions Can Be Approximated

$$\int_{|x| \leq r} (f(x) - f_n(x))^2 dx \leq \frac{(2Cr)^2}{n} \quad (1)$$

- **The more neurons allowed, the smaller the error.**
  - Error decreases as  $\mathcal{O}(1/n)$  — very fast convergence!
- **The smoother the function (the smaller  $C$ ), the smaller the error.**
  - The constant  $C$  measures the "complexity" of  $f$  in the frequency domain.
- **The larger the domain (the larger  $r$ ), the greater the error.**
  - Need more neurons to approximate over larger regions.
- **Approximation is in average (in  $L^2$ -norm).**
  - Not pointwise — some points may have larger errors.

**Takeaway** One-hidden-layer neural networks can approximate *any* smooth function arbitrarily well, with error decreasing as  $\mathcal{O}(1/n)$ . This is the theoretical foundation for deep learning!

# All Sufficiently Smooth Functions Can Be Approximated

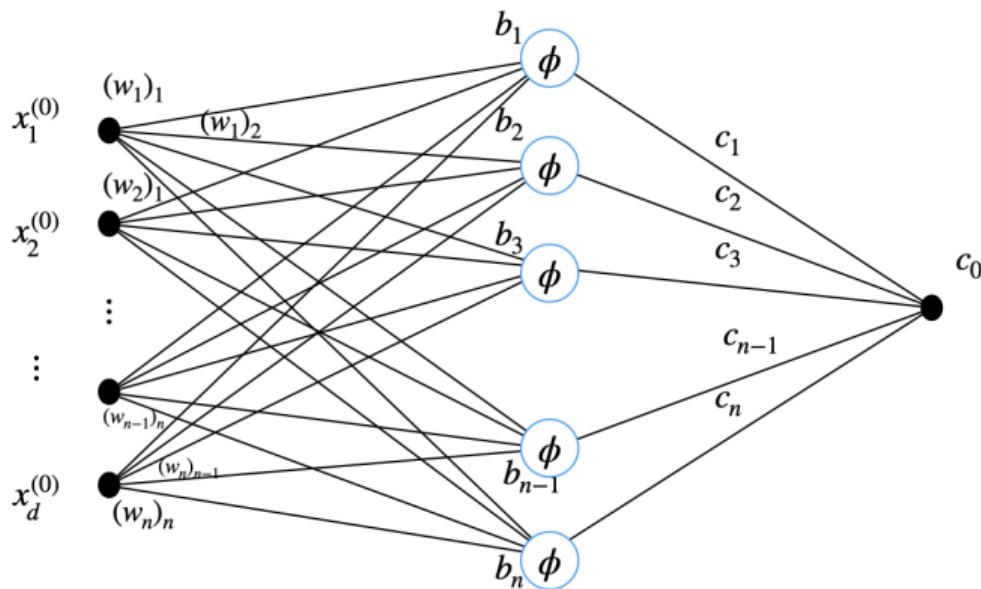
$$\int_{|x| \leq r} (f(x) - f_n(x))^2 dx \leq \frac{(2Cr)^2}{n} \quad (1)$$

- **The more neurons allowed, the smaller the error.**
  - Error decreases as  $\mathcal{O}(1/n)$  — very fast convergence!
- **The smoother the function (the smaller  $C$ ), the smaller the error.**
  - The constant  $C$  measures the "complexity" of  $f$  in the frequency domain.
- **The larger the domain (the larger  $r$ ), the greater the error.**
  - Need more neurons to approximate over larger regions.
- **Approximation is in average (in  $L^2$ -norm).**
  - Not pointwise — some points may have larger errors.
- **Applicable for any "sigmoid-like" activation function.**
  - Any  $\phi$  with  $\lim_{x \rightarrow -\infty} \phi(x) = 0$  and  $\lim_{x \rightarrow \infty} \phi(x) = 1$ .

**Takeaway** One-hidden-layer neural networks can approximate *any* smooth function arbitrarily well, with error decreasing as  $\mathcal{O}(1/n)$ . This is the theoretical foundation for deep learning!

# The Function $f_n$ is a One-Hidden-Layer NN with $n$ Nodes

$$f_n(\mathbf{x}) = \sum_{j=1}^n c_j \phi(\mathbf{x}^\top \mathbf{w}_j + b_j) + c_0 = \mathbf{c}^\top \phi(\mathbf{W}^\top \mathbf{x} + \mathbf{b}) + c_0$$

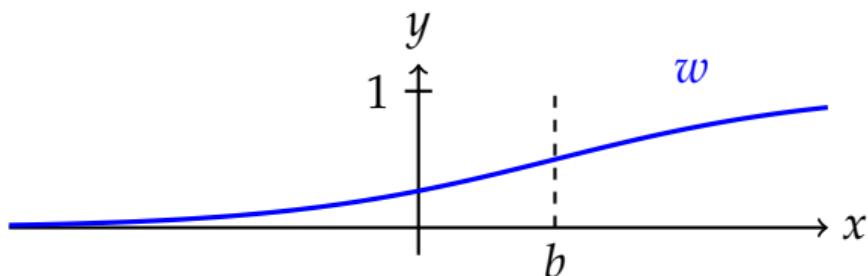


- **Input:**  $\mathbf{x} \in \mathbb{R}^d$
- **Weights:**  $\mathbf{w}_j \in \mathbb{R}^d$
- **Biases:**  $b_j \in \mathbb{R}$
- **Output weights:**  $c_j \in \mathbb{R}$
- **Activation:**  $\phi$  (sigmoid-like)
- **Output:** scalar  $f_n(\mathbf{x})$

# Intuition: Constructing Functions (Step 1)

How can we build any function using Sigmoids?

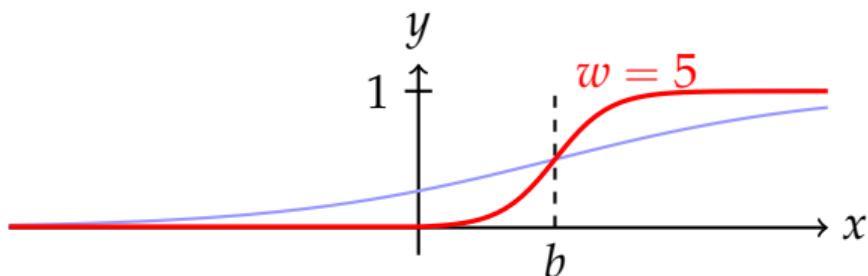
- Consider a neuron with weight  $w$  and bias  $-wb$ :  $y = \sigma(w(x - b))$ .



# Intuition: Constructing Functions (Step 1)

How can we build any function using Sigmoids?

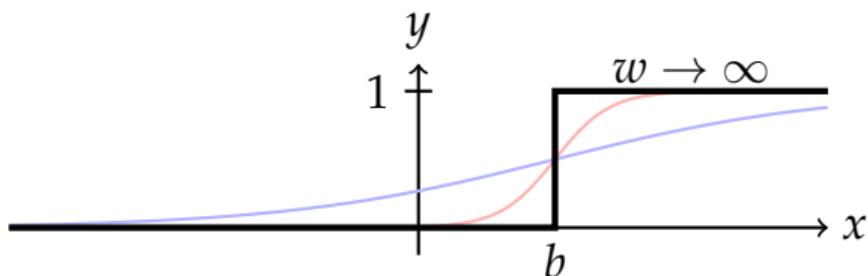
- Consider a neuron with weight  $w$  and bias  $-wb$ :  $y = \sigma(w(x - b))$ .
- As  $w \rightarrow \infty$ , the sigmoid becomes a **Step Function** at  $x = b$ .



# Intuition: Constructing Functions (Step 1)

How can we build any function using Sigmoids?

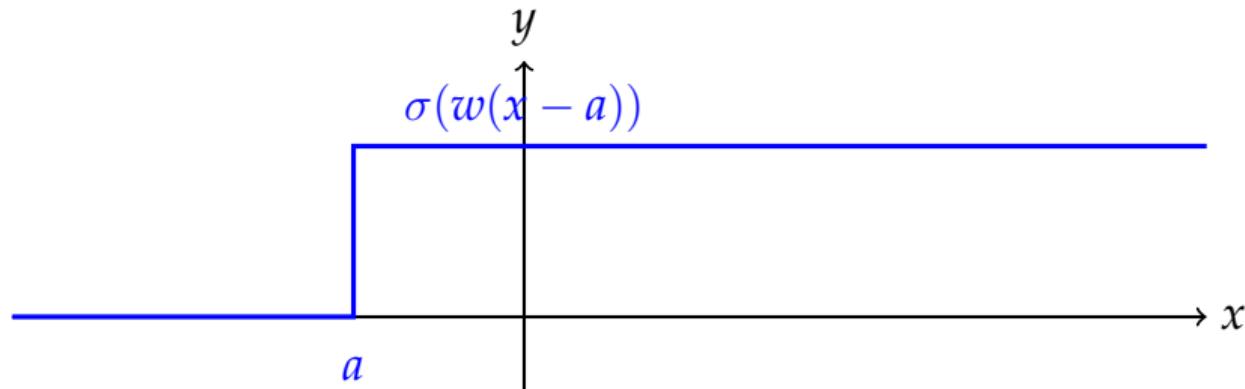
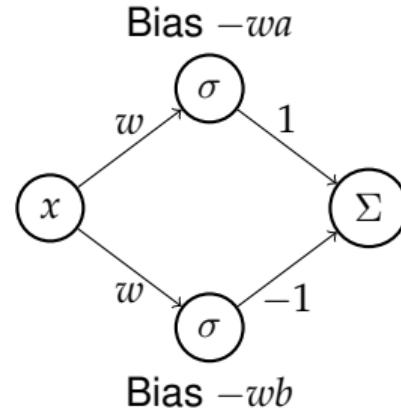
- Consider a neuron with weight  $w$  and bias  $-wb$ :  $y = \sigma(w(x - b))$ .
- As  $w \rightarrow \infty$ , the sigmoid becomes a **Step Function** at  $x = b$ .



# Intuition: Constructing Functions (Step 2)

- We can subtract two step functions to create a Rectangle (or "bump"):

$$\text{Bump}(x) \approx \sigma(w(x - a)) - \sigma(w(x - b))$$

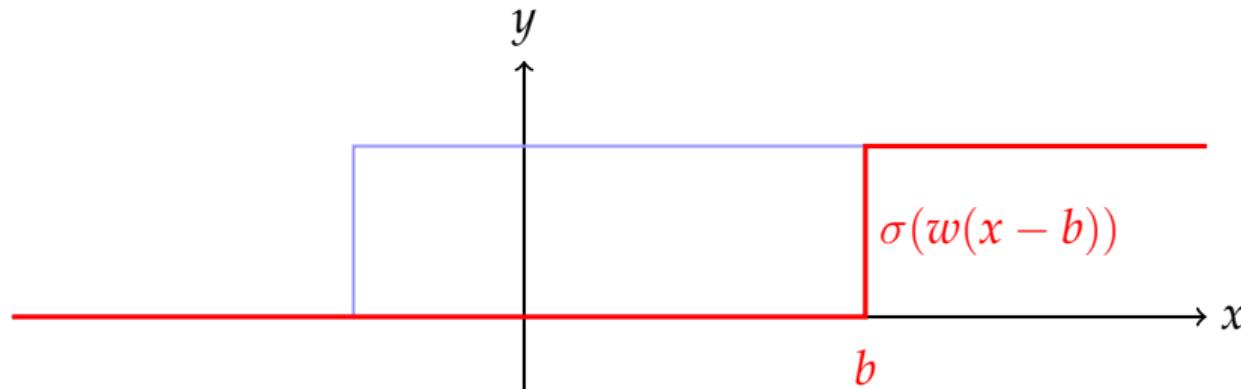
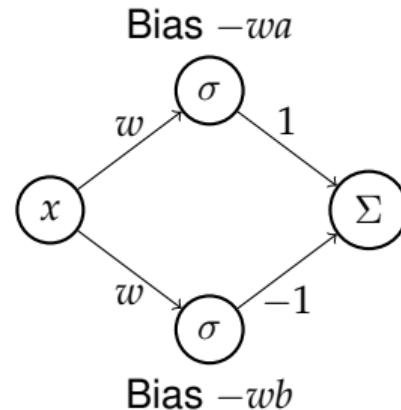


# Intuition: Constructing Functions (Step 2)

- We can subtract two step functions to create a Rectangle (or "bump"):

$$\text{Bump}(x) \approx \sigma(w(x - a)) - \sigma(w(x - b))$$

- This requires **2 neurons** in the hidden layer.

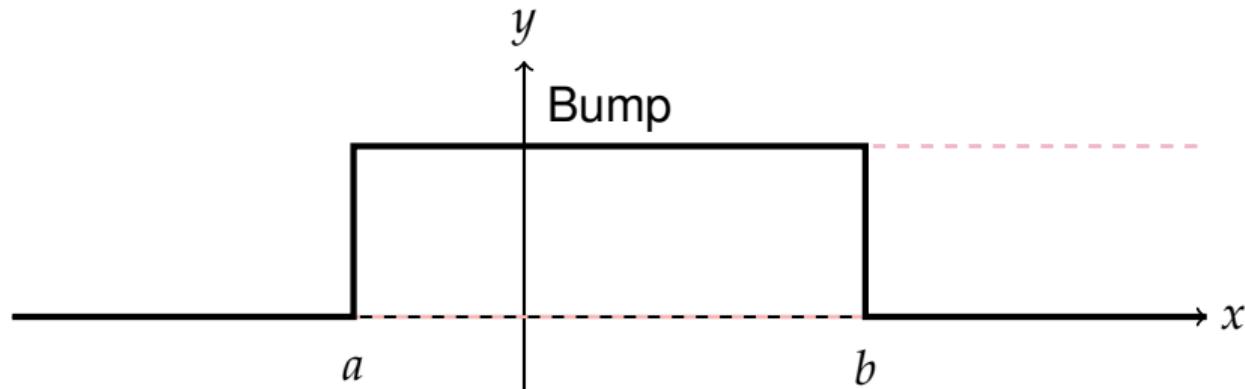
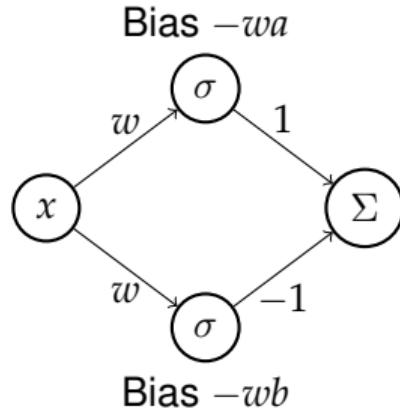


# Intuition: Constructing Functions (Step 2)

- We can subtract two step functions to create a **Rectangle** (or "bump"):

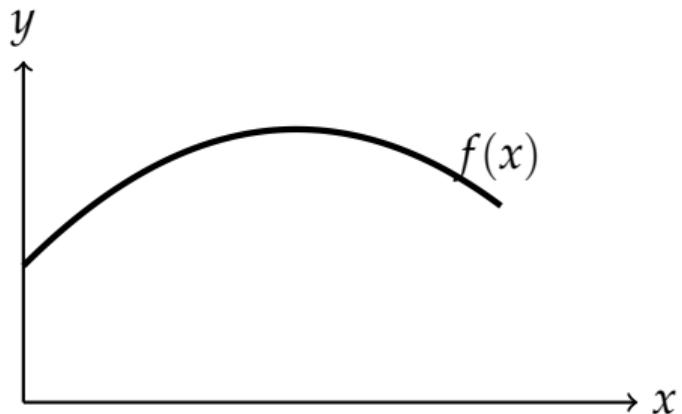
$$\text{Bump}(x) \approx \sigma(w(x - a)) - \sigma(w(x - b))$$

- This requires **2 neurons** in the hidden layer.
- We can control the width ( $b - a$ ) and position of the bump.



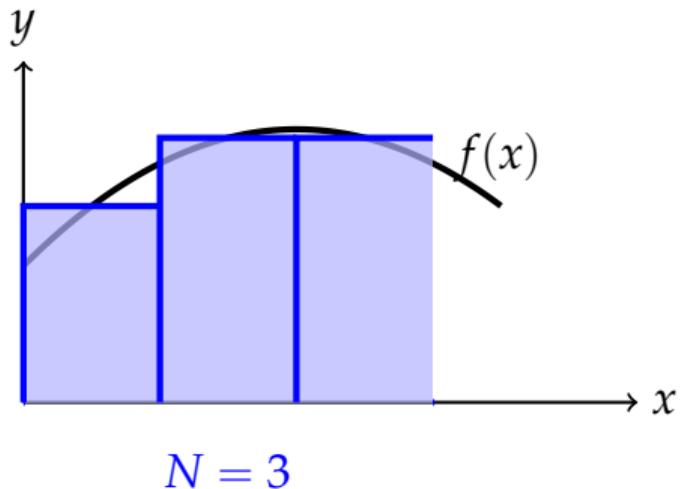
## Intuition: Constructing Functions (Step 3)

- Any continuous function can be approximated by a sum of rectangles ([Riemann Sum](#)).



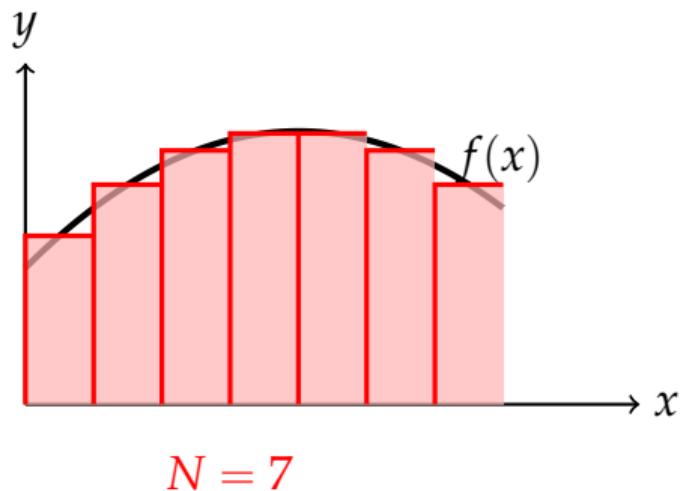
## Intuition: Constructing Functions (Step 3)

- Any continuous function can be approximated by a sum of rectangles ([Riemann Sum](#)).
- Since a NN can build any rectangle, it can sum them up to approximate the function.



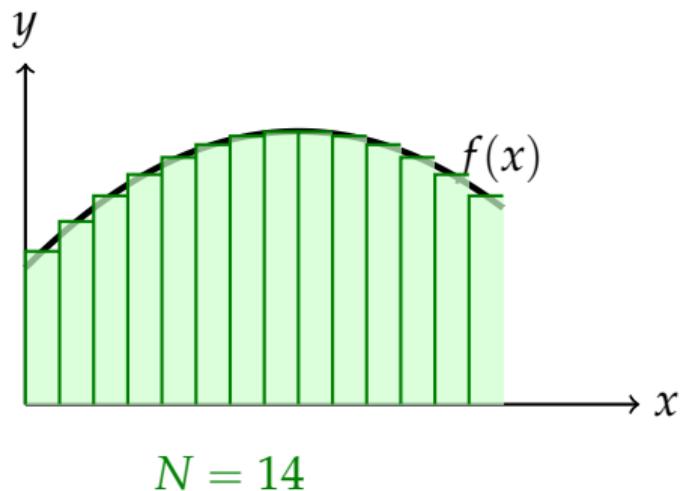
## Intuition: Constructing Functions (Step 3)

- Any continuous function can be approximated by a sum of rectangles ([Riemann Sum](#)).
- Since a NN can build any rectangle, it can sum them up to approximate the function.
- **Conclusion:** A 1-hidden layer NN with  $2K$  nodes can approximate a function represented by  $K$  Riemann rectangles.



## Intuition: Constructing Functions (Step 3)

- Any continuous function can be approximated by a sum of rectangles ([Riemann Sum](#)).
- Since a NN can build any rectangle, it can sum them up to approximate the function.
- **Conclusion:** A 1-hidden layer NN with  $2K$  nodes can approximate a function represented by  $K$  Riemann rectangles.

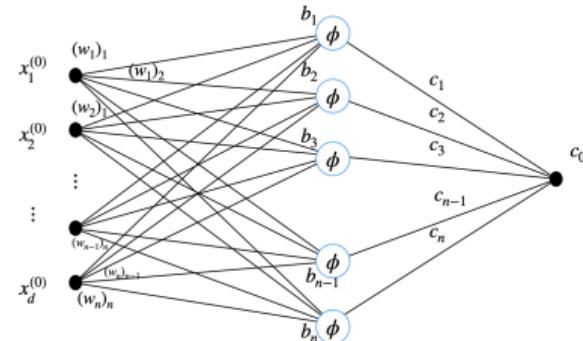


# Conclusion in the 1D Case

- 1 Approximate the function in the Riemann sense by a sum of  $k$  rectangles.
- 2 Represent each rectangle using **two nodes** in the hidden layer.
- 3 Compute the sum of all nodes in the hidden layer (considering appropriate weights and signs) to get the final output.

## Remarks:

- The same intuition applies to any sigmoid-like function.
- This is an intuitive explanation, not a quantitative one.
- The weights  $w$  must be large to mimic the step function.

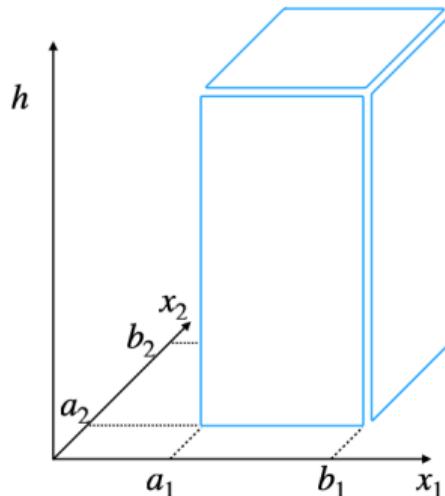


**Takeaway 1D Construction:** Any continuous 1D function can be approximated using a one-hidden-layer network with  $2k$  sigmoid neurons (for  $k$  rectangles). **The key:** sigmoids approximate step functions!

# Larger Dimension: $d = 2$

## Extension to 2D:

- Approximate the function by 2D box functions.
- Constructing a 2D box is harder than 1D. We need to intersect two infinite "strips".



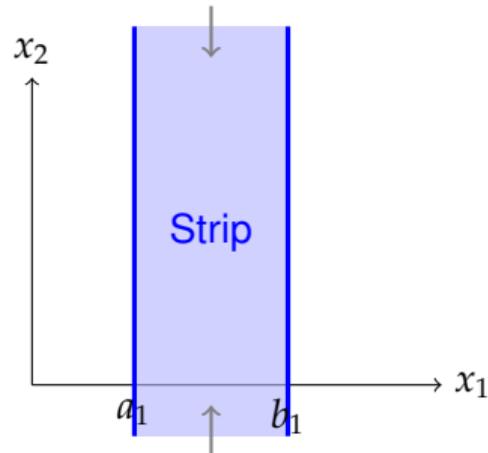
# Step 1: Create Infinite Strip in $x_1$ Direction

$$(x_1, x_2) \mapsto \phi(w(x_1 - a_1)) - \phi(w(x_1 - b_1)) \quad (2)$$

**Result:**

- Value  $\approx 1$  when  $a_1 \leq x_1 \leq b_1$
- Value  $\approx 0$  otherwise
- Unbounded in  $x_2$  direction

**Network:** 2 sigmoid neurons in hidden layer.



## Step 2: Create Infinite Strip in $x_2$ Direction

$$(x_1, x_2) \mapsto \phi(w(x_2 - a_2)) - \phi(w(x_2 - b_2)) \quad (3)$$

**Result:**

- Value  $\approx 1$  when  $a_2 \leq x_2 \leq b_2$
- Value  $\approx 0$  otherwise
- Unbounded in  $x_1$  direction

**Network:** 2 more sigmoid neurons (total: 4 in hidden layer).



## Step 3: Combine Strips to Form a Cross

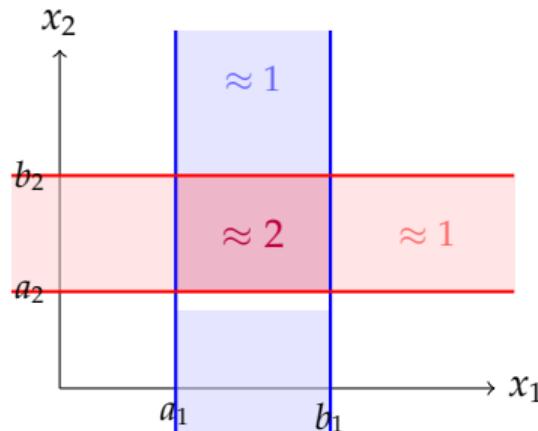
$$(x_1, x_2) \mapsto \phi(w(x_1 - a_1)) - \phi(w(x_1 - b_1)) + \phi(w(x_2 - a_2)) - \phi(w(x_2 - b_2)) \quad (4)$$

**Result:**

- Value  $\approx 2$  in rectangle (both strips active).
- Value  $\approx 1$  in arms (one strip active).
- Value  $\approx 0$  elsewhere.

**Problem:** Unwanted infinite arms! **Solution:**

Threshold at  $c \in (1, 2]$ .



## Step 4: Thresholding (Construction vs. Existence)

### Thresholding Operation:

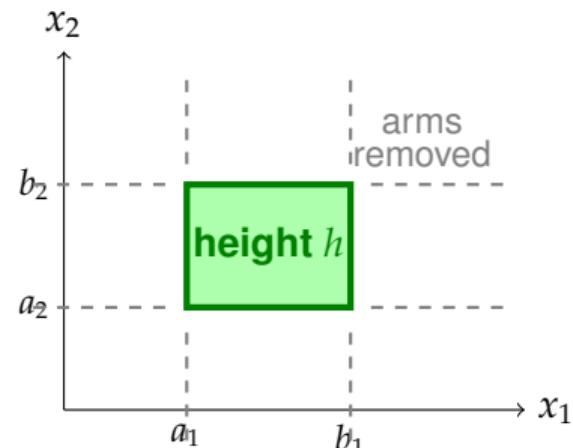
- Compose cross with  $\mathbf{1}_{x>c}$  for  $c \in (1, 2]$ .
- We approximate  $\mathbf{1}_{x>c}$  using a steep sigmoid:

$$\phi(w \cdot (x - c))$$

- Implementation:** This effectively uses a **2nd hidden layer** to clean up the shape.

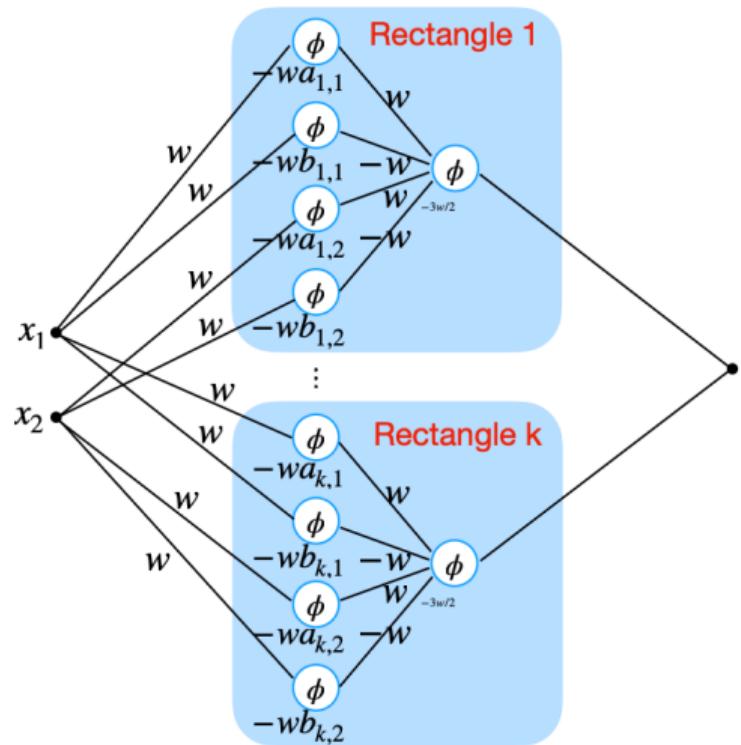
**Result:** Clean rectangle!

$$f(\mathbf{x}) \approx \begin{cases} h & \text{if } \mathbf{x} \in \text{Box} \\ 0 & \text{otherwise} \end{cases}$$

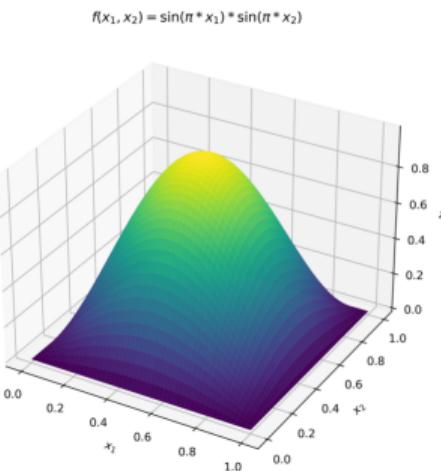
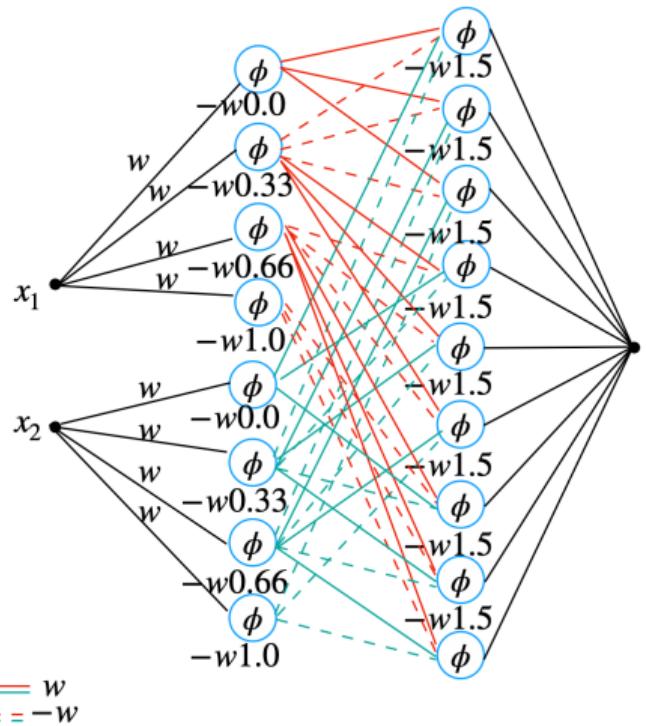


# Deeper Networks: Efficiency

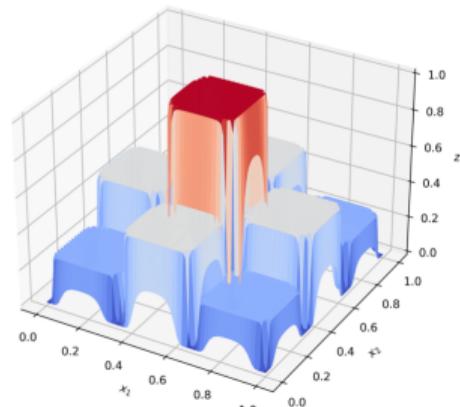
- While 1 layer is *sufficient* (Barron), it may require exponentially many neurons.
- Deep networks (2+ layers) can construct complex shapes (like clean rectangles) more efficiently.
- We can implement multiple rectangles with 2 hidden layers to approximate complex 2D functions.



# Example: Approximate $f(x) = \sin(\pi x_1) \cdot \sin(\pi x_2)$ with 9 Rectangles

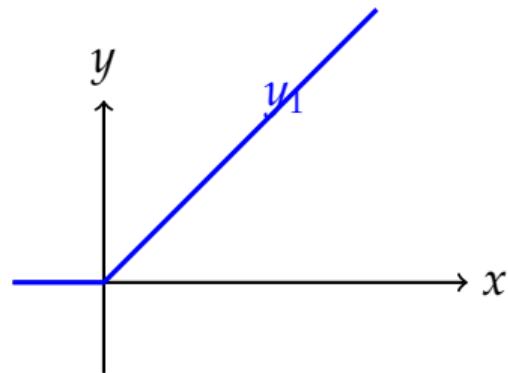


Neural Network Approximation



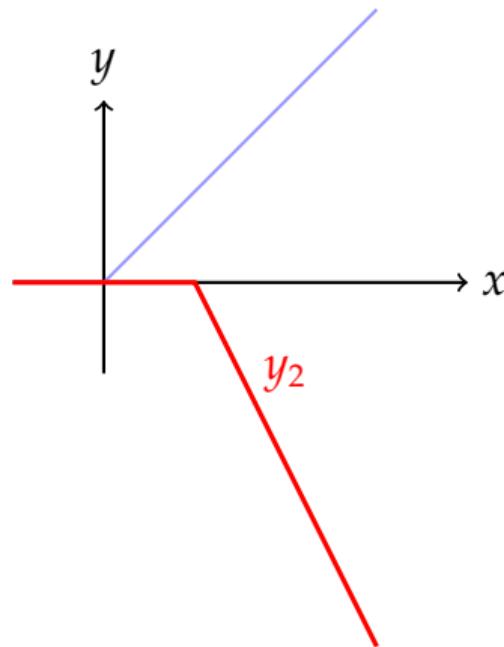
# Approximation with ReLUs

- What about ReLU?  $\phi(x) = \max(0, x)$ .



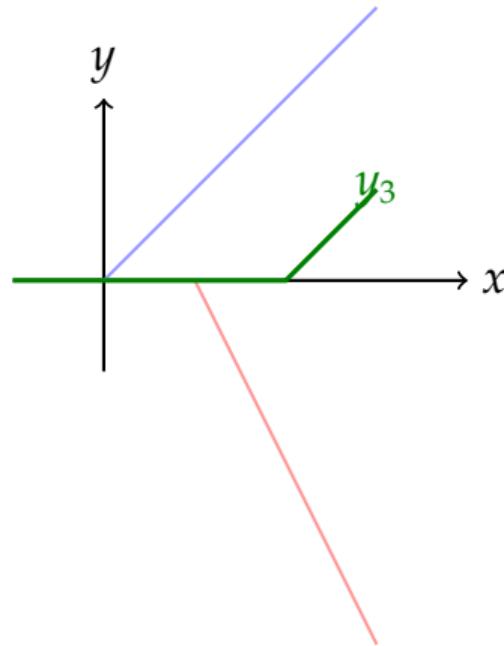
# Approximation with ReLUs

- What about ReLU?  $\phi(x) = \max(0, x)$ .
- A ReLU unit introduces a "kink" (change in slope).



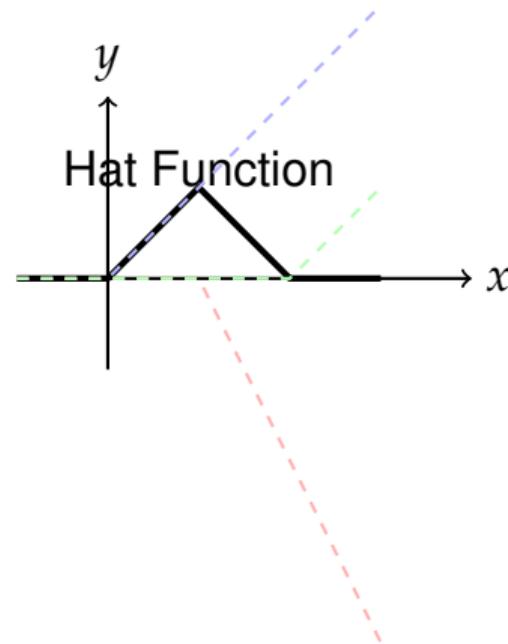
# Approximation with ReLUs

- What about ReLU?  $\phi(x) = \max(0, x)$ .
- A ReLU unit introduces a "kink" (change in slope).
- Sum of ReLUs → Piecewise Linear (PWL) Function.



# Approximation with ReLUs

- What about ReLU?  $\phi(x) = \max(0, x)$ .
- A ReLU unit introduces a "kink" (change in slope).
- Sum of ReLUs  $\rightarrow$  Piecewise Linear (PWL) Function.
- Any continuous function can be approximated by a piecewise linear function.



# $\ell_2$ vs $\ell_\infty$ Approximations

- **$\ell_2$ -norm (average approx.):**
  - Measures the average deviation between the true function and the approximation.
  - Barron's result guarantees this for 1-layer nets.

# $\ell_2$ vs $\ell_\infty$ Approximations

- **$\ell_2$ -norm (average approx.):**
  - Measures the average deviation between the true function and the approximation.
  - Barron's result guarantees this for 1-layer nets.
- **$\ell_\infty$ -norm (pointwise approx.):**
  - Measures the **maximum** error at any point:  $\sup_x |f(x) - \hat{f}(x)|$ .
  - Ensures uniformly small error across the domain.

# $\ell_2$ vs $\ell_\infty$ Approximations

- **$\ell_2$ -norm (average approx.):**
  - Measures the average deviation between the true function and the approximation.
  - Barron's result guarantees this for 1-layer nets.
- **$\ell_\infty$ -norm (pointwise approx.):**
  - Measures the **maximum** error at any point:  $\sup_x |f(x) - \hat{f}(x)|$ .
  - Ensures uniformly small error across the domain.
- $\Rightarrow$  **Question:** Can we get  $\ell_\infty$  approximation with Neural Networks?  
Yes, by exploiting Piecewise Linear (PWL) functions!

# $\ell_\infty$ Approximation with Piecewise Linear Functions

**Def:** piecewise linear (PWL) function:

$$q(x) = \sum_{i=1}^m (a_i x + b_i) \mathbf{1}_{r_{i-1} \leq x < r_i},$$

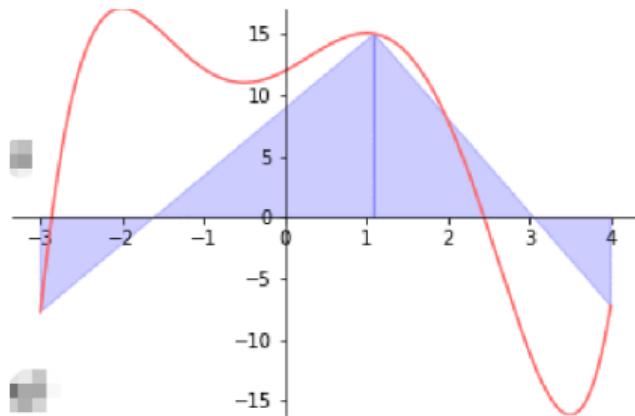
with  $a_i r_i + b_i = a_{i+1} r_i + b_{i+1}$  (continuity).

**$\ell_\infty$ -approximation result** (Shekhtman, 1982):

Let  $f$  be a continuous function on  $[c, d]$ . For all  $\varepsilon > 0$ , there exists a piecewise linear function  $q$  such that:

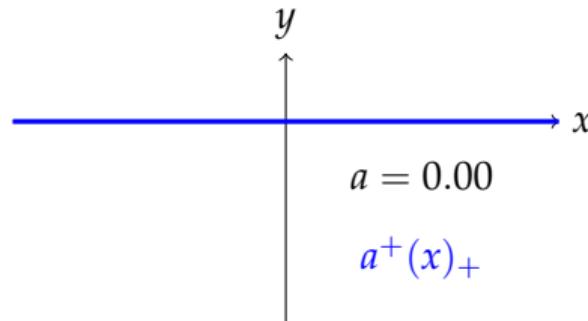
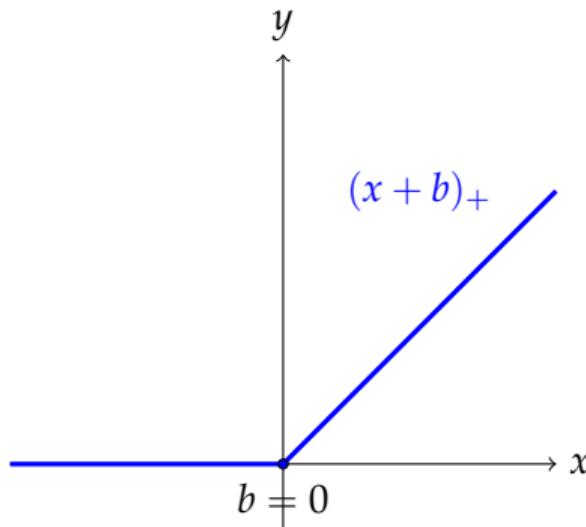
$$\sup_{x \in [c, d]} |f(x) - q(x)| \leq \varepsilon$$

⇒ Goal: Approximate PWL functions with a ReLU NN.



# ReLU Activation and PWL Functions

**Role of the bias and weight in ReLU function**  $(ax + b)_+ = \max\{0, ax + b\}$ :



**The bias  $b$  determines the position of the kink.**

**The weight  $a$  determines the slope.**

A linear combination of ReLUs  $\sum_{i=1}^m \tilde{a}_i(x - \tilde{b}_i)_+$  is a piecewise linear function.

# Piecewise Linear Functions Can Be Written as Combination of ReLU

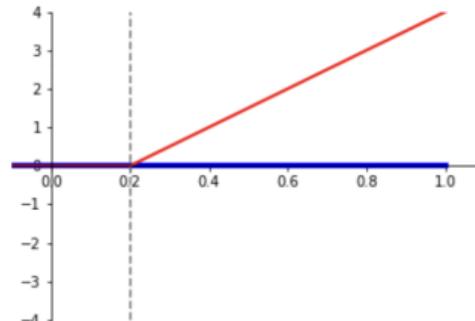
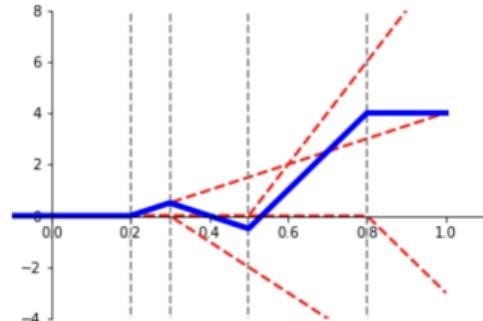
**Claim 1:** Any PWL  $q(x)$  can be rewritten as:

$$q(x) = \tilde{a}_1 x + \tilde{b}_1 + \sum_{i=2}^m \tilde{a}_i (x - \tilde{b}_i)_+$$

where  $\tilde{a}_1 = a_1, \tilde{b}_1 = b_1, a_i = \sum_{j=1}^i \tilde{a}_j$  and  $\tilde{b}_i = r_{i-1}$ .

**Proof sketch:**

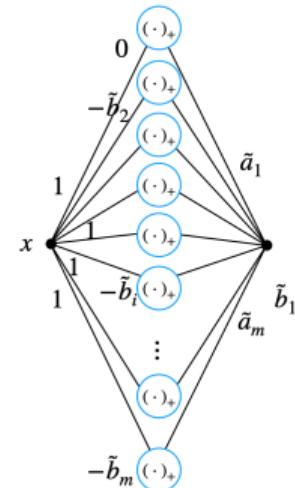
- How do we get a new segment with slope  $a$  starting at  $r > \max(\tilde{b}_j)$ ?
- Get the kink at  $r$  by setting  $\tilde{b}_{i+1} = r$  and slope by additionally canceling existing slope i.e.  $\tilde{a}_{i+1} = a - \sum \tilde{a}_i$ .



# PWL as One-Hidden Layer NN with ReLU

**Claim 2:**  $q$  can be implemented as a one-hidden-layer NN with ReLU activation. Each term corresponds to one node:

- Bias  $-\tilde{b}_i$
- Output weight  $\tilde{a}_i$



The term  $\tilde{a}_1x + \tilde{b}_1$  corresponds to one node:

- Bias  $\tilde{b}_1$ : bias of the output node.
- Term  $\tilde{a}_1x = \tilde{a}_1(x)_+$ , assuming  $x \geq 0$  (if input domain bounded).

**Key Result** Any  $\ell_\infty$  approximation via PWL can be implemented with a one-hidden-layer ReLU network!

# Proof of the Equivalent Formulation

**Show:** The two representations are equivalent:

$$q(x) = \sum_{i=1}^m (a_i x + b_i) \mathbf{1}_{r_{i-1} \leq x < r_i}$$

$$r(x) = \tilde{a}_1 x + \tilde{b}_1 + \sum_{i=2}^m \tilde{a}_i (x - \tilde{b}_i)_+$$

where  $\tilde{a}_1 = a_1$ ,  $\tilde{b}_1 = b_1$ ,  $a_i = \sum_{j=1}^i \tilde{a}_j$ , and  $\tilde{b}_i = r_{i-1}$ .

- **For**  $x \in [0, r_1]$ :

$$(\tilde{a}_1, \tilde{b}_1) = (a_1, b_1) \implies q(x) = a_1 x + b_1 = \tilde{a}_1 x + \tilde{b}_1 = r(x)$$

(because  $\tilde{b}_2 = r_1$ , so no ReLU terms activate)

- **For**  $x \in [r_1, r_2]$ :

$$\begin{aligned} r(x) &= \tilde{a}_1 x + \tilde{b}_1 + \tilde{a}_2 (x - r_1)_+ \\ &= a_1 x + b_1 + (a_2 - a_1)(x - r_1) = a_2 x + b_1 - (a_2 - a_1)r_1 = q(x) \end{aligned}$$

Note:  $r'(x) = a_2$  and  $r(r_1) = q(r_1)$ , so  $r(x) = q(x)$  for  $x \in [r_1, r_2]$ .

# Proof by Induction

**Induction Hypothesis:** Assume  $r(x) = q(x)$  for  $x \in [0, r_{i-1}]$ .

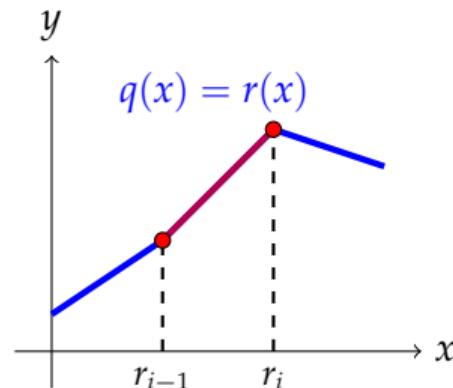
**Induction Step:** Show  $r(x) = q(x)$  for  $x \in [r_{i-1}, r_i]$ .

For  $x \in [r_{i-1}, r_i]$ :

$$\begin{aligned} r(x) &= \tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^m \tilde{a}_j (x - \tilde{b}_j)_+ \\ &= \tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^i \tilde{a}_j (x - \tilde{b}_j) = \sum_{j=1}^i \tilde{a}_j x + \dots \end{aligned}$$

**Thus:**

- $r'(x) = \sum_{j=1}^i \tilde{a}_j = a_i$  (correct slope)
  - $r(r_{i-1}) = q(r_{i-1})$  (correct starting point)
- $\implies r(x) = q(x)$  for  $x \in [r_{i-1}, r_i]$



The segment in red shows the interval  $[r_{i-1}, r_i]$ .

**Key insight:** Two affine functions with the same starting point and slope are identical.

# Table of Contents

- ① Introduction and Motivation
- ② Neural Networks: Basic Structure
- ③ Universal Approximation Theorem
- ④ Training Neural Networks

# The Learning Problem

Given a dataset  $\mathcal{D} = \{(\mathbf{x}_n, y_n)\}_{n=1}^N$ , we want to find parameters  $\theta = \{\mathbf{W}^{(l)}, \mathbf{b}^{(l)}\}$  that minimize a loss function:

$$\min_{\theta} \mathcal{L}(\theta) = \frac{1}{N} \sum_{n=1}^N \ell(y_n, f(\mathbf{x}_n; \theta))$$

- **Regression:** MSE Loss  $\ell(y, \hat{y}) = (y - \hat{y})^2$ .
- **Classification:** Cross-Entropy Loss  $\ell(y, \hat{y}) = -\sum_k \mathbf{1}(y = k) \log \hat{y}_k$ .

# Backpropagation

- We use Gradient Descent (or SGD) to optimize  $\theta$ :

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

**Takeaway** **Backpropagation** makes training deep networks practical by efficiently computing gradients via the chain rule. Modern frameworks provide **automatic differentiation**, making implementation straightforward!

# Backpropagation

- We use Gradient Descent (or SGD) to optimize  $\theta$ :

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

- How to compute gradients  $\nabla_{\theta} \mathcal{L}$  for a deep net?

**Takeaway** **Backpropagation** makes training deep networks practical by efficiently computing gradients via the chain rule. Modern frameworks provide **automatic differentiation**, making implementation straightforward!

# Backpropagation

- We use Gradient Descent (or SGD) to optimize  $\theta$ :

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

- How to compute gradients  $\nabla_{\theta} \mathcal{L}$  for a deep net?
- **Backpropagation:** Efficient application of the Chain Rule.
  - **Forward Pass:** Compute activations layer by layer.
  - **Backward Pass:** Propagate error signals ( $\delta$ ) from output to input.

**Takeaway** **Backpropagation** makes training deep networks practical by efficiently computing gradients via the chain rule. Modern frameworks provide **automatic differentiation**, making implementation straightforward!

# Backpropagation

- We use [Gradient Descent](#) (or SGD) to optimize  $\theta$ :

$$\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}(\theta)$$

- How to compute gradients  $\nabla_{\theta} \mathcal{L}$  for a deep net?
- **Backpropagation:** Efficient application of the [Chain Rule](#).
  - **Forward Pass:** Compute activations layer by layer.
  - **Backward Pass:** Propagate error signals ( $\delta$ ) from output to input.
- Modern frameworks (PyTorch, TensorFlow) do this automatically ([Autodiff](#)).

**Takeaway** **Backpropagation** makes training deep networks practical by efficiently computing gradients via the chain rule. Modern frameworks provide **automatic differentiation**, making implementation straightforward!

# Summary

- **Structure:** NNs are composed of layers of neurons.

Linear Transform( $\mathbf{Wx} + \mathbf{b}$ ) → Non-linear Activation( $\phi$ ) → ...

Non-linearity is essential for learning complex patterns!

- **Representation Power:** NNs are Universal Approximators.

- **Barron's Theorem:** One hidden layer can approximate any smooth function with  $\mathcal{O}(1/n)$  error.
- **1D:** Sigmoids create bumps → Riemann sums.
- **2D:** Deeper networks (2 layers) are more efficient for constructing clean shapes.
- **ReLU:** Equivalent to piecewise linear functions ( $\ell_\infty$  approx).

- **Training:** Optimized via SGD using gradients computed by Backpropagation.

**Final Takeaway** Neural networks combine **universal approximation capability** with **efficient training** via backpropagation. This makes them the most powerful and practical machine learning models for complex, high-dimensional data!